Surface Orders, Cyclic Time and the Arithmetic Spectral Bridge

Dessins, resolutions, Auslander–Reiten theory, quantum models and verifiable research

Amelie Schreiber

Research development edition

November 3, 2026

Reader's guide and mathematical status

This book develops the surface-order and cyclic-time constructions of the supplied *Emergent cyclic time* manuscripts together with the subsequent projective-resolution, arithmetic and spectral work. The originals remain immutable. Their final unconditional GRH claims do not survive the identified arithmetic and operator-theoretic gaps. Here every construction has a specified category, coefficient field and determinant target. Established results, conditional implications, computed experiments and open comparisons play different roles.

The first two parts give the shared article constructions and their expanded proofs. The next part develops entire AR components, corrects the face-unfolding and Hom-control claims, and relates valid meshes to derived triangles and determinantal semi-invariants. The physics part retains the original topics while replacing invalid positivity, hardware and continuum inferences by proved finite statements or explicit research models. Computational chapters supply trained finite neural architectures, higher-genus dessins and a local interactive companion. The final comparison problem remains an independent arithmetic theorem, not a consequence of the pictures or learned parameters.

A reader may start with cellular dessins and the node, then projective resolutions, then the AR atlas. The more advanced arithmetic chapters use character sums and Fredholm determinants. The physics chapters can be read with finite-dimensional linear algebra and elementary quantum mechanics. Each computational claim has a script, seed and stated backend. The genus-two and genus-three sheet views are projections; the torus views are topological embeddings, not asserted uniformizing metrics.

The provenance manifest records the source hashes and chapter inventory. Its location is `book/source-provenance.json`. The review records state the audit coverage. The supplied 2053 file is byte-identical to the pinned source called COMPLETED in the source clone; the other supplied COMPLETED file is a shorter, different version. This naming discrepancy is retained in the provenance record. No filename certifies a theorem.

The companion is served on this machine at `http://127.0.0.1:8765`. It contains adjustable dessins, category-labelled AR diagrams, quantum carrier circuits, and saved training curves. The private repository remains the authoritative source. Selected lemmas are checked in Lean4; a full formalization of this entire book has not yet been completed.

Contents

I Constructions and statements **1**

1 Cellular dessins and their infinite orders **2**

2 Projective resolutions, dual dessins and determinants **4**
2.1 Auxiliary staircase resolutions over the order 5
2.2 A monotone quantum comparison with exact finite ranks 6

3 Determinant semi-invariants and Schur constructions **8**
3.1 Young tableaux and a cyclic representation stratum 9

4 Lattice almost-split sequences and genuine diamond quotients **11**
4.1 Polynomial strings, completed lattices and the Fukaya category 12

5 Leavitt relations, the constrained Fock algebra and spectrum **14**
5.1 A graph presentation and an actual pullback 15
5.2 A separate diagonal-fiber C* completion retaining the pairing 16

6 Arithmetic attachment, descent and the remaining comparison **18**
6.1 A constructive arithmetic benchmark on the theta dessin 19
6.2 A genuine cellular carrier, and what still fails for the order 19
6.3 Arbitrary projective actions: an obstruction and a genuine determinant repair 20
6.4 Tame descent: a surviving order and a sharper obstruction 22
6.5 Whole components, derived endpoints and determinant relations 24

7 A genuine theta covering and its categorical limits **25**

8 Verification and research frontier **27**

II Proofs and calculations **29**

9 Expanded dessin, resolution and derived calculations **30**
9.1 The cycle order in coordinates 30
9.2 The two-arm complex and its calculations 30

10 Semi-invariants, tableaux and explicit intertwiners **32**
10.1 Derived information retained by a presentation 32
10.2 One Young operator and one LR decomposition in coordinates 32
10.3 Two precise functorial limitations 33

11 Full lattice and diamond diagrams **35**
11.1 The nodal exact diagram 35
11.2 Chebyshev geometry and the exact quotient 35
11.3 Nine actual strings and all maps 36
11.4 The pure-injective quotient and what its dimension means 37

12 The operator extension in detail **38**
12.1 A corrected cycle lift and the integer boundary map . . . 38
12.2 Spectral domain and arithmetic trace . . . 38

13 Abel determinant estimates and the arithmetic program **40**
13.1 Precisely how classical and recent sources enter . . . 41
13.2 A concrete sequence of remaining constructions . . . 42

14 Graph pullbacks and completions with specified fibers **43**
14.1 The complete Fock graph presentation . . . 43
14.2 A genuine application of the supplied pullback theorem . . . 43
14.3 Diagonal residue obstruction and the corrected pullback . . . 44
14.4 Generation, faithfulness and the exact failure of extension . . . 45
14.5 Invariant basis number and the Leavitt comparison . . . 46

15 String finiteness and the smooth quotient surface **47**

16 Arithmetic character sums and the cellular determinant **49**
16.1 The uniform calculation and its exceptional cases . . . 49
16.2 Deck descent and the actual cellular basis . . . 50
16.3 An explicit chain lift and determinant cancellation . . . 50
16.4 A decisive restriction at inert primes . . . 51
16.5 The remaining order and global spectral obligations . . . 52

17 The theta order interface beyond tensor insertion **53**
17.1 Normal forms and the special fiber . . . 53
17.2 All projectives and all linear endomorphisms . . . 54
17.3 An actual integral determinant functor of multiplicity one . . . 55
17.4 The exact rotation obstruction and its allowed escape . . . 55

18 Explicit tame descent and the full-deck coefficient test **57**
18.1 The action on genuine formal series . . . 57
18.2 A projective symmetry survives on the endomorphism order . . . 58
18.3 Determinants, Schur weights and the boundary correction . . . 58
18.4 Why the local Jacobi plane does not descend by this repair . . . 59
18.5 The full orbit is the normalization, and arbitrary kernels still fail . . . 59

19 Literal formal and computational coverage **62**

20 Explicit covering calculations and derived transport **64**

III AR components, derived categories and Hom-control **66**

21 AR geometry and determinant relations **67**
21.1 What an AR picture specifies . . . 67
21.2 The corrected elliptic-theta face example . . . 68
21.3 Hom control and coverings . . . 70
21.4 Homogeneous tubes . . . 71
21.5 Nodal lattice AR geometry . . . 72
21.6 An HLS component in a dessin quotient . . . 73
21.7 Derived AR triangles . . . 75
21.8 Ringel's wind-wheel tile . . . 76
21.9 Meshes and semi-invariants . . . 77
21.10Physical and arithmetic scope . . . 79

22 Coverings, equivariant AR theory and graded surfaces 80
22.1 Reading the three new sources together . . . 80
22.2 A cover that really preserves the relations . . . 80
22.3 Equivariant modules, not an unproved folding rule . . . 81
22.4 Derived and symplectic data needed for a determinant . . . 81
22.5 From a Hom decomposition to a semi-invariant calculation . . . 82

IV Quantum models, trained renormalization and physical realizations 83

23 Positive evolution, clocks and holonomy 84
23.1 States and two different logarithms . . . 84
23.2 Total positivity, wiring and physical positivity . . . 85
23.3 Weyl clocks, Fourier phases and twisted successors . . . 85
23.4 Infinite clocks: rotations and odometers are distinct limits . . . 86
23.5 Holonomy, DG compatibility and persistence . . . 87

24 Entanglement and trained quantum renormalization 88
24.1 An entangled pair and the finite first law . . . 88
24.2 Correlation geometry and the gravitational research question . . . 89
24.3 Pooling is a channel: an exact trainable example . . . 89
24.4 Trained QGCN and QGAT on ribbon graphs . . . 90
24.5 Associative memory with a correct measurement model . . . 91
24.6 Evidence for trained quantum renormalization . . . 91

25 Noise, molecular spins and hyperbolic-code architectures 92
25.1 Ancillas, noise and measurement resources . . . 92
25.2 The peptide-spin evidence and its limits . . . 93
25.3 Control calculations: rotations and exchange . . . 93
25.4 A proposed molecular-design program . . . 94
25.5 Hyperbolic kagome geometry and the assembly problem . . . 95
25.6 CSS codes from boundary maps . . . 95
25.7 The arithmetic comparison is still an additional task . . . 96

V Computational and learned constructions 98

26 Higher-genus dessins and reproducible finite carriers 99
26.1 A hyperelliptic family beyond the torus . . . 99
26.2 Larger genera and infinite hyperbolic windows . . . 100
26.3 What a three-dimensional view can claim . . . 101
26.4 Finite clocks and cross-backend agreement . . . 101

27 Trained quantum coarse-graining architectures 103
27.1 Exact teacher data from Ising decimation . . . 103
27.2 Shared quantum graph convolution . . . 104
27.3 QCNN compression with a genuine pooling channel . . . 104
27.4 A larger genus-four nine-qubit example . . . 105
27.5 Quantum-feature graph attention . . . 105
27.6 Photonic Gaussian pooling . . . 105
27.7 Reproduction and limits of the evidence . . . 106

28 Reading an architecture from its tensors 108
28.1 Three distinct kinds of evidence . . . 108
28.2 Shapes explain the operation . . . 108
28.3 Quantum and Gaussian inspection . . . 109
28.4 A reproducible way to explore . . . 109

29 ToricGT and mirror interfaces **111**
29.1 A toric chart inside an actual dessin representation stratum . . . 111
29.2 Tropical charts and exactly what their certificates prove . . . 112
29.3 The quantum phase block and its cocycle . . . 113
29.4 A quantum–classical ToricGT extension with auditable interfaces . . . 114
29.5 Mirror symmetry supplies category comparisons, not a spectral shortcut . . . 115
29.6 What would make the learned component relevant to GRH? . . . 116
29.7 An executed hybrid spectral experiment . . . 116

30 Persistence and explicit resolutions **118**
30.1 The polynomial ring encodes the two commuting directions . . . 118
30.2 A complete three-grade example . . . 119
30.3 A bifiltered CW realization and the syzygies it creates . . . 120
30.4 Schur modules organize the generators and the first syzygies . . . 121
30.5 Fitting data, exactness and specialization . . . 121
30.6 A correct finite space of complexes and its incidence model . . . 122
30.7 Coefficients in a genuine infinite dessin order . . . 123
30.8 Quantum data, monotonicity and the next usable interface . . . 123
30.9 The executable staircase and quantum laboratory . . . 124
30.10What the polynomial Lean resolution actually verifies . . . 125

31 The remaining arithmetic theorem and a verifiable research loop **126**
31.1 A precise positive spectral target . . . 126
31.2 A surviving short-row reduction . . . 126
31.3 Formal coverage and the next research step . . . 127

Part I

Constructions and statements

Chapter 1

Cellular dessins and their infinite orders

We start with the surface-order constructions of Schreiber [1, 2], retaining the original sources unchanged. The generic-resolution source [3] is background, not a substitute for exactness of a particular complex. All primary algebra examples below come from explicit clean Belyi maps. Finite-dimensional objects occur as modules or specified quotient strata of those infinite algebras, not as replacements for them. The initial arithmetic target is a primitive nonprincipal Dirichlet L-function. Finite-order Hecke characters, Dedekind zeta functions and general Artin representations require separate comparison theorems; Artin holomorphy is not assumed.

A clean dessin on a compact oriented surface is specified by a nonempty finite dart set H, a fixed-point-free involution α, and a permutation σ such that $\langle \alpha, \sigma \rangle$ is transitive. Contract the white vertices of degree two. The resulting cellular graph has vertices the σ-orbits, edges the α-orbits, and faces the ϕ-orbits, where

$$\tau = \sigma^{-1}, \qquad \phi = \alpha\sigma^{-1}, \qquad 2 - 2g = \#\operatorname{cyc}(\sigma) - |H|/2 + \#\operatorname{cyc}(\phi).$$

Write $[h] = \{h, \alpha h\}$. Our quiver has vertices these edges and arrows $a_h : [h] \to [\tau h]$. Paths are composed in travel order: ab traverses a first. Set

$$A = kQ/(a_h a_{\phi h} : h \in H), \quad J = A_{>0}, \quad \Lambda = \varprojlim A/J^m.$$

Each vertex has two incoming and two outgoing arrows. Each arrow has precisely one permitted successor $a_{\tau h}$ and one forbidden successor $a_{\phi h}$; these are distinct because α has no fixed points. The same holds for predecessors. This is the *locally gentle* condition. Every permitted cycle repeats indefinitely, so A is infinite dimensional. We reserve "gentle" without qualification for the convention additionally imposing finite dimension.

Proposition 1.1 (The actual split order)**.** *For a permitted cycle c of length ℓ, let B_c be the hereditary cycle order in $M_\ell(k[[u_c]])$: diagonal and forward matrix entries are in $k[[u_c]]$, and entries crossing the chosen cycle cut are in $u_c k[[u_c]]$. Then Λ is the suborder of $B = \prod_c B_c$ whose two diagonal constant terms labelled $h, \alpha h$ agree. With $R = k[[u]]$ acting by $u \mapsto (u_c I)_c$, it is finite free of rank $\sum_c \ell(c)^2$; the conductor contains uB.*

Proof. In a cycle order a path of length less than ℓ is the corresponding matrix unit, multiplied by u_c exactly when it crosses the cut; another full turn multiplies it by u_c. Thus completed paths in each permitted branch give all entries of B_c. An arrow belongs to only that branch,

whereas $e_{[h]}$ has diagonal entries at both occurrences $h, \alpha h$. Hence the only condition on branchwise formal coefficients is the displayed equality of diagonal constants. This proves both injectivity and surjectivity of the path-normal-form identification. Multiplication respects it: switching branches across a paired edge is precisely a forbidden pair. The quotient B/Λ is killed by u, so the submodule has the same generic rank as B. It is torsion free and finite over the DVR R, hence free of that rank. The path and u-adic topologies agree: $uA \subset J$, and if L is the longest permitted cycle, any path of length at least $L+1$ is u times a shorter nonempty path. Thus $J^{L+1} \subset uA$. This also justifies the completion used in the identification. □

The path-degree variable q below is not u: one full turn has path degree $\ell(c)$ but u-degree one. Formal, Laurent and operator-norm completions will likewise remain distinct.

Here are five explicit global inputs. Permutations list darts starting at zero; $E = |H|/2$.

Surface and Belyi map β	σ	α	$(g, \operatorname{rank}_R \Lambda)$
$\mathbb{P}^1$, $4z(1-z)$	(0)(1)	(01)	$(0, 2)$
$\mathbb{P}^1$, $(z+1)^2/(4z)$	(01)	(01)	$(0, 4)$
$y^2 = x^3 - 1$, x^3	(012)(345)	(03)(14)(25)	$(1, 18)$
$\mathbb{P}^1$, $4z^3/(1+z^3)^2$	(012)(345)	(03)(15)(24)	$(0, 18)$
$y^2 = x^3 - x$, x^2	(0123)	(02)(13)	$(1, 16)$

The first is the nodal interval: $A = k[x, y]/(xy)$ and $\Lambda = k[[x, y]]/(xy)$. The second has $A = k\langle x, y\rangle/(x^2, y^2)$, still infinite dimensional. The two theta dessins have the same two permitted three-cycles but different face cycles, respectively (051324) and (04)(13)(25). The elliptic square has one permitted and one forbidden four-cycle. For example x has pole order two on the cubic $y^2 = x^3 - 1$; x^3 has ramification $(3, 3)$ above zero, $(2, 2, 2)$ above one since $x^3 - 1 = y^2$, and (6) above infinity. The rational theta map has passport $(3, 3), (2, 2, 2), (2, 2, 2)$. These ramification computations, and the analogous direct zeros/poles in the other rows, verify cellular Belyi inputs rather than merely abstract quivers.

Chapter 2

Projective resolutions, dual dessins and determinants

Use right modules $P_i = e_i A$, $S_i = P_i/P_i J$. Internal shift $M\langle n\rangle$ places its old generator in degree n. Let $\mathcal{R}_n(i)$ be the *indexing words* of length n from i with every adjacent pair forbidden, including e_i for $n = 0$. Such a word is usually zero in A; it still labels a summand.

Theorem 2.1 (Exact resolutions)**.** *There is a minimal graded projective resolution, with $P_0(i) = e_i A$ and for $n \geq 1$,*

$$P_n(i) = \bigoplus_{p=a_1\cdots a_n \in \mathcal{R}_n(i)} e_{t(a_n)} A\langle n\rangle, \qquad d_n(x)_p = a_n x \text{ in the summand } a_1 \cdots a_{n-1}.$$

The augmentation $P_0(i) \to S_i$ is the quotient. For the completed order the identical finite-rank projective terms and formal-series maps resolve the residue simple. Ext and Tor against residue simples count forbidden words with the prescribed endpoints. In the graded case Ext has matching internal shift n.

Proof. The maps are right linear and degree zero. Adjacent differentials multiply by a forbidden pair. The augmentation kernel is the direct sum of aA over initial arrows. For any arrow a, multiplication $e_{t(a)}A \to e_{s(a)}A$ kills exactly the permitted paths beginning with its forbidden successor $f(a)$. All other basis paths survive to distinct basis paths. Thus its kernel is $f(a)A$. For a fixed prefix, different last arrows have disjoint surviving supports; there is no cancellation between summands. This proves exactness inductively. For formal series the same assertion holds coefficient by coefficient and identifies the entire closed kernel, not just its dense polynomial subspace. Terms are idempotent projectives and all entries are arrows. Consequently applying Hom or tensor with a residue simple kills every differential, proving minimality and the Ext/Tor count. In A minimality is graded; it does not require that J be its ordinary Jacobson radical. □

The two arms never terminate in this closed class. Their labelled patterns follow ϕ-cycles. An ungraded positive syzygy tail can have a shorter period by permutation of summands: the nodal interval has labelled period two but unordered tail period one; the elliptic theta has labelled period six and unordered tail period three. A simple need not be isomorphic to a positive syzygy; for the node it is R-torsion, whereas the positive syzygies are lattices.

Let $C_{ij}(q) = \sum_{d\geq 0} \dim_k(e_i A_d e_j) q^d$ and $E_{ij}(q) = \sum_{n\geq 0} (-1)^n \dim_k \operatorname{Ext}^n_A(S_i, S_j) q^n$. Degreewise Euler characteristic of the resolution gives $EC = I$: every coefficient is finite, because a summand in homological degree n starts in internal degree n. Constant terms are I, so the inverse is two-sided. Already signed $E(q)$ must not be evaluated at $-q$ again.

Proposition 2.2 (Dual dessin and cycle determinant)**.** *The complementary-relation algebra is another closed dessin algebra, with $\alpha^{\#} = \alpha$, $\sigma^{\#} = \phi^{-1}$, permitted successor ϕ, and forbidden successor τ. It exchanges graph vertices and faces and preserves genus. Its Cartan matrix satisfies $E(q) = C_{A^{\#}}(-q)$, with the opposite convention if required for Yoneda composition. For independent commuting arrow weights w_a,*

$$\det C_A(\mathbf{w}) = \frac{\prod_{f\in \mathrm{cyc}(\phi)}(1-(-1)^{|f|}w(f))}{\prod_{c\in \mathrm{cyc}(\sigma)}(1-w(c))}.$$

Proof. $\alpha(\sigma^{\#})^{-1} = \alpha\phi = \tau$, so the claimed relations and face data follow directly; the connectedness group is unchanged. To prove the determinant, put $S_{v,a} = w_a$ when $s(a) = v$, and $T_{a,v} = 1$ when $t(a) = v$, zero otherwise. On arrow states let $U_{a,b} = w_b$ for permitted adjacency and $V_{a,b} = w_b$ for forbidden adjacency. Then $TS = U + V$ and the path sum is $C = I + S(I - U)^{-1}T$. Sylvester's identity gives $\det C = \det(I + V)/\det(I - U)$. Each transition is a weighted permutation, whose cycle determinant is respectively $1 - (-1)^{\ell}w(c)$ and $1 - w(c)$. Everything holds in formal power series, hence as rational functions. This supplies a direct proof of the Bessenrodt–Holm formula in the closed dessin class [4]. □

The five univariate determinants are, in the order above,

$$\frac{1+q}{1-q},\quad \frac{1+q}{1-q},\quad \frac{1-q^6}{(1-q^3)^2},\quad \frac{(1-q^2)^3}{(1-q^3)^2},\quad 1.$$

The last cancellation does not imply finite global dimension: its forbidden resolution is infinite. These rational functions have cycle poles/zeros, not arithmetic prime Euler factors.

For clarity, a complex is a sequence with $d^2 = 0$, its homotopy category identifies maps differing by $dh + hd$, and $D(A)$ inverts maps inducing cohomology isomorphisms. A perfect complex is represented by a bounded complex of finite projectives. Our infinite simple resolutions compute derived Hom and tensor but are generally not perfect. The two-term presentation $[P_1 \to P_0]$ is perfect even if its kernel is nonzero. Quadratic linearity is a Koszul statement; it is not by itself an equivalence between all completed derived categories. Likewise a finite-dimensional gentle graded-curve theorem [6] cannot simply be transferred to these closed infinite orders. Here the surface walks explicitly compute the resolution endpoints: τ turns round a graph vertex and ϕ follows a face. For longer string objects all maps can be computed from the occurrence matrices; a universal geodesic-intersection classification for this completed category remains a separate theorem.

2.1 Auxiliary staircase resolutions over the order

The two filtration parameters below are not dessin arrows. Let $R = k[x, y]$ for a field k, and let $I = (g_0, \dots, g_{m-1})$, $m \geq 2$, where $g_i = x^{a_i}y^{b_i}$,

$$a_0 > \cdots > a_{m-1} = 0, \qquad 0 = b_0 < \cdots < b_{m-1}.$$

The positive outer exponents make R/I finite length. Write $R(-a, -b)$ for a free generator born in bidegree (a, b).

Proposition 2.3 (Adjacent syzygies)**.** *A minimal bigraded free resolution of $M = R/I$ is*

$$0 \longrightarrow \bigoplus_{i=0}^{m-2} R(-a_i, -b_{i+1}) \xrightarrow{d_2} \bigoplus_{i=0}^{m-1} R(-a_i, -b_i) \xrightarrow{d_1} R \longrightarrow M \longrightarrow 0,$$

where $d_1(e_i) = g_i$ *and* $d_2(f_i) = -y^{b_{i+1}-b_i}e_i + x^{a_i-a_{i+1}}e_{i+1}$. *Deleting row* i *from* d_2 *gives determinant* $(-1)^i g_i$. *The Hilbert series numerator is* $1 - \sum_i s^{a_i}t^{b_i} + \sum_{i=0}^{m-2} s^{a_i}t^{b_{i+1}}$.

Proof. Each column maps to zero. In a relation $\sum h_i g_i = 0$, all terms except the last contain $x^{a_{m-2}}$. Coprimality with $y^{b_{m-1}}$ implies $h_{m-1} = x^{a_{m-2}}q$. Subtract $qd_2(f_{m-2})$ from the coefficient vector; its last component vanishes. Factor $x^{a_{m-2}}$ from the remaining generators. They form a smaller staircase with the same strict inequalities, ending in a pure y-power. Induction expresses the remaining relation in the preceding adjacent columns. The base of one generator has no relations, since R is a domain. Hence $\ker d_1 = \operatorname{im} d_2$. If $d_2(q) = 0$, its first row gives $-y^{b_1}q_0 = 0$; successive rows then force all $q_i = 0$. This proves injectivity. Every differential entry has positive degree, so reduction modulo (x, y) kills the differentials and the resolution is minimal. In the row-deleted matrix the unique nonzero permutation term uses the negative y entries before row i and the positive x entries after it. Their exponents telescope to b_i, a_i, respectively, and the sign is $(-1)^i$. Finally take the degreewise Euler characteristic; the free summands have Hilbert denominator $(1 - s)(1 - t)$. □

For $I = (x^2, xy, y^2)$ the underlying sequence is also kernel checked in Lean 4 over the actual ring MvPolynomial(Fin 2, k). The verification proves the kernel/image equalities, the stated ideal image, surjective quotient augmentation, projective free terms, and vanishing of both induced Hom differentials to the constant-coefficient residue field. It uses regular multiplication on the free modules, distinct from the residue action on k. General staircases, graded minimality and the surface-order coefficient cone remain separately scoped prose arguments.

At grade $\nu \in \mathbb{N}^2$, M_ν has dimension one exactly when no generator degree is componentwise at most ν. Its comparable-grade map has rank one exactly when both endpoint grades satisfy this condition. A CW realization takes one loop and attaches a disk once along it at each generator degree. The active indices form an interval: the x inequality selects a final interval, the y inequality an initial interval. With c active disks, the cellular map $k^c \to k$ has every entry one. Thus $\dim H_1 = \mathbf{1}_{c=0}$ and $\dim H_2 = \max(c - 1, 0)$, the latter counting active adjacent syzygies. This auxiliary CW complex is not the dessin surface.

For the actual nodal order $\Lambda = k[[\xi, \eta]]/(\xi\eta)$ put $u = \xi + \eta$ and $E = \Lambda/u^2\Lambda$. The basis $(1, \xi, u, u\xi)$ of E satisfies $\xi^2 = u\xi$, $u^2 = 0$, $\eta = u - \xi$. Tensoring the resolution above with $0 \to \Lambda \xrightarrow{u^2} \Lambda \to E \to 0$ gives, over $\mathcal{B} = \Lambda[x, y]$, free ranks $(1, m + 1, 2m - 1, m - 1)$. Indeed the scalar resolution stays exact after tensoring over the field; its degree-zero module is free over Λ on the standard monomials. Multiplication by u^2 is injective there and on all free terms, so the long exact homology sequence of its mapping cone gives positive-degree homology zero and degree-zero module $E \otimes_k M$. In contrast, as an R-module this object has ranks $4(1, m, m - 1)$ in its minimal resolution. Neither finite-length E nor these parameter rings replace the infinite parent dessin algebra or its lattice category.

2.2 A monotone quantum comparison with exact finite ranks

Proposition 2.4 (One common channel)**.** *Let* T *be a positive trace-preserving linear map between finite complex matrix algebras. For density matrices* ρ, σ, *their trace distance* $d(\rho, \sigma) = \frac{1}{2}\|\rho - \sigma\|_1$ *satisfies* $d(T\rho, T\sigma) \le d(\rho, \sigma)$. *For finitely many labelled states, keep fixed vertex birth grades* e_i. *At* (r, e) *include vertices with* $e_i \le e$ *and edges with distance at most* r, *and take the flag complex. Then* T *gives an identity-on-labels inclusion from the input to the output bifiltration, natural in* (r, e).

Proof. Write the Hermitian difference as $A = A_+ - A_-$ with $A_\pm \ge 0$. Positivity, the trace norm triangle inequality and trace preservation give $\|TA\|_1 \le \|TA_+\|_1 + \|TA_-\|_1 = \operatorname{tr} A_+ + \operatorname{tr} A_- = \|A\|_1$. Thus every input edge is an output edge; the fixed births preserve vertices, and all cliques are

preserved. Increasing either parameter preserves the simplices, and every map is the same label inclusion, so the squares commute. Complete positivity is sufficient but is not needed for this norm argument. □

For a finite inclusion $K \subset L$, work over a field and embed a basis Z_K of source one-cycles into the target edge space. If B_L is the matrix of target two-boundaries, the induced H_1 rank is

$$\rho_1(K, L) = \operatorname{rank}[B_L \; Z_K] - \operatorname{rank} B_L.$$

Its image is $(Z_K + \operatorname{im} B_L)/\operatorname{im} B_L$; source boundaries already map into target boundaries. This proves the formula, including empty complexes. Computing a flag two-skeleton suffices for H_0, H_1, but does not compute the full flag complex's H_2. Integer ceilings of scaled exact distances preserve the inclusion. For numerical distances, frozen integer labels define an exact combinatorial diagram only after their inequality is checked; this is not an interval certificate of the underlying quantum distances. None of these finite ranks identifies a completed arithmetic determinant or proves GRH.

Chapter 3

Determinant semi-invariants and Schur constructions

Assume k algebraically closed of characteristic zero for invariant-ring statements. A finite-dimensional representation of the *infinite* A has spaces N_i, arrow matrices $N(a)$ satisfying the monomial relations, and base-change group $G = \prod_i \mathrm{GL}(N_i)$. With column vectors, $N(a) \mapsto g_{t(a)} N(a) g_{s(a)}^{-1}$. Let $P_0 = \bigoplus P_i^{a_i}$, $P_1 = \bigoplus P_i^{b_i}$, and $f : P_1 \to P_0$ be a finite path-matrix map. When $\sum a_i \dim N_i = \sum b_i \dim N_i$, define

$$c_f(N) = \det D_f(N), \qquad D_f(N) = \mathrm{Hom}_A(f, N).$$

It is intrinsically a determinant-line section; chosen bases give this scalar.

Proposition 3.1 (Weights and exact homological scope)**.** $c_f(gN) = \prod_i \det(g_i)^{b_i - a_i} c_f(N)$. *Restricted determinants of this form span all semi-invariants on a closed G-stable relation or nilpotence locus. Set $M = \operatorname{coker} f$. On the square locus,*

$$c_f(N) \neq 0 \quad \Longleftrightarrow \quad \mathrm{Hom}(M, N) = 0 \quad \Longleftrightarrow \quad \mathrm{Hom}(M, N) = \mathrm{Ext}^1(M, N) = 0.$$

If f is injective, $\operatorname{coker} D_f = \mathrm{Ext}^1(M, N)$ *and higher Ext vanish. Without injectivity neither that cokernel equality nor higher vanishing follows.*

Proof. Naturality gives $D_f(gN) = G_1 D_f(N) G_0^{-1}$, where G_j repeats the relevant g_i. Take determinants. The contragredient coordinate-ring action has the inverse character. Schofield–van den Bergh's spanning theorem allows oriented cycles [36, Theorem2.3]. Exactness of invariants for reductive G in characteristic zero lifts a fixed-weight function from any closed stable locus to the ambient representation space; restrict the spanning determinants. Injectivity of the path map need not survive that restriction. Finally $\ker D_f = \mathrm{Hom}(M, N)$. Writing $K = \operatorname{im} f$, the exact sequence with projective P_0 expresses $\mathrm{Ext}^1(M, N)$ as $\mathrm{Hom}(K, N)/\operatorname{im} \mathrm{Hom}(P_0, N)$; the epimorphism $P_1 \to K$ embeds this quotient in $\operatorname{coker} D_f$. A square finite matrix with zero kernel has zero cokernel. If f is injective it is the actual length-one resolution, giving the stronger assertion. □

For the node, $f = \times x$, $M = \Lambda/(x)$, $N = \Lambda/(x, y^2)$ give $\dim \operatorname{coker} D_f = 2$ but $\dim \mathrm{Ext}^1(M, N) = 1$. For the elliptic square name arrows so that $a, b : 0 \to 1$, $c, d : 1 \to 0$, permitted cycle *acbd* and forbidden cycle *adbc*. With $f = a$ and $N(a) = 1$, $N(b) = N(c) = N(d) = 0$, its Hom matrix is [1], yet the subsequent forbidden resolution gives $\mathrm{Ext}^{2,3}(M, N) = k$. Thus nonvanishing is not all-degree derived orthogonality.

Proposition 3.2 (What completion does to invariants)**.** *A continuous d-dimensional representation of Λ kills J^d. On the reduced nilpotent representation variety of fixed dimension vector every invariant regular function is constant. Nonzero-weight semi-invariants need not be constant.*

Proof. Continuity with discrete finite-dimensional target kills some J^m; its Loewy filtration has at most d nonzero strict steps. Choose vertexwise complements to this filtration. An arrow sends a layer p into layers $q \geq p+1$. Scaling layer p by z^p conjugates each component by z^{q-p}. At $z = 0$ the representation tends to the zero-arrow representation. Every invariant therefore has the same value at all points; reducedness and the Nullstellensatz give the ring assertion. In the square example the scalar $N(a)$ varies on a nilpotent locus and has weight g_1/g_0, so semi-invariants survive. Formal entries in a projective map are evaluated on this locus by truncating modulo J^d, a legitimate finite polynomial operation. □

On the node, $u = x+y$ is regular and A is free of rank two over $k[u]$. The polynomial presentation $\times(u-\lambda)$ for $\lambda \neq 0$ is injective and its determinant on N is $\det(u_N - \lambda I)$, the characteristic polynomial of u_N up to the sign $(-1)^d$. But $u-\lambda$ is a unit in Λ, with its geometric-series inverse, so those modules do not extend continuously to the completion. An R-free lattice and its u-nilpotent quotients are different categories. On a lattice a separately specified endomorphism F has a meaningful relative characteristic polynomial over R; the scalar base action uI alone does not create Frobenius.

3.1 Young tableaux and a cyclic representation stratum

For a nonempty partition $\lambda \vdash m$, choose a numbering T of its boxes. On $V^{\otimes m}$ let $Y_T = R_T C_T$, the unnormalized row sum followed by the signed column sum. The classical identity $Y_T^2 = h_\lambda Y_T$, where h_λ is the hook product, makes Y_T/h_λ an idempotent in characteristic zero. Its image is $S_\lambda(V)$, one Schur copy, rather than the whole symmetric-group isotypic component [31, 32]. Its character is s_λ, computed by semistandard tableaux (weak rows, strict columns), and

$$\dim S_\lambda(k^d) = \prod_{(i,j)\in\lambda} \frac{d+j-i}{h_{ij}}.$$

Since $f^{\otimes m}$ commutes with the place permutations, it induces $S_\lambda(f)$ and respects composition and zero maps. Applying this at every vertex and arrow therefore gives a genuine representation of the monomial dessin algebra: $S_\lambda(N(b))S_\lambda(N(a)) = S_\lambda(0) = 0$ for a forbidden pair. The empty partition does not have this property: it sends zero arrow maps to the identity of k.

The Littlewood–Richardson rule gives natural intertwiners

$$S_\lambda(V) \otimes S_\mu(V) \cong \bigoplus_\nu S_\nu(V)^{\oplus c^\nu_{\lambda\mu}}.$$

Here $c^\nu_{\lambda\mu}$ counts semistandard tableaux of skew shape ν/λ and content μ whose row-reading word is Yamanouchi. Naturality makes the same decompositions intertwine all the arrows. The tensor operation is vertexwise with synchronous arrow maps, not $\otimes_A$ and not a Hopf-algebra tensor structure. There is no automatic tensor unit preserving these zero relations. These functors are not additive or exact; they cannot simply be applied degreewise to a projective resolution. They also need not preserve the fixed R-action on lattices.

In the elliptic theta set one of the two permitted three-cycle branches to zero. The other has arbitrary maps $X_1 : V_0 \to V_2$, $X_2 : V_2 \to V_1$, $X_3 : V_1 \to V_0$, with all spaces of dimension d. This is a closed representation stratum of our infinite dessin algebra. Put $M = X_3X_2X_1$.

Theorem 3.3 (Cyclic invariants and a local positive kernel)**.** *The invariant ring on this stratum is $k[c_1(M), \dots, c_d(M)]$, where c_j are characteristic coefficients. Its total-arrow-degree Hilbert series is $\prod_{j=1}^{d}(1-t^{3j})^{-1}$. Common multidegree (m,m,m) has basis $\chi_\lambda(M) = \operatorname{tr} S_\lambda(M)$, $\lambda \vdash m$, $\ell(\lambda) \le d$; other multidegrees have no invariants. Products obey the Littlewood–Richardson rule. Over $\mathbb{C}$, if $\|M_i\| \le R_0$ and $0 \le r$ and $rR_0^2 < 1$, the kernel*

$$K_r(M_i, M_j) = \det(I - rM_i \otimes \overline{M_j})^{-1} = \sum_\lambda r^{|\lambda|} \chi_\lambda(M_i)\overline{\chi_\lambda(M_j)}$$

is absolutely convergent and positive semidefinite.

Proof. Cycle-trace generation for quivers, including relation quotients, is the Le Bruyn–Procesi theorem [29, 30]. Every cycle here is a power of the full three-cycle, up to cyclic rotation. In characteristic zero Newton identities give the generators c_j. They are algebraically independent: take $X_1 = X_2 = I$ and arbitrary X_3. Scalar changes at the three vertices force equal arrow multidegrees. Restriction to that same section identifies degree (m,m,m) with degree-m symmetric polynomials in eigenvalues, whose Schur basis and product rule give the assertion. For the kernel use the ordinary Cauchy identity, first as a formal identity in eigenvalues. On Schur modules the operator norm of $S_\lambda(M)$ is at most $R_0^{|\lambda|}$, so its trace is bounded by $\dim S_\lambda(\mathbb{C}^d) R_0^{|\lambda|}$. The dimension Cauchy identity bounds the absolute series by $(1 - rR_0^2)^{-d^2}$. Each summand is a rank-one Gram kernel, proving positivity by absolute summation. □

The term for the empty partition is a constant scalar feature, not an allowed module-valued tensor unit. In dimension two,

$$(\operatorname{tr} M)^2 = \operatorname{tr} \operatorname{Sym}^2 M + \det M, \qquad \operatorname{tr} \operatorname{Sym}^2 M = \tfrac{1}{2}((\operatorname{tr} M)^2 + \operatorname{tr} M^2).$$

More generally $\det(I - zM) = \sum_j (-z)^j \operatorname{tr} \bigwedge^j M$ and its reciprocal is $\sum_m z^m \operatorname{tr} \operatorname{Sym}^m M$. These are actual determinant/Schur connections, not identifications with an arithmetic L-function. On continuous finite-length Λ-modules the monodromy is nilpotent and all positive-degree characters vanish. The displayed Cauchy positivity is local representation-theoretic positivity, not the global Weil form.

Chapter 4

Lattice almost-split sequences and genuine diamond quotients

For the interval dessin let $X = (x)$, $Y = (y)$ in $\Lambda = k[[x, y]]/(xy)$, $R = k[[u]]$, $u = x + y$. A lattice is a finitely generated Λ-module free over R.

Theorem 4.1 (The nodal lattice calculation)**.** *Every lattice is $\Lambda^r \oplus X^a \oplus Y^b$. The almost-split sequences in this lattice category are*

$$0 \to Y \hookrightarrow \Lambda \xrightarrow{\times x} X \to 0, \qquad 0 \to X \hookrightarrow \Lambda \xrightarrow{\times y} Y \to 0.$$

In particular $\tau X = Y$, $\tau Y = X$. Multiplication by u between two copies of either row, followed by reduction modulo u, gives an actual nonzero exact three-by-three diagram.

Proof. Normalize M *inside its total quotient module*: $\overline{M} = BM$, $B = k[[x]] \times k[[y]]$. The conductor gives $u\overline{M} \subset M \subset \overline{M}$. Both branch projections of $W = M/u\overline{M} \subset k^p \oplus k^q$ are onto. Intersect W with the two factors; its remaining quotient identifies their two residual quotients. Separate bases for these pure factors and the common quotient put W into pure first lines, pure second lines and diagonal pairs. Lift branch bases over the DVRs. These are respectively the summands X, Y, Λ. Their endomorphism rings are local.

Exactness of the first sequence is $\operatorname{ann} x = Y$. It does not split: a map $X \to \Lambda$ lands in $\operatorname{ann} y = X$, and its composite with $\times x$ is a nonunit in $\operatorname{End} X = k[[x]]$. Every map $\Lambda \to X$ factors through this epimorphism, $\operatorname{Hom}(Y, X) = 0$, and every nonunit endomorphism of X is divisible by x and factors through $X \hookrightarrow \Lambda \to X$. Classification handles all finite sums; a unit X component would be a retraction. This is right almost splitness. The dual argument for nonsections out of Y proves left almost splitness. Endomorphisms fixing either structure map have constant term one, hence are units, proving minimality. Swap x, y for the other sequence. Finally all three terms are R-free. Multiplication by u is injective and reduction preserves exactness, proving the diagram assertion. □

The lattice AR quiver has arrows $Y \to \Lambda \to X$ and $X \to \Lambda \to Y$; the mesh composites are zero. Its projective vertex Λ is not the endpoint of another nonsplit AR sequence. AR meshes do not mean that all rows and columns of any drawn grid are short exact.

For a second, explicit family fix $n \geq 3$ and the rational Belyi map

$$\beta_n(z) = \frac{1}{1 - T_n(z)^2} = \frac{1}{(1 - z^2)U_{n-1}(z)^2},$$

where T_n, U_{n-1} are Chebyshev polynomials. It has passport $(2n), (2^n), (2^{n-1}, 1, 1)$. Its contracted dual graph is a path, identifying its genus-zero dessin. The associated infinite algebra has vertices

$1, \dots, n$, endpoint loops l, r, opposite arrows $a_i : i \to i+1$, $b_i : i+1 \to i$, and precisely the relations $l^2, r^2, a_i b_i, b_i a_i$. Its single permitted cycle is $l a_1 \cdots a_{n-1} r b_{n-1} \cdots b_1$; the order rank is $(2n)^2$.

Retain one arrow at each edge in an orientation Ω, and kill its opposite. The quotient is exactly the string algebra $H_n(\Omega)$ of Huang–Lin–Su [46]; the killed-arrow ideal contains J^{n+2}, so it is also an actual quotient of Λ_n. Inflation identifies its module category with the full abelian quotient-annihilator subcategory $\mathcal{C}_\Omega$ of Λ_n-modules. This is generally not extension closed in the whole order category. The imported AR property is intrinsic to $\mathcal{C}_\Omega$, not automatically an ambient lattice or finite-length AR property. The extended edition gives nine explicit string modules, all twelve maps and the four exact meshes from the supplied diamond resource, now situated in this actual dessin quotient.

The other supplied resources also have a precise home. For the interval order the subcategory annihilated by J^2 is $\operatorname{Mod} B_0$, $B_0 = \Lambda/J^2$. This is a domestic string quotient, not itself a gentle algebra. Inflation is full and preserves purity, and $\operatorname{Zg} B_0$ is a closed subspace of $\operatorname{Zg} \Lambda$. Ringel's infinite-word/pure-injective program and its proof by Puninski–Prest apply there [43, 44]; the isolated band xy^{-1} gives Krull–Gabriel and relative Cantor–Bendixson dimension two by [45, Theorem10.19]. This is not an ambient rank calculation for all Λ-modules. The parameter in $x = J_2, y = \lambda J_2$ survives even though all cycle traces vanish; the pp condition $\exists v\ (z = xv \ \wedge \ (y - \lambda x)v = 0)$ detects it. These modules are u-torsion and are not nonzero order lattices. Infinite-string and locally finite AR theorems must retain their distinct finiteness hypotheses; a surviving permitted cycle has infinitely many paths at a vertex.

4.1 Polynomial strings, completed lattices and the Fukaya category

Crawley–Boevey's theorem [38, Theorems1.1–1.4] applies directly to the infinite polynomial dessin algebra: its finitely controlled modules decompose into strings and finite-dimensional bands; its pointwise artinian modules additionally allow primitive injective bands. For our finite quiver, these conditions mean finitely generated and artinian respectively. The words may be infinite, but their modules have algebraic direct-sum bases. This is not a classification of every module or of arbitrary formal completions.

In fact no nonzero completed-order lattice, restricted to A, meets either finiteness condition. The polynomial A is finite over $k[u]$: remove full turns from sufficiently long paths as in Proposition 1.1. If an R-free lattice were finitely generated over A, it would be finite torsion free, hence free, over the PID $k[u]$. Multiplication by $1 - u$ would then fail to be onto, whereas it is a unit over $R = k[[u]]$. Also $M \supsetneq uM \supsetneq u^2M \supsetneq \cdots$ excludes artinianity. For a concrete indecomposable outside the ordinary string/band list take $k[[x]]$ with $y = 0$: its endomorphisms are multiplication by formal series, so its endomorphism ring is local. Its x is injective and nonsurjective, while $1 - x$ is invertible; these properties exclude every pure-x ordinary string or band. The extended edition gives the full word argument.

The Fukaya connection requires equally precise categories. Lekili–Polishchuk explicitly allow infinite-dimensional locally gentle algebras, but their equivalence is for *perfect dg modules* of a homologically smooth algebra, with specified grading, line field and stops [66, Definition3.1.1, Lemma3.1.3, Theorem3.2.2]. Every nonempty closed dessin here has both permitted and forbidden cycles, so its polynomial algebra is neither proper nor smooth by their criterion. This does not exclude appropriate singular or completed variants; it prevents a direct use of that particular theorem. Infinite word, graded curve and geodesic are not interchangeable without a surface realization and boundary/grading data.

There is a genuine positive example inside our setting: the polynomial interval quotient $A/(y) =$

$k[x]$ is still infinite dimensional, is smooth, and has no forbidden cycle. Its perfect category has the partially wrapped annulus model with one stop total. In degree zero $\mathrm{Perf}(k[x]) = D^b(\mathrm{mod}_{\mathrm{fg}}\, k[x])$. Inflation of this quotient is fully faithful on ordinary modules, but not on derived categories: $\mathrm{Ext}^2_A(k[x], k[x]) = k$, whereas $\mathrm{Ext}^2_{k[x]}(k[x], k[x]) = 0$. The former follows from the alternating y, x resolution. Thus this quotient bridge does not identify the whole closed dessin's perfect category with the annulus category.

Finally, Leavitt path algebras *can* fail invariant basis number, but need not. The permitted-state Laurent algebras here have finite-dimensional matrix quotients, and therefore invariant basis number. The polynomial and completed dessin algebras have their vertex-simple quotients, giving the same conclusion. The complex Leavitt algebra of G_α below is a dense *-subalgebra of its graph C*-algebra, on that specified graph; it is not the original gentle algebra with ordinary CK1 imposed.

Chapter 5

Leavitt relations, the constrained Fock algebra and spectrum

Work now over $\mathbb{C}$. Ordinary Leavitt and graph C* constructions [47, 48] cannot simply impose our original forbidden products: if $ab = 0$ and CK1 says $a^*a = p_{t(a)}$, then $0 = a^*ab = b$. Every arrow is a forbidden successor, hence every arrow and then every vertex is zero. CK2 is not needed. The ordinary graph-Toeplitz CK1 setting has the same collapse.

The valid permitted-arrow-state graph instead consists of the disjoint directed σ-cycles. Its Leavitt algebra is $\bigoplus_c M_{|c|}(k[z, z^{-1}])$, and its graph C* algebra is $\bigoplus_c M_{|c|}(C(\mathbb{T}))$. These forget the α-gluing; both theta dessins give the same answer. They are not the completed original order. Nor is the ordinary Toeplitz algebra of a cycle of length $\ell > 1$ equal to $M_\ell(\mathcal{T})$: their K_0 ranks are ℓ and one respectively [50].

Let $\mathcal{H} = \ell^2(\mathcal{P})$ for all permitted finite paths, including vacua e_v. Define $L_a\delta_p = \delta_{ap}$ when the prefix is permitted and composable, and zero otherwise. Let P_v project onto paths starting at v, $Q_a = L_aL_a^*$ and v_v onto the vacuum. Then

$$L_aL_{\phi(a)} = 0, \quad L_a^*L_a = P_{t(a)} - Q_{\phi(a)} = v_{t(a)} + Q_{\tau(a)}, \quad P_v - \sum_{s(a)=v} Q_a = v_v.$$

These follow on basis vectors. The different initial projection avoids CK1 collapse. Write $\mathcal{T}_A = C^*(P_v, L_a)$, a constrained-path Fock C* algebra in the monomial/subproduct setting [51, 52]. The polynomial A acts faithfully: applying a finite path sum to the sum of vacua recovers its coefficients. This representation need not extend to all formal Λ series: $\sum_{m\geq 0} c^m$ along a cycle would send its vacuum to a nonsquare-summable sum.

Theorem 5.1 (Compact extension and incidence index)**.** *Let $\mathcal{H}_v$ be the terminal-path sector. There is an exact sequence*

$$0 \longrightarrow \bigoplus_{v\in Q_0} \mathcal{K}(\mathcal{H}_v) \longrightarrow \mathcal{T}_A \longrightarrow \bigoplus_{c\in \mathrm{cyc}(\sigma)} M_{|c|}(C(\mathbb{T})) \longrightarrow 0.$$

With $\delta[u] = [1 - W^*W] - [1 - WW^*]$, its boundary map is $\delta(c) = -\sum_{h\in c} e_{[h]}$, minus the unsigned incidence map B_Γ of the cellular graph. Thus, for its E edges,

	$K_0(\mathcal{T}_A)$	$K_1(\mathcal{T}_A)$
Γ *bipartite*	$\mathbb{Z}^{E+1}$	$\mathbb{Z}$
Γ *nonbipartite*	$\mathbb{Z}^E \oplus \mathbb{Z}/2$	0.

Proof. All prefixes preserve the terminal sectors. For paths p, q ending at v, $L_p v_v L_q^* = |\delta_p\rangle\langle\delta_q|$. Therefore the vacuum ideal is exactly the displayed sum of compacts: it contains these matrix units and that sum is itself an ideal. In the quotient, L_a has initial projection $Q_{\tau(a)}$ and range Q_a, a graph CK family for the permitted-state cycles. The Q_a have infinite rank and nonzero quotient images; the length gauge action descends. Gauge-invariant uniqueness [48, Theorem2.1] identifies the quotient as claimed.

For a cycle c based at arrow a, lift its corner unitary by $W = L_c Q_a + (1 - Q_a)$. Then $W^*W = 1$ and $1 - WW^* = Q_a - L_c Q_a L_c^*$ projects onto the first $|c|$ paths starting with a. Their terminal vertices count exactly its incident darts, including twice for a loop. This proves the index formula. The six-term exact sequence gives $K_1 = \ker B_\Gamma$ and $0 \to \operatorname{coker} B_\Gamma \to K_0 \to \mathbb{Z}^{\#\,\mathrm{cyc}\,\sigma} \to 0$. The free last term splits abstractly. Spanning-tree integer elimination gives Smith entries 1 repeated $V - 1$ times and one zero in the bipartite case; an odd cycle instead adds the entry two. This proves the groups. □

The nodal interval has $K_0 = \mathbb{Z}^2$, $K_1 = \mathbb{Z}$; the single-loop sphere has $\mathbb{Z} \oplus \mathbb{Z}/2$, zero. Both theta dessins have $\mathbb{Z}^4$, $\mathbb{Z}$. Coarse K-theory does not recover genus or the arithmetic factors.

5.1 A graph presentation and an actual pullback

The Fock algebra itself has a graph presentation, on a different graph from the original gentle quiver. Let G_α have vertices q_h for $h \in H$ and sinks b_e for $e \in Q_0$, and edges

$$s_h : q_h \to q_{\tau h}, \qquad t_h : q_h \to b_{[\tau h]}.$$

Proposition 5.2. *$C^*(G_\alpha) \cong \mathcal{T}_A$, by $q_h \mapsto Q_{a_h}$, $b_e \mapsto v_e$, $s_h \mapsto L_{a_h} Q_{a_{\tau h}}$ and $t_h \mapsto L_{a_h} v_{[\tau h]}$. The integer graph relation at h is $q_h - q_{\tau h} - b_{[\tau h]}$. Eliminating these relations recovers $K_1 = \ker B_\Gamma$ and $K_0 \cong \operatorname{coker} B_\Gamma \oplus \mathbb{Z}^{\#\,\mathrm{cyc}\,\sigma}$.*

Proof. The two edges at q_h have orthogonal initial pieces and orthogonal ranges summing to Q_{a_h}, by the Fock initial-projection identity. All vertex projections are nonzero. The length gauge acts with degree one on both edge families, so gauge-invariant uniqueness applies. They generate since $L_{a_h} = s_h + t_h$ and $P_e = v_e + \sum_{[h]=e} Q_{a_h}$. In the integer matrix a kernel coefficient is constant around each permitted cycle; its sink equations are exactly $B_\Gamma c = 0$. In the cokernel, eliminate all but one state per cycle; its remaining relation is the sum of its incident sink generators. The remaining cycle generators are free, proving the stated groups. □

For the nodal interval, G_α has two loop vertices, each with an exit to the common vacuum sink. Chirvasitu–Hajac–Tobolski's actual pullback theorem [65, Theorem3.4] gives

$$\{T \in \mathcal{T}_A : \pi(T) \in \mathbb{C}^2\} \cong C^*(E_1), \qquad \mathbb{C}^2 \hookrightarrow C(\mathbb{T})^2 \text{ by constants},$$

where E_1 has the same three vertices and countably many arrows from each loop vertex to the sink, with no loops. The arrow numbered $n \geq 0$ maps to loop^n followed by exit. This algebra is AF, with $K_0 = \mathbb{Z}^3$, $K_1 = 0$. The extended edition verifies every hypothesis. The theorem does not turn arbitrary cycle-vertex identifications into pullbacks: its first graph is acyclic and its path functor must cover all pointed paths. Its sink-extension Theorem4.3 has further injection and path-coverage requirements.

5.2 A separate diagonal-fiber C* completion retaining the pairing

The formal order's residue gluing cannot be copied by taking diagonals of arbitrary full matrices: $\operatorname{diag}(E_{12}E_{21}) \neq \operatorname{diag}(E_{12})\operatorname{diag}(E_{21})$. For $\ell > 1$ there is no nonzero *-homomorphism from $M_\ell(C(\mathbb{T}))$ to a commutative C*-algebra, as restriction to its constant matrix units shows. Restrict the distinguished fiber first. Define

$$D_\ell = \{f \in C(\mathbb{T}, M_\ell(\mathbb{C})) : f(1) \text{ is diagonal}\}, \qquad \mathcal{D}_H = \Big(\prod_c D_{|c|}\Big) \underset{\mathbb{C}^H}{\times} \mathbb{C}^{Q_0},$$

where the first map evaluates diagonal entries at one and the second repeats λ_e at the two darts paired by α.

Theorem 5.3 (A faithful polynomial pullback model)**.** *There is a split exact sequence*

$$0 \to \bigoplus_c C_0(\mathbb{T} \setminus \{1\}, M_{|c|}) \to \mathcal{D}_H \to \mathbb{C}^{Q_0} \to 0,$$

so $K_0(\mathcal{D}_H) = \mathbb{Z}^{|Q_0|}$ and $K_1(\mathcal{D}_H) = \mathbb{Z}^{\#\operatorname{cyc}\sigma}$. The polynomial dessin algebra embeds faithfully and generates this C-algebra with its adjoints. This embedding does not extend to a unital algebra map from all of Λ.*

Proof. The section consists of the constant paired diagonal projections. The kernel is the displayed vanishing fiber, a sum of matrix suspensions, giving the groups by split exactness and suspension. For an explicit embedding set $g(z) = |1 - z|^2/4$, choose one cut arrow per permitted cycle, and send a_h to $g(z)z^{\epsilon_h}E_{h,\tau h}$ in its branch, where $\epsilon_h = 1$ at the cut and zero otherwise. Send e_e to the constant paired diagonal projection. Every forbidden product is zero. A full cycle at h maps to $w_c(z)E_{hh}$, $w_c = g^{|c|}z$. This scalar function is injective on the circle: away from its unique zero at one its argument recovers z. Continuous functional calculus therefore supplies every $f(z)E_{hh}$ with $f(1) = 0$. Multiplying by path images and approximating division away from one supplies every vanishing off-diagonal entry. Together with the constant paired diagonals these generate $\mathcal{D}_H$. Path normal forms in a fixed entry are a nonzero prefactor times a polynomial in w_c; its infinite range proves independence. Evaluation at one first separates the zero-length coefficients. Finally for $z_0 \neq 1$, $1 - w_c(z_0)^{-1}c_h$ is a unit in the formal order, but its proposed image is singular at z_0. A unital map preserves units, proving nonextension. □

This is the C* algebra generated by a faithful polynomial representation and its adjoints, distinct from the completed order and from $\mathcal{T}_A$. The norm closure of the polynomial image alone need not contain those adjoints. Their K groups already differ. The pairing belongs to the labelled presentation; an unmarked isomorphism class need not recover the ribbon ordering. No arithmetic descent, Frobenius action or completed-L positive determinant is inferred from this additional valid gluing model.

Proposition 5.4 (The degree spectral triple)**.** *Define $N\delta_p = |p|\delta_p$ on $\{\xi : \sum_p |p|^2|\xi_p|^2 < \infty\}$. With the polynomial creation/annihilation *-algebra this is an ungraded spectral triple, p-summable exactly for $p > 1$. With N_+^{-s} zero on the kernel,*

$$\operatorname{Tr} e^{-tN} = E + \frac{2E}{e^t - 1}, \quad \operatorname{Tr} N_+^{-s} = 2E\zeta(s), \quad \operatorname{Tr}\big(\chi(N)N_+^{-s}\big) = 2EL(s,\chi)$$

for $t > 0$ and $\operatorname{Re} s > 1$ in the last two formulas, where χ is an externally chosen Dirichlet character.

Proof. There are E vacua and $2E$ paths of every positive length: each starting arrow has a unique continuation. The real diagonal operator is self-adjoint on the stated maximal domain; finite supports are a core and its resolvent is compact. Prefixes preserve the domain and $[N, L_a] = L_a$, $[N, L_a^*] = -L_a^*$, so polynomial commutators are bounded. The trace formulas and summability reduce to the corresponding scalar series. The character multiplier is bounded; membership in $\mathcal{T}_A$ is not asserted. □

Length parity anticommutes with the creators but commutes with N, so it is not an even-spectral-triple grading. The ordinary odd Fredholm class of this positive N is trivial: $N(1 + N^2)^{-1/2} - I$ is compact. The spectrum depends only on E; zeros of the scalar zeta trace are not eigenvalues of N.

Chapter 6

Arithmetic attachment, descent and the remaining comparison

The compact ideal gives a precise home for an arithmetic determinant, with no claim that the arithmetic is derived from the dessin. Fix one terminal sector. Its two paths of each length $n \geq 1$ give orthonormal vectors $\xi_n = (\delta_{p_{n,1}} + \delta_{p_{n,2}})/\sqrt{2}$ and an isometry $U : \ell^2(\mathbb{N}_{\geq 1}) \to \mathcal{H}_v$.

For a nonprincipal character χ put $A_n = \sum_{j \leq n} \chi(j)$, $B_0 = \sup |A_n| \leq q_\chi$, and, on $2^k \leq n < 2^{k+1}$, set

$$w_n = \sqrt{n}(k+1), \quad u_n = A_n/w_n, \quad v_n(s) = w_n(n^{-s} - (n+1)^{-s}), \quad \varphi_s(z) = \sum_n v_n(s) z_n - z_1.$$

Here $u_1 = 1$. For $\operatorname{Re} s > 0$, $u, v(s) \in \ell^2$ and the pairing is complex linear, with no conjugation of s. Define

$$\widetilde{K_s} = U(u \otimes \varphi_s)U^*, \qquad \det_F(I + \widetilde{K_s}) = L(s, \chi).$$

Indeed dyadic blocks give $\|u\|^2 \leq 2B_0^2$, and on a compact set $\operatorname{Re} s \geq \sigma_0 > 0$, $|s| \leq M$, the squared v tail is bounded by $M^2 \sum_{k \geq J} (k+1)^2 2^{-2\sigma_0 k}$. Thus the rank-one family is trace-norm holomorphic. Its determinant is $1 + \varphi_s(u)$, equal by Abel summation to L. Replacing v by $P_N v$ and adding $A_N(N+1)^{-s} e_1^*$ to the functional gives exactly every raw partial sum S_N, with compact-uniform trace-norm convergence; $\|UXU^*\|_1 = \|X\|_1$. For primitive χ, completion replaces the functional by $g(s) \sum v_n(s) z_n - z_1$, with the standard nonzero gamma factor g, rather than multiplying the whole old functional. The extended edition supplies the quantitative estimates.

This inserts a scalar encoding into one actual dessin compact ideal; it is not known to be A-linear, a derived chain map or Frobenius compatible. A common fixed positive metric cannot make this same fixed-range family symmetric for all critical heights: symmetry forces its functionals to be proportional, whereas the coefficient ratio at heights $0, t \neq 0$ is asymptotic to $(\sigma + it)n^{-it}/\sigma$, with two opposite subsequential limits. Different positive realizations remain possible.

Proposition 6.1 (Complex dessin data do not determine arithmetic descent)**.** *The rational Belyi pairs $E_d : y^2 = x^3 - d$, $\beta_d = x^3/d$, for $d = 1, 2$, are isomorphic over $\mathbb{C}$ but have distinct Frobenius polynomials at the good prime seven: $1 - 4T + 7T^2$ and $1 - T + 7T^2$.*

Proof. Choose $a^3 = d$, $b^2 = d$. The projective isomorphism $(x, y) \mapsto (ax, by)$ from E_1 to E_d preserves the Belyi map. The degree-six clean passports are the elliptic theta passport already computed. Directly counting y for $x = 0, \ldots, 6$ gives respectively $(0, 1, 1, 0, 1, 0, 0)$ and $(0, 0, 0, 2, 0, 2, 2)$. Including infinity yields four and seven points, hence $a_7 = 8 - \#E_d(\mathbb{F}_7) = 4, 1$. Both discriminants $-432d^2$ are nonzero modulo seven. □

Any invariant of the complex dessin or its undecorated order/Fock algebra alone is the same for these two inputs, so cannot recover both factors without descent data. This does not obstruct a theory including a specified rational model and compatible Galois action. Elliptic H^1 has weight one and unshifted center $s = 1$, unlike Dirichlet center $1/2$; the example is not a counterexample to GRH.

6.1 A constructive arithmetic benchmark on the theta dessin

The preceding twists can be distinguished constructively once their finite-field descent data are retained. This closes the elementary local-factor computation, not the global positive-determinant problem. Let $\kappa = \mathbb{F}_q$, $\operatorname{char}\kappa \neq 2, 3$ and $d \in \kappa^*$. Put $E_d : y^2 = x^3 - d$ and $a_q(d) = q + 1 - \#E_d(\kappa)$. For $q \equiv 1 \pmod 3$, fix a nontrivial cubic character ρ and the quadratic character η, both extended by zero at zero. Define

$$J = \sum_{u\in\kappa} \rho(u)\eta(1-u), \qquad j = -\eta(-1)J, \qquad \lambda_d = \eta(d)\rho(d), \quad \pi_d = \lambda_d j.$$

Theorem 6.2 (Uniform local factor)**.** *For every such prime power and nonzero d,*

$$a_q(d) = \pi_d + \overline{\pi_d}, \quad |\pi_d|^2 = q, \quad P_d(T) := 1 - a_q(d)T + qT^2 = (1 - \pi_d T)(1 - \overline{\pi_d}T).$$

If $q \equiv 2 \pmod 3$, instead $a_q(d) = 0$. The split-case values lie in $\mathbb{Z}[\omega]$, $\omega^2 + \omega + 1 = 0$.

Proof. Cube fibers have size $1 + \rho(t) + \overline{\rho(t)}$, including at zero; the first term is the literal constant one, not the zero-extended trivial character. Counting square fibers gives $a_q(d) = -\sum_x \eta(x^3 - d)$. Grouping by $t = x^3$ removes the constant term since $\sum_t \eta(t - d) = 0$. Substitution $t = du$ gives $\sum_t \rho(t)\eta(t - d) = \rho(d)\eta(-d)J$; the other term is its conjugate. This proves the signed trace formula.

For a nontrivial additive character ψ, write $G(\chi) = \sum_x \chi(x)\psi(x)$. Character orthogonality gives $|G(\chi)|^2 = q$ for nontrivial χ: substitute $x = uy$ in its squared norm; the inner sum over $y \neq 0$ is $q - 1$ at $u = 1$ and -1 otherwise. Grouping $G(\rho)G(\eta)$ by $x + y$ gives $JG(\rho\eta)$, the zero-sum fiber vanishing since $\rho\eta$ is nontrivial. The three Gauss norms yield $|J|^2 = q$. Neither purity nor RH was assumed. The remaining signs and character values have modulus one. For $q \equiv 2 \pmod 3$ cubing is bijective, so the quadratic-character sum is zero. Finally all split-case summands are signs times cube roots of unity. These are the classical Gauss–Jacobi methods [68, pp.499–502], with our normalization proved explicitly. □

At $q = 7$, choose $\rho(3) = \omega$. Then $j = 3 + 2\omega$, $\lambda_2 = \omega^2$, and $\pi_2 = -1 - 3\omega$, giving the two factors already counted. At $q = 13$, $\rho(2) = \omega$ gives $J = 1 + 4\omega$ but $j = -1 - 4\omega$; this detects the essential sign hidden by the first example.

6.2 A genuine cellular carrier, and what still fails for the order

Assume $q \equiv 1 \pmod 3$ in this cellular construction. The elliptic theta dessin has cellular chains $C_2 = \mathbb{Q}$, $C_1 = \mathbb{Q}^3$, $C_0 = \mathbb{Q}^2$. Orient the three edges from v_+ to v_-; then $\partial_1(e_i) = v_- - v_+$ and $\partial_2 = 0$. Thus H_1 is the sum-zero plane, with basis $e_0 - e_1, e_1 - e_2$. Rotation σ acts there by

$$C = \begin{pmatrix} 0 & -1 \\ 1 & -1 \end{pmatrix}, \qquad G = \begin{pmatrix} 2 & -1 \\ -1 & 2 \end{pmatrix}.$$

G is the actual Euclidean edge Gram matrix in this basis. The deck involution α acts as $-I$. Consequently the six-element deck group $\langle \sigma, \alpha \rangle$ acts on this homology and permutes the original quiver and its entire projective resolution.

Write $\pi_d = A_d + B_d\omega$. Its multiplication action on this *actual cellular* plane is

$$M_d = A_d I + B_d C = \begin{pmatrix} A_d & -B_d \\ B_d & A_d - B_d \end{pmatrix}, \qquad \det(I - TM_d) = P_d(T), \qquad M_d^t G M_d = qG.$$

This constructs the local matrix from finite-field character sums and the dessin's deck action, rather than fitting its trace. Positivity follows from $2x^2 - 2xy + 2y^2 = (x-y)^2 + x^2 + y^2$. The normalized matrix is G-unitary, not asserted self-adjoint. The extended edition supplies a rational cellular chain lift, its exact determinant cancellation, and Kummer descent compatibility on homology. Its odd-over-even determinant ratio is the classical finite-field curve zeta function

$$Z(E_d/\kappa, T) = \frac{P_d(T)}{(1-T)(1-qT)}.$$

Identification with the geometric ℓ-adic Frobenius polynomial uses the standard elliptic trace and determinant theorem; a canonical integral comparison or an order-derived Frobenius functor has not been constructed here.

There is also a precise obstruction to extending this same real two-dimensional carrier to $q \equiv 2 \pmod 3$. Frobenius then conjugates the cubic deck action. The relation $MC = C^2 M$ forces $M = \left(\begin{smallmatrix} a & b \\ a+b & -a \end{smallmatrix}\right)$, whose determinant is $-a^2 - ab - b^2 \le 0$, contradicting the required positive determinant q. This excludes that real carrier with its fixed deck action, not complex or enlarged Hilbert–Pólya constructions. The extended edition gives a compatible complex local replacement.

For the following insertion test use the characteristic-zero theta order $\Lambda_{\mathbb{Q}}$ over $R_{\mathbb{Q}} = \mathbb{Q}[[u]]$, with residue simple $S_e = \mathbb{Q}$. The limitation is concrete. The deck action has finite order and cannot alone supply the norm-$\sqrt{q}$ Jacobi coefficient. Inserting M_d into $S_e \otimes \mathbb{Q}^2$ gives an actual residue-module endomorphism with determinant P_d, but that simple is not perfect over the order and choosing it breaks full deck symmetry. Inserting the same matrix $\Lambda_{\mathbb{Q}}$-linearly into the regular projective instead gives the $R_{\mathbb{Q}}$-determinant P_d^{18}, or P_d^6 on $e\Lambda_{\mathbb{Q}} \otimes \mathbb{Q}^2$. Taking a root is not an established determinant-line comparison; a deck-semilinear lift is a different map. These surviving multiplicity/equivariance obligations prevent promoting the local computation to the required arithmetic order functor.

For the global comparison fix a nonzero rational d and use its good reductions. The finite-field curve zeta has center $1/2$ after $T = q^{-s}$. The *global elliptic* L-function over $\mathbb{Q}$ has unshifted center one, whereas Dirichlet L-functions have center $1/2$. Local weight-one purity does not determine either global zero set. The elementary local bound proves absolute convergence of the elliptic Euler product on $\operatorname{Re} s > 3/2$; the new local identities give no determinant-preserving continuation with the desired global positive spectrum.

6.3 Arbitrary projective actions: an obstruction and a genuine determinant repair

The preceding tensor test has a stronger replacement. Keep the characteristic-zero completed elliptic theta order Λ over $R = \mathbb{Q}[[u]]$, and write $K = \mathbb{Q}((u))$. All determinants in the next theorem are *ordinary underlying-R determinants*; they are not reduced norms.

Theorem 6.3 (Projective determinant multiplicity)**.** *Let P be any finitely generated projective right Λ-module and $f \in \mathrm{End}_\Lambda(P)$. Then $P \simeq \bigoplus_{i=0}^{2} (e_i\Lambda)^{n_i}$. Let $A_i \in M_{n_i}(\mathbb{Q})$ be the induced maps on its three top multiplicity spaces. Then*

$$\det_R(1 - Tf)\Big|_{u=0} = \left(\prod_{i=0}^{2} \det(1 - TA_i)\right)^6.$$

For any bounded complex of these projectives and any Λ-linear chain endomorphism, its alternating determinant ratio specializes to a sixth power in $\mathbb{Q}(T)$. Generically, the ordinary ratio is a cube in $K(T)$.

Proof. The arrow ideal is the Jacobson radical: its elements admit geometric-series inverses after subtracting from one, and the quotient is $\mathbb{Q}^3$. A projective's top lifts to a map from the displayed sum; Nakayama gives surjectivity, projectivity splits it, and another application of Nakayama kills the kernel.

Put $B = \Lambda/u\Lambda$. Each e_iB has path layers of dimensions $1, 2, 2, 1$: the identity, two length-one paths, two length-two paths, and one length-three path. The two full cycles based at i add to ue_i, so become negatives; all length-four paths vanish. A positive-length entry of a projective homomorphism strictly raises path length. Consequently the associated graded map has six copies of $A = \mathrm{diag}(A_0, A_1, A_2)$, giving the formula by triangular block determinants. Base change of determinants justifies specialization. Multiplicativity along filtrations is ordinary linear algebra, consistent with [69, Tags0FJ9,0GSX]; no purity assumption enters.

Each term's determinant has constant term one in T, so every denominator has nonzero specialization. Taking the finite alternating product therefore preserves the sixth power, even after cancellation. Finally $\Lambda_K = M_3(K) \times M_3(K)$; a projective's branch modules are sums of three-dimensional simples. A commuting endomorphism acts on each multiplicity space, repeating every factor three times. This proves the generic assertion. □

Thus no such ordinary perfect-complex ratio can equal the constant-in-u quadratic $P_d(T)$, or the full local curve zeta $P_d(T)/((1-T)(1-qT))$. For P_d, the valuation at infinity is -2, not a multiple of six. For the zeta ratio, $P_d(1) = q + 1 - a_q > 0$ by the local bound $|a_q| \le 2\sqrt{q} < q + 1$, so its pole at $T = 1$ is simple. This also covers square prime powers where the numerator may have a double root. It is an obstruction for the specified ordinary determinant and linear category, not all spectral or determinant theories.

There is nevertheless a *genuine integral functor* removing this multiplicity. Choose a permitted branch and its cycle cut. Its projection $\Lambda \to B_+ \subset M_3(R)$ makes the column lattice $W_+ = R^3$ a left Λ-module. Define

$$\Phi_+ : \mathrm{Perf}(\Lambda) \longrightarrow \mathrm{Perf}(R), \qquad P^\bullet \longmapsto P^\bullet \otimes_\Lambda W_+.$$

Bounded projectives compute this derived tensor functor; flatness of W_+ on all modules is unnecessary. Since $e_iW_+ \simeq R$, its determinant on $e_i\Lambda \otimes_{\mathbb{Q}} V$ with external action $1 \otimes M_d$ is exactly P_d, whereas on the regular module it is P_d^3. This is a defined change of determinant functor, not an arbitrary polynomial root. It does not construct the arithmetic action independently: the already computed M_d is still the input.

The price of this repair is precise. Branch swap α exchanges W_+ and W_-; pairing them yields P_d^2 on that basic-projective input. The chosen branch determinant cannot intertwine α on all endomorphisms: the corner unit $e_i + c_i^+$ acting on $e_i\Lambda$, with c_i^+ a plus full cycle, acts on its determinant by $1 + u$, whereas its α-image acts by one. Also no rank-one-per-vertex column lattice supported on one branch admits an R-linear lift of the rotation σ fixing u. Its three nonzero arrow

scalars have nonnegative u-valuations summing to one; an invertible cyclic lift would make all three equal. This contradiction concerns that lattice, not every determinant-level rotation comparison. Ramification $u = v^3$ *and a new lattice* with all three arrow scalars v admits such a rotation, changing the stipulated base/lattice data. A descent-compatible arithmetic determinant retaining the required symmetry and a global positive realization remains open.

6.4 Tame descent: a surviving order and a sharper obstruction

We now test the proposed ramified repair rather than assume that its rotation descends. Extend coefficients explicitly to $k = \mathbb{Q}(\omega)$, $\omega^2 + \omega + 1 = 0$. Put $R = k[[u]]$, $S = k[[v]]$, $u = v^3$, and let $g(v) = \omega v$, fixing k. On the balanced plus column $W = S^3$, let $C\epsilon_i = \epsilon_{i+1}$ and

$$\Theta = Dg, \qquad D = \operatorname{diag}(1, \omega^2, \omega).$$

The plus arrows act by $vE_{i,i-1}$, and minus arrows by zero. Direct calculation gives $\Theta^3 = 1$, compatibility of Θ with this Λ_R-action, and

$$\Theta C = \omega^2 C\Theta, \qquad L = W^{\Theta} = R\epsilon_0 \oplus Rv\epsilon_1 \oplus Rv^2\epsilon_2.$$

Thus invariants restore precisely the original cut lattice. The extension $R \to S$ is finite *free* of rank three and faithfully flat; the group quotient is not a torsor at $v = 0$. The comparison $S \otimes_R L \to W$ has matrix $\operatorname{diag}(1, v, v^2)$, determinant u and cokernel $S/(v) \oplus S/(v^2)$. Hence this semilinear group action alone is not effective descent of the column to R. This is the explicit local root-stack situation, consistent with the trivial-stabilizer criterion for vector-bundle descent [70, Theorem10.3]; no stack theorem is needed for the displayed calculation.

Two positive constructions survive. The determinant line $\bigwedge_S^3 W$ has trivial D and C actions and does descend, but the natural map

$$\det_R(L) \longrightarrow (\det_S W)^{\Theta}$$

has image uR inside R, not an isomorphism. Determinant and coarse invariants cannot be interchanged while dropping this nonunit comparison. Also the *carrier* endomorphism algebra $\operatorname{End}_S(W)^{\Theta}$, conjugated to the invariant basis, is exactly the hereditary order with entries R on and below the diagonal and uR above. Its rotation descends as $\operatorname{Ad}(A)$, where

$$A = \begin{pmatrix} 0 & 0 & u \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix}, \qquad A^3 = uI.$$

The algebra automorphism has order three although A is not an integral unit. This is not $\operatorname{End}_\Lambda(W)$ or a descended arithmetic cohomology group.

The natural arithmetic attempt fails: $aI + bC$ does not commute with Θ when $b \neq 0$, as in both prime-seven Jacobi examples. The original cellular sum-zero plane is not Θ-stable either. Replacing C by vC does give the descended operator A on L, but

$$\det(1 - T(aI + bA)) = (1 - aT)^3 - ub^3T^3,$$

whose special fiber is $(1 - aT)^3$, not the required local quadratic. These carrier maps are not automatically Λ-linear arithmetic endomorphisms.

There is a stronger category-level conclusion.

Theorem 6.4 (Full-deck coefficient kernels)**.** *Let W_0 be a finite R-free left Λ_R-module on which the central u acts by the given scalar. Suppose it has invertible R-linear intertwiners lifting the order automorphisms σ and α. Then, for some $m \geq 0$, every $e_i W_0$ has rank $6m$. For any finite projective P and Λ_R-linear $f : P \to P$, let A_i be its top blocks. The genuine coefficient functor $\Phi_{W_0} = - \otimes^{\mathbf{L}}_{\Lambda} W_0$ satisfies*

$$\det_R(1 - T\Phi_{W_0}(f))\Big|_{u=0} = \left(\prod_i \det_k(1 - TA_i)\right)^{6m}.$$

Bounded perfect ratios retain this exponent. Thus this whole class of full-deck, fixed-center finite free coefficient kernels cannot supply either the single local quadratic or the full local curve zeta.

Proof. In the generic fiber intersect W_0 with its two central branch summands to obtain finite free integral lattices $W_\pm$; an integral direct-sum decomposition is not assumed. Rotation preserves them and branch swap exchanges them. On a branch let the common vertex rank be r. The three arrow matrices have cyclic product uI_r. Their determinant valuations are nonnegative, sum to r, and rotation intertwiners make them equal. Therefore $3 \mid r$. Branch swap makes the two ranks equal, so the full vertex rank is $6m$.

For the determinant formula, filter W_0/uW_0 by the arrow radical, whose fourth power vanishes. A positive-length entry of f raises the filtration; a top block A_i repeats across its i-th graded pieces a total of $\dim_k e_i(W_0/uW_0) = 6m$ times. Triangular determinants give the assertion. Finite alternating products preserve it. Quadratic degree and the uncancelled simple pole of the full local zeta exclude the targets as before. □

This theorem is sharp: $W_0 = \Lambda_R$ has vertex rank six. It requires neither left projectivity nor left flatness of a general kernel. It excludes this coefficient category, not arbitrary semilinear correspondences, twisted determinant lines or infinite spectral constructions. The extended edition identifies the full character-orbit repair with the hereditary normalization and explains why Schur commutator cancellation is insufficient. All determinant multiplicities here are relative to the stated k-valued base; restricting a rational quadratic factor to the former $\mathbb{Q}$-base squares it. No arithmetic descent or global positive determinant has been recovered by this ramification.

The recent NCG literature supplies concrete comparison targets, with strong hypotheses. Connes–Consani construct finite abelian covers of the adelic Riemann sector with prime-orbit length $\log p$ and arithmetic Frobenius monodromy [53, Theorem3.12]. The prolate/local-factor and moment constructions [54, 55] give particular finite-prime analytic models, not a uniform positive comparison for all completed L-functions. The zeta-spectral-triple determinant theorem [56, Theorem5.10] assumes a specified restricted Weil form, simple even bottom eigenfunction and boundary normalization. Its infinite comparison is still an additional step. The distributional real-zero theorem [57, Theorem6.1] concerns a particular interval quadratic form, essential self-adjointness on its trigonometric core and an even simple ground state; a Fock vacuum alone verifies none of those arithmetic identifications. Deninger's periodic-packet construction and real-cohomology obstruction [60, 62] further specify, rather than remove, these issues. Full theorem-level reading coverage and exact versions are recorded alongside this draft.

A complete approach must construct arithmetic descent and prime actions on actual order/derived or analytic objects, prove determinant-line and local-factor compatibility, control regularization and conductor/gamma normalization, and establish a positive realization for that *same* function. Finite-field cohomological trace formulas and Deligne purity supply the classical model [20, 21]; alternating duality alone is not positivity or purity. In a number-field construction these are obligations, not imported conclusions.

For a sharply testable final analytic target, put $F_\chi(z) = \Lambda(1/2 + iz, \chi)$. Suppose entire functions $h_m(z)\det(zI - H_m)$, with H_m finite self-adjoint and h_m entire nowhere zero, converge locally uniformly on $\mathbb{C}$ to this nonzero completed function. Hurwitz on every disk disjoint from $\mathbb{R}$ then proves all its zeros real, hence the Dirichlet critical-line assertion. Establishing this convergence with independently constructed H_m is exactly the missing load-bearing claim; a rank-one encoding or a reformulation does not establish it. Hecke, Dedekind and Artin extensions must state their analytic scope anew.

6.5 Whole components, derived endpoints and determinant relations

The full node lattice category has precisely the three indecomposables X, Λ, Y and arrows $X \to \Lambda \to Y$, $Y \to \Lambda \to X$. Its stable category has only X, Y, no irreducible arrows, and suspension exchanging them; its AR triangles have zero middle term. For the polynomial node at a smooth support $x = \lambda + t$, $y = 0$, $\lambda \neq 0$, the modules $T_r = k[t]/t^r$ form an ambient homogeneous tube:

$$0 \longrightarrow T_r \xrightarrow{(\pi,t)} T_{r-1} \oplus T_{r+1} \xrightarrow{(t,-\pi)} T_r \longrightarrow 0, \quad T_0 = 0, \quad \tau T_r = T_r.$$

The shift changes the derived placement, not the quasi-length. These nonzero-central-support objects do not extend to continuous finite-length modules of the arrow-adic completion. At its singular point even the candidate pure-branch tube fails ambient almost-splitness: $T_2^y \to k$ cannot factor through $T_2^x \to k$, since every intervening map lands in the x-socle, killed by the last map.

The nine HLS strings above sit in an entire intrinsic quotient component $T_{p,q}$, $(p, q) \in \mathbb{Z}^2$, with arrows increasing one coordinate and $\tau T_{p,q} = T_{p-1,q-1}$ [46]. Its central object $E = M(ab)$ has infinite projective dimension: restriction to the first corner gives the dual-number simple with its infinite loop resolution, while restricted projectives are free or zero. Its module AR sequence therefore yields a distinguished derived triangle but not a derived AR triangle ending at E. Ringel's four-quarter wind-wheel tile [43], with infinite-word diagonals and mixed center, is a different category's quilt; it is not silently identified with that HLS component or the Belyi surface.

There is no general equation "AR mesh equals polynomial semi-invariant syzygy." The intrinsic $k(1 \to 2)$ quotient of the infinite Chebyshev algebra has AR row $0 \to S_2 \to E \to S_1 \to 0$ and a free presentation determinant z on dimension $(1, 1)$, so its semi-invariant ring $k[z]$ has no generator relation. A genuine framed elliptic-square stratum has $Z = [x_1, x_2, y_1, y_2] \in \mathrm{Mat}_{2\times 4}$ with six minors and relation

$$p_{12}p_{34} - p_{13}p_{24} + p_{14}p_{23} = 0.$$

These minors have the same determinant character for the target GL_2 with framed source; they are not all semi-invariants for the full unframed vertex group. In contrast, the nodal lattice AR row gives a genuine relative determinant identity over R: multiplication by $a + bx$ on the basis $(y, 1)$ has matrix $\left(\begin{smallmatrix} a & -b \\ 0 & a+ub \end{smallmatrix}\right)$ and characteristic factor $(1 - Ta)(1 - T(a + ub))$. The endpoints have infinite order-projective dimension but are R-perfect.

Finally, the three-vertex theta algebra has no unital surjection onto a six-vertex radical-square-zero circular algebra: its semisimple image has at most three primitive idempotent images. Duplicating face vertex spaces instead gives an exact faithful unfolding which is not full; it does not inherit quotient or AR-transfer claims. Zwara Hom-control means that the Hom-dimension defect is bilinear on K_0, not merely a covering direct-sum formula [72, 73]. The book supplies explicit fixed-dimension counterexamples, corrected circular-complex incidence resolutions, complete drawn component windows and all category hypotheses. These additions do not supply the arithmetic determinant comparison.

Chapter 7

A genuine theta covering and its categorical limits

Use column matrices, hence covariant quiver representations (left modules); opposite categories translate the right-module convention in [74]. Let A be the infinite elliptic theta dessin algebra: vertices $i \in \mathbb{Z}/3$, arrows $a_i, b_i : i \to i+1$, and all mixed length-two paths zero. For the arrow ideal J and $L \geq 2$ put $B_L = A/J^L$. This is a finite string quotient of our actual infinite gentle algebra; for $L > 2$ its length-L relations prevent treating it as quadratic gentle.

Proposition 7.1 (A locally bounded integer cover)**.** *The category $\widetilde{B}_L$ with vertices $r \in \mathbb{Z}$, two arrows $a_r, b_r : r \to r+1$, and the same mixed and length-L zero relations, has orbit B_L under $r \mapsto r+3$. For finite-support representations, push-down F is exact and*

$$\operatorname{Hom}_{B_L}(FM, FN) \simeq \bigoplus_{g \in \mathbb{Z}} \operatorname{Hom}_{\widetilde{B}_L}(M, T_g N), \qquad (T_g N)_r = N_{r-3g}.$$

Only finitely many summands are nonzero. In general F is neither full, dense, nor Hom-controlled in Zwara's sense.

Proof. A nonzero path is constant-branch with length less than L. Thus the incoming and outgoing total Hom dimensions at a cover vertex are $2L-1$, and $\dim B_L = 3(2L-1)$. Paths and relations lift uniquely; the deck action is free and has the stated orbit. Push-down sums spaces of equal residue. A morphism decomposes uniquely into blocks with target lift index minus source index equal to $-3g$. The arrow commutation equations preserve that difference, proving the formula; finite supports bound the possible differences. Direct sums of vector spaces preserve exact sequences. The following explicit witnesses prove the three failures. □

For $L = 5$, let P be the cover interval on $0, 1, 2, 3, 4$ with all its a-maps one and b-maps zero, and let S be the sum of the five simples at the same vertices. They have the same K_0 class, but

	$\dim \operatorname{End}_{\widetilde{B}_5}$	$\dim \operatorname{End}_{B_5}(F(-))$	defect
S	5	9	4
P	1	2	1

Indeed the pushed-down simple multiplicities are $(2, 2, 1)$, whereas $\operatorname{End}_{B_5}(FP) = kI \oplus kX^3$ for its single nilpotent five-step chain X. The extra endomorphism X^3 proves non-fullness; the unequal defects at the same class exclude a bilinear K_0 defect. To use Zwara's finite-algebra definition literally, restrict the source to the finite convex interval $[0, 4]$, then extend by zero before push-down. This exact functor has the same endomorphism spaces and K_0 comparison.

Proposition 7.2 (An indecomposable module outside this cover image)**.** *Over a characteristic-zero field, B_3 has a nine-dimensional indecomposable representation which is not the push-down of any finite-support $\widetilde{B}_3$ representation.*

Proof. Label its basis by $0, \dots, 8$ with vertex residues $(0, 1, 0, 1, 2, 1, 2, 0, 2)$, and use matrix units E_{ij} acting on columns. Set

$$X = E_{10} + E_{32} + E_{43} + E_{65} + E_{76} + E_{08},$$
$$Y = E_{12} + E_{45} + E_{78}.$$

Their restrictions between adjacent vertex spaces give the a_i, b_i maps. They satisfy $XY = YX = X^3 = Y^3 = 0$, hence every length-three path is zero. The endomorphism ring preserving vertex spaces is $kI \oplus \mathfrak{r}$, where

$$\mathfrak{r} = \langle E_{02} + E_{13}, E_{15}, E_{35} + E_{46}, E_{48}, E_{68} + E_{70}, E_{72} \rangle, \qquad \mathfrak{r}^3 = 0.$$

Solving the arrow commutation equations gives exactly these seven basis elements. The ring is local: an element $cI + n$ is invertible exactly when $c \neq 0$, by the finite geometric series. It has no nontrivial idempotents, so the representation is indecomposable.

A finite-support lift would define an Euler operator D acting by r on the summand at lift vertex r. Then $[D, X] = X$ and $[D, Y] = Y$. Every displayed nonzero arrow is the unique nonzero entry in its row and column in its own matrix. Its commutator entry is therefore $D_{ii} - D_{jj}$. Add the six X equations and subtract the three Y equations. Each diagonal term cancels, leaving $0 = 6 - 3 = 3$, a contradiction. This argument allows arbitrary D, not just a diagonal matrix in the displayed basis. □

Corollary 7.3 (Higher-genus and derived non-density)**.** *For the clean dessin with rotations $\sigma = (0 \cdots m-1)(m \cdots 2m-1)$ and $\alpha = \prod_{i=0}^{m-1}(i\, i+m)$, odd $m \geq 3$, and genus $(m-1)/2$, the analogous A_m/J^3 has a $3m$-dimensional representation outside its integer-cover image in characteristic zero. Its bounded-derived push-down is not dense either.*

Proof. Use basis $v_0, \dots, v_{3m-1}$, assign vertices $i, i+1, i$ to $v_{3i}, v_{3i+1}, v_{3i+2}$ modulo m, and set

$$Xv_{3i} = v_{3i+1}, \quad Xv_{3i+2} = v_{3i+3}, \quad Yv_{3i+2} = v_{3i+1},$$

with indices modulo $3m$ and all other values zero. Again mixed products and cubes vanish. An Euler operator would yield $2m - m = m = 0$ by the same telescoping calculation, impossible in characteristic zero. If a bounded finite-support complex pushed to this module in degree zero, exactness would make its zeroth cohomology a module lift, a contradiction. □

The source's derived result is a *precovering* in the stated bounded categories [74, Theorem 4.28]; density remains an additional condition. Neither this result nor the GLS equivariant classification [75] identifies our completed order with a GLS algebra. Graded surface/Fukaya models [76] likewise require their own dg category, stops and grading. These categorical tools do not construct the independent positive arithmetic determinant.

Chapter 8

Verification and research frontier

The two nodal/invariant Lean files construct the node as an actual subring of $k[[X]] \times k[[X]]$ with equal constants, prove its annihilators, regular common parameter, augmentation, alternating projective resolution and a bijective pair of power-series coordinates. A second file proves CK1 collapse under the stated permutation/coverage hypotheses, determinant base change, actual 2×2 symmetric-square functoriality and relation preservation, a Schur character identity and the fixed-parameter lattice failure. These thirty printed theorem types have only standard Lean foundations. They do not formalize general dessins, abstract Ext/AR/Schur–Weyl, C* exactness, Fredholm determinants, arithmetic descent or GRH. Earlier formal layers remain preserved with their literal scopes.

Exact-rational diagnostics independently check the displayed dessins, path identities, Schur and Young examples and the nine-object grid. These finite checks supplement the proofs, rather than prove infinite spectral or arithmetic assertions. The next arithmetic target is an exact descent/Frobenius-compatible construction on the labelled order/Fock data, with a determinant-preserving positive comparison. The local Schur kernel and the simple degree spectrum alone do not supply it. All source clones and earlier research are retained; the earlier general focused PDFs are archived at the preceding commit.

Ten further declarations in `DessinPullbackComponents.lean` check actual matrix-unit identities, nonmultiplicativity of raw diagonal extraction, impossibility of a unital two-state matrix residue into a nonzero commutative ring, multiplicativity on diagonal fibers, and genuine equal-image pullback subring/projection components. The analytic suite then contained187printed declarations at unchanged pins. The circle-function C* model, functional calculus, generation, graph uniqueness and K-theory remain prose/external-theorem arguments, not hidden assumptions of these component checks.

Sixteen additional declarations in `ThetaJacobi.lean` check genuine finite-field MulChar Jacobi conjugation, norm and rescaling with explicit nontriviality hypotheses; cubic–quadratic noncancellation; the cyclotomic multiplication determinant and positive Gram identity; and actual theta cellular boundary, homology-inclusion, restriction and rational chain-map matrices. The full analytic suite now contains203declarations. The point-count proof, deck/Kummer comparison, cohomological trace interpretation and global arithmetic comparison are not thereby fully formalized.

The final three `ThetaJacobi.lean` declarations verify the actual conjugation matrix identities and the real two-dimensional inert-prime determinant obstruction. They do not assume any unknown zero location or exclude complex/enlarged spectral realizations.

`ThetaProjectiveMultiplicity.lean` adds five declarations. Arbitrary upper-block matrices with six identical diagonal blocks have the sixth-power determinant and characteristic-factor

formulas. A literal power-series constant-coefficient map specializes matrix determinants, for each scalar parameter; the coefficientwise polynomial/order interface remains prose. A nonzero quadratic cannot be a rational-function sixth power, proved by polynomial degrees. Three genuine power series with product X cannot have the cyclically equivalent residue-vanishing conditions forced by invertible rank-one arrow intertwiners. The derivation of the path blocks from the full theta order, projective decomposition, derived branch functor and deck intertwiners is not fully formalized. The complete analytic suite now has208printed declarations; no hidden arithmetic positivity assumption occurs.

`ThetaTameDescent.lean` adds twelve printed component declarations (and one auxiliary polynomial lemma). They use genuine formal-power-series rescaling and its coefficient operator, proving the coefficient action, invariant-coefficient vanishing and the tame operator's cube. Actual three-state matrices satisfy the clock/rotation commutator, comparison determinant and closed-fiber equations, weighted-cycle cube and full characteristic polynomial, noncommutation of a nonconstant rotation polynomial, trivial whole-column determinant actions, and the same-operator comparison for vC. The classification of all fixed series by residues, full order action and hereditary invariant algebra, stack interpretation, valuation/rank theorem, general coefficient-functor filtration and GRH are not thereby fully formalized. No unknown positivity premise is assumed.

Part II

Proofs and calculations

Chapter 9

Expanded dessin, resolution and derived calculations

9.1 The cycle order in coordinates

Number a permitted cycle in travel order $0, \dots, \ell-1$. Map an arrow $i \to i+1$ to $E_{i,i+1}$ for $i < \ell-1$, and the final arrow to $u_c E_{\ell-1,0}$. A path from i to j then gives $u_c^r E_{ij}$, where r counts cut crossings. These are linearly independent path normal forms, and every positive-length permitted path has exactly one such expression on exactly one branch. The full-cycle sum in a branch is $u_c I$. At a paired edge its two diagonal occurrences come from a single original idempotent; no positive-length coefficients are identified. This proves the order description including self-pairings within one cycle. Although the suborder is generically a product of full matrix algebras, it is not their unglued product at the special fiber.

For the interval this becomes the literal fiber product

$$\Lambda = \{(f,g) \in k[[X]]^2 : f(0) = g(0)\}, \qquad x = (X,0),\ y = (0,X),\ u = (X,X).$$

Every element has a unique expression $a(u) + b(u)x$: take $a = g$ and $b = (f-g)/X$. Thus it is R-free on $1, x$ with $x^2 = ux$ and $y = u - x$. The annihilator of x consists of pairs $(0,g)$ with $g(0) = 0$, exactly $y\Lambda$; the other branch is symmetric. Multiplication by u is injective, coefficientwise. These are actual formal power-series statements, not conclusions from finite truncations.

9.2 The two-arm complex and its calculations

For an arrow a write $f(a)$ for its unique forbidden successor. The summand in $P_n(i)$ indexed by the initial arrow a has endpoint $t(f^{n-1}(a))$. Its map multiplies by that last arrow. Thus the resolution can be computed by iterating a finite permutation, while the projectives themselves remain infinite-dimensional. The Hom complex against a module N is the direct sum of the endpoint spaces with differentials given by the same arrow matrices. Tensor against a left module instead uses the same path entries acting on the left. No rank or exactness claim for an arbitrary N is inferred by counting endpoints alone.

For the nodal residue field the maps are particularly explicit:

$$\cdots \xrightarrow{\mathrm{diag}(x,y)} \Lambda^2 \xrightarrow{\mathrm{diag}(y,x)} \Lambda^2 \xrightarrow{(x \quad y)} \Lambda \longrightarrow k \longrightarrow 0.$$

If $xz_1 + yz_2 = 0$, its first branch says the first branch of z_1 is zero, and its second branch says the second branch of z_2 is zero. Hence $z_1 = yw_1$, $z_2 = xw_2$. The same annihilator calculation proves

the alternating kernels in all degrees. Projectivity is genuine: the terms are regular modules or their finite direct sums. It is not a finite free resolution, despite the period of its positive tail.

For the elliptic theta the quiver vertices are $v_0 = \{0, 3\}$, $v_1 = \{1, 4\}$, $v_2 = \{2, 5\}$. The arrows a_0, a_3 go $v_0 \to v_2$, a_2, a_5 go $v_2 \to v_1$, and a_1, a_4 go $v_1 \to v_0$. The two forbidden arms from v_0 are

$$a_0, a_5, a_1, a_3, a_2, a_4, \ldots, \qquad a_3, a_2, a_4, a_0, a_5, a_1, \ldots.$$

Their top pairs are $(v_2, v_2), (v_1, v_1), (v_0, v_0)$, repeated. The three steps exchange the arms. Their distinct top pairs exclude shorter unordered positive periods. The spherical theta instead has forbidden two-cycles and successive top pairs $(v_2, v_1), (v_0, v_0)$, giving period two. Thus even identical hereditary overorders need not give identical resolutions.

Quadratic Koszul duality here has a directly checkable local meaning: relations complement the two possible successors. The dual dessin replaces the graph rotation by ϕ^{-1}. In the Ext algebra against A/J, concatenating the forbidden-word basis gives the product when the join is forbidden; a permitted join gives zero. One can see this by lifting a length-r basis class to the resolution: discard the prescribed initial r letters and identify the remaining summand with its suffix resolution, taking zero on other prefixes. These prefix-discard comparison maps commute with the suffix differentials, with the standard shift signs when viewed as shifted complexes. The concatenated indexing word exists exactly when its join is forbidden; evaluation of the two basis classes on that word gives one. If the join is permitted there is no such indexing word and the product is zero. Composition therefore concatenates indexing words, with the chosen right-module/Yoneda opposite convention. This is a graded algebra statement; it neither makes the unbounded resolution perfect nor constructs an equivalence of all formal-completion derived categories.

For an arbitrary bounded projective complex $P^\bullet$, its determinant line over a commutative coefficient ring is $\bigotimes_n (\det P^n)^{(-1)^n}$, with determinant duals for negative exponents. A chain automorphism acts on this line; a quasi-isomorphism gives the corresponding determinant-functor isomorphism. For a general endomorphism F, the expression for $1 - zF$ uses the locus where it is invertible and gives a rational determinant-line action/trivialization, not an everywhere invertible line map at singular parameters. Additivity is compatible with exact triangles under the usual determinant-functor conventions [17, 18]. In our order context an R-linear determinant uses the finite free underlying R-modules; it must specify a separate endomorphism. Determinants of a differential presentation, alternating determinants of a chain action, regularized analytic determinants and the Cartan path matrix are four distinct constructions. Their equality requires a proof, not their shared name.

The geometric model literature [6, 8, 9] is useful for identifying which crossings and gradings would have to be represented by a functor. To apply it here one must specify the category of polynomial strings, completed lattices or finite-length quotients; construct its surface/curve objects; and prove that degree-r intersections compute the Hom group into shift r, with local systems and boundary exceptions included. We have supplied explicit walks and maps in the examples, not claimed this general completed-order theorem. On a finite word, a map can already be computed rigorously without such a theorem: assign a coefficient matrix between equal-vertex occurrences and solve $F_{t(a)} N(a) = N'(a) F_{s(a)}$ for every arrow. The diamond calculations below implement exactly this system.

Chapter 10

Semi-invariants, tableaux and explicit intertwiners

10.1 Derived information retained by a presentation

Let $K = \operatorname{im} f$. The injection $\operatorname{Hom}(K, N) \to \operatorname{Hom}(P_1, N)$ descends to the Ext injection used in Theorem 3.1; its kernel modulo $\operatorname{im}\operatorname{Hom}(P_0, N)$ is zero because the map $P_1 \to K$ is onto. The missing part of the cokernel detects maps on P_1 not descending through K. This explains exactly why an injective resolution map is needed for its full Euler interpretation.

In the node example the resolution of $M = \Lambda/(x)$ has maps $x, y, x, y, \ldots$. On $N = k[y]/(y^2)$ the Hom cochain maps are $0, y, 0, y, \ldots$. Thus $\operatorname{Hom}(M, N) = N$ and $\operatorname{Ext}^1(M, N) = \ker y$ have dimensions two and one; the cokernel of the first zero map is two-dimensional. In the elliptic-square example the first four Hom maps are $1, 0, 0, 0$, periodically. Hence degrees zero and one vanish, degrees two and three are k, and degree four vanishes. This also makes clear why a nonzero scalar semi-invariant does not give a derived orthogonality theorem.

With coordinate action $(gF)(N) = F(g^{-1}N)$ the Schofield weight is $a - b$, the negative of the point-evaluation weight. For a single arrow it is the same elementary distinction as the action on $\operatorname{Hom}(V, W)$ versus on its coordinate functions $V \otimes W^*$. The Cauchy decomposition for coordinate functions is

$$\operatorname{Sym}^m(V \otimes W^*) \cong \bigoplus_{\lambda \vdash m} S_\lambda(V) \otimes S_\lambda(W)^*.$$

A relation quotient can kill these components; its ambient multiplicities do not descend unchanged. Acyclic-quiver semi-invariant/LR results such as [35] retain their acyclicity assumptions. The explicit two-arrow slice below is acyclic, whereas the whole theta quiver is not.

10.2 One Young operator and one LR decomposition in coordinates

For $\lambda = (2, 1)$ number the first row $1, 2$ and second row 3. Then

$$Y = (I + P_{12})(I - P_{13}), \quad Y^2 = 3Y, \quad \dim \operatorname{im}(Y \mid (k^2)^{\otimes 3}) = 2.$$

The hook product is three; its Specht dimension is $3!/3 = 2$, also two here for a different reason. The normalized projector extracts one $S_{(2,1)}$ copy. The trace $\operatorname{tr}((Y/3)M^{\otimes 3})$ is therefore the Schur character, not twice that character. This fixes the common row/column normalization ambiguity.

Write $t = \operatorname{tr} M$, $\Delta = \det M$ for a two-dimensional endomorphism. Its complete characters obey $h_0 = 1$, $h_1 = t$, $h_m = th_{m-1} - \Delta h_{m-2}$, and

$$\chi_{(a,b)}(M) = \Delta^b h_{a-b}(M), \quad a \geq b \geq 0.$$

In particular $h_2 = t^2 - \Delta$, $h_3 = t^3 - 2t\Delta$, and $\chi_{(2,1)} = t\Delta$. Thus $th_2 = h_3 + \chi_{(2,1)}$ is the LR identity $V \otimes \operatorname{Sym}^2 V \cong \operatorname{Sym}^3 V \oplus S_{(2,1)}V$, with $S_{(2,1)}V \cong \det V \otimes V$.

An explicit natural splitting avoids relying just on characters. In characteristic zero multiplication $m : V \otimes \operatorname{Sym}^2 V \to \operatorname{Sym}^3 V$ has natural section

$$s(v_1v_2v_3) = \tfrac{1}{3}\sum_{i=1}^{3} v_i \otimes \prod_{j\neq i} v_j.$$

The map $j : \bigwedge^2 V \otimes V \to V \otimes \operatorname{Sym}^2 V$ sends $(v \wedge w) \otimes z$ to $v \otimes wz - w \otimes vz$. In dimension two it is injective: for a basis e_1, e_2 its images at $z = e_1, e_2$ have distinct independent coordinates. They lie in $\ker m$, whose dimension is $6 - 4 = 2$. Hence (s, j) is the required isomorphism. Tensorial definitions make both maps commute with every linear endomorphism, including singular arrow maps. This is a genuine LR intertwiner on our dessin representation stratum.

For a concrete theta representation take

$$X_1 = \begin{pmatrix} 1 & 1 \\ 0 & 1 \end{pmatrix}, \quad X_2 = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}, \quad X_3 = \begin{pmatrix} 1 & 0 \\ 1 & 1 \end{pmatrix}.$$

Then $M = X_3X_2X_1 = \left(\begin{smallmatrix} 2 & 2 \\ 2 & 3 \end{smallmatrix}\right)$, $t = 5$, $\Delta = 2$, $h_2 = 23$, $h_3 = 105$ and $\chi_{(2,1)} = 10$. The identity is $115 = 105 + 10$. Since the other permitted branch is zero, every forbidden product is zero; this is an actual representation of the infinite elliptic dessin algebra.

In the invertible-arrow open stratum one can gauge X_1, X_2 to the identity, leaving conjugation of M. The invariant-ring proof in the core works also at singular arrows by cycle-trace generation, not by assuming that this gauge choice extends there. The degree-one monodromy coefficients have total arrow degree three; this is why the Hilbert-series denominator uses $3j$, not j.

For the Cauchy kernel choose the Schur subspace inside the tensor power with its inherited Hilbert norm. $M^{\otimes m}$ preserves it and has norm at most $\|M\|^m$. This proves the trace bound for nonnormal matrices as well as diagonal ones. Formal Cauchy identities extend by polynomial continuity in matrix entries; the stated absolute bound then justifies analytic summation. Gram positivity does not assert entrywise positivity or positivity of an unrelated arithmetic form. For three exact test matrices M, $\left(\begin{smallmatrix} 1 & 1 \\ 0 & 2 \end{smallmatrix}\right)$, $\left(\begin{smallmatrix} 0 & 1 \\ -1 & 1 \end{smallmatrix}\right)$, the choice $r = 1/100$ satisfies $rR_0^2 \leq 21/100 < 1$ using Frobenius norms, and all seven principal minors of the rational kernel matrix are nonnegative.

10.3 Two precise functorial limitations

At any actual vertex simple S of A, the split exact sequence $0 \to S \to S \oplus S \to S \to 0$ uses first inclusion and second projection. Termwise Sym^2 has dimensions $1, 3, 1$ at that vertex; its middle kernel contains the mixed monomial as well as the first square, but its first image contains only the first square. Middle homology is S. Thus even split exactness is not preserved. Genuine derived Schur constructions instead use mixed tensor degrees, signed total differentials and suitable flat/projective replacements [32, 33]; their membership in a particular order category must still be checked.

The R-lattice $M = Re_x \oplus Re_y$ with $x = \operatorname{diag}(u, 0)$, $y = \operatorname{diag}(0, u)$ has $x + y = uI$. On $\operatorname{Sym}^2_R M$ the images are $\operatorname{diag}(u^2, 0, 0)$ and $\operatorname{diag}(0, 0, u^2)$. Their sum kills the mixed vector, hence equals

neither uI nor u^2I. This is a specific failure to retain the original central action, not a failure of the polynomial monomial representation functor. Even a common ramified base substitution does not repair this mixed vector.

Semi-invariants retain useful moduli on nilpotent modules. In the theta quiver keep only $a_0, a_3 : v_0 \to v_2$, put their scalar values $(x, y) = (1, \lambda)$ and all other arrows zero. The v_1 simple splits off. Every length-two path acts by zero, so these are continuous completed modules. The two scalars have the same weight g_{v_2}/g_{v_0}, and y/x is a rational invariant on $x \neq 0$, giving the projective pencil $[x : y]$. It is not a regular invariant on the whole affine locus, whose orbit closures meet zero.

For the interval quotient take $x = J_2$, $y = \lambda J_2$. With its top and socle specified, the separated Kronecker coordinates are again $(1, \lambda)$; the relation to the same-weight ratio is exact. The pp formula in the core is stronger than merely noticing a ratio: on M_λ its solution set is $\operatorname{im} x$, whereas on M_μ, $\mu \neq \lambda$, the equation $(\mu - \lambda)xv = 0$ forces $z = xv = 0$. Different parameters are therefore distinguishable by actual positive-primitive conditions even though every cycle trace is zero.

Chapter 11

Full lattice and diamond diagrams

11.1 The nodal exact diagram

The normalization proof of Theorem 4.1 uses BM inside the generic module. Replacing it without justification by $B \otimes_\Lambda M$ could introduce torsion and would not prove the same statement. Once the subdirect residue subspace is decomposed, the branch basis changes lift over complete DVRs; taking inverse images modulo the conductor gives the direct decomposition, not just a decomposition of residues.

Here is the promised diagram, with the terminal short exact column for every position:
The canonical diagram appears in Chapter 21; the proof and calculation below are retained.

Each vertical column is completed by zero above and below. In the bottom row identify $Y/uY = k$ via generator y, $X/uX = k$ via x, and $\Lambda/u\Lambda = k[\epsilon]/(\epsilon^2)$ with $x = \epsilon$, $y = -\epsilon$. Its first map is $1 \mapsto -\epsilon$, its second the augmentation. It is nonsplit as a dual-number sequence. This explicit order reduction, not an unrelated finite algebra, explains its appearance.

For a further Ext calculation, resolve $X \cong \Lambda/(y)$ with successive maps $y, x, \ldots$. On Y the first is multiplication by y and the second zero, so $\mathrm{Ext}^1(X, Y) = Y/yY = k$; on X the first is zero and the second injective, so $\mathrm{Ext}^1(X, X) = 0$. Interchanging branches gives the other two cases. The two nonsplit lattice rows realize the nonzero cross classes. Their Hom spaces are finite over R, generally infinite over k.

11.2 Chebyshev geometry and the exact quotient

The identity $1 - T_n(z)^2 = (1 - z^2)U_{n-1}(z)^2$ follows from $z = \cos\theta$ and the sine/cosine multiple-angle formulas, hence as a polynomial identity. $T_n' = nU_{n-1}$. Therefore β_n has its only zero at infinity, of order $2n$; its n points over one are double since $\beta_n - 1 = T_n^2/(1 - T_n^2)$; and its poles are simple at ± 1 and double at the $n - 1$ interior roots of U_{n-1}. The derivative has no additional critical values. The total ramification is $(2n - 1) + n + (n - 1) = 4n - 2$, as required for a degree-$2n$ rational Belyi map. The dual graph has $n + 1$ vertices of degrees $1, 1, 2, \ldots, 2$, and n edges; it is a connected path. Its planar ribbon embedding is unique, so this identifies the particular nested-loop dessin, not only its passport.

In the quotient retaining orientation Ω, any surviving path traverses each chain edge at most once and either endpoint loop at most once. Its length is at most $n + 1$; thus the killed-arrow ideal contains J^{n+2}. The quotient of A_n and of its completion is the same $H_n(\Omega)$. For $n = 3$ with forward chain arrows, its path counts in lengths zero through four are $3, 4, 3, 2, 1$, total thirteen. These

counts describe a quotient representation category of the infinite order, not a change of primary object.

Inflation is full because the two algebras have exactly the same actions on these annihilated objects. It preserves internal exact sequences. It does not preserve arbitrary ambient extension closure: the two adjacent simples are annihilated by the killed reverse arrow, but the two-dimensional extension having only that arrow nonzero is not. Likewise endpoint-local freeness over $k[l]/l^2$ or $k[r]/r^2$ in [46] is not R-freeness. To transfer an almost-split property to a larger category one must prove the requisite approximation/factorization statements; locally finite quiver criteria [40, 41, 42] do not apply just because each vertex has two arrows. The permitted cycles violate interval/Hom finiteness.

11.3 Nine actual strings and all maps

Take $n = 3$ with $1 \xrightarrow{a} 2 \xrightarrow{b} 3$, loops l, r and $l^2 = r^2 = 0$ in the quotient. Here b denotes the retained forward a_2, not a reverse arrow of the original doubled chain. The following words are in travel order. A superscript $-$ means a formal inverse in a word, not an inverse algebra element.

Object	word	start	dimension vector
A	$l^-abrb^-a^-l^-$	1	$(4,2,2)$
B	l^-ab	1	$(2,1,1)$
C	l^-abr^-	1	$(2,1,2)$
D	$abrb^-a^-l^-$	1	$(3,2,2)$
E	ab	1	$(1,1,1)$
F	abr^-	1	$(1,1,2)$
G	$r^-b^-a^-labrb^-a^-l^-$	3	$(4,3,4)$
H	$r^-b^-a^-lab$	3	$(2,2,3)$
I	$r^-b^-a^-labr^-$	3	$(2,2,4)$

A basis vector is assigned to each visited position $0, \dots, \ell(w)$ at its vertex. A direct letter sends its preceding vector to its following vector; an inverse letter sends the following vector to the preceding vector. Sum these elementary maps for each arrow and put all unmentioned entries zero. This defines every matrix, including repeated loops, and verifies the relations. Set the killed reverse arrows to zero to obtain actual Λ_3-modules.

Their signed grid is

The canonical diagram appears in Chapter 21; the proof and calculation below are retained.

Horizontal projections keep the initial word positions and kill the remaining positions. Horizontal inclusions append the final r^-, keeping old positions. Upper vertical projections delete the initial l^-: old position $j \geq 1$ maps to position $j-1$, and old position zero maps to zero. Lower vertical inclusions prepend $r^-b^-a^-l$: position j maps to $j+4$. Multiply these maps by the indicated signs. For every arrow these formulas commute with the string action, including the two positions adjacent to a cut or join; those are exactly where the loop inverse and chosen orientation matter.

Every small square has two length-two composites whose sum is zero. The exact sequences are

$$0 \to A \to B \oplus D \to E \to 0, \qquad 0 \to B \to C \oplus E \to F \to 0,$$
$$0 \to D \to E \oplus G \to H \to 0, \qquad 0 \to E \to F \oplus H \to I \to 0.$$

For example the maps out of A retain its initial positions $0,1,2,3$ in B and delete its first position in D, so their joint kernel is zero. The maps to E are onto, and the dimension vectors give $\dim A + \dim E = \dim B + \dim D$ vertexwise. The zero composite plus this equality proves exactness.

The other three use precisely the same keep/delete or append/prepend rules; their vertexwise dimension equalities and jointly injective/surjective maps prove exactness. The shipped exact checker constructs all matrices and verifies these assertions independently over $\mathbb{Q}$.

The stronger almost-split property is imported by identifying $E = w = ab$ as the minimal string of type $(2, 2)$ and these exact hook/cohook words with [46, Proposition3.12], using that paper's Propositions2.6/2.9 and Lemma3.10. Opposite-module conventions are reconciled by reversing travel/composition conventions; the closed doubled quiver is isomorphic to its opposite by exchanging the two arrows at each chain edge. Thus intrinsically $\tau E = A$, $\tau F = B$, $\tau H = D$, $\tau I = E$. The rows and columns of this grid are *not* generally short exact, unlike the nodal diagram. Two small cautions in that source are not needed here: a tube-rank label in Example3.9(3) is inconsistent with its listed mouth, and a general homogeneous-tube statement must not be read as a converse classifying every such object as a band. Neither is used to import these four meshes.

11.4 The pure-injective quotient and what its dimension means

For $B_0 = \Lambda/J^2 = k\langle x, y\rangle/(x, y)^2$, words alternate direct and inverse letters because every direct/direct and inverse/inverse adjacent pair is forbidden. Up to rotation and inverse its only primitive band is xy^{-1}; the bridge quiver has two oriented band vertices with no connecting nontrivial bridge. The longest path has length zero. The domestic string results [44, 45] therefore give relative dimension $0 + 2 = 2$, not an automatic dimension three for every one-domestic algebra. Finite-dimensional points have relative CB rank zero; the appropriate Prüfer/adic and one-sided periodic string points have rank one; the unique generic point has rank two. These statements concern the quotient category with its inherited pp structure.

Indeed a pp formula is a finite system of linear equations with existentially quantified unknowns. On a module annihilated by J^2 every coefficient reduces to B_0, so purity and pure-injectivity are preserved in both directions under inflation. The condition $MJ^2 = 0$ is a definable closed condition, giving the closed Ziegler subspace. Its relative ranks need not be ranks in the entire Λ spectrum. This subcategory is not extension closed: $\Lambda/(x^3, y)$ gives an extension of objects killed by J^2 that is not itself killed by J^2. Since u^2 annihilates every object of the subcategory, it contains no nonzero R-free lattice. Ringel's representation-theoretic node construction is also not a Morita equivalence identifying the geometric nodal order with the Kronecker algebra. Infinite-dimensional polynomial string classifications [38] retain finitely-controlled or pointwise-artinian assumptions; an arbitrary complete lattice does not automatically satisfy them as an algebraic A-module.

Chapter 12

The operator extension in detail

12.1 A corrected cycle lift and the integer boundary map

In the permitted-state graph each length-ℓ cycle has a full matrix-corner Laurent/circle coordinate. Its chosen K_1 generator is the full cycle at one state plus the complementary identity. In Fock space $L_c^* L_c = Q_a + v_{s(a)}$. Thus the tempting lift $L_c + (1 - Q_a)$ has overlapping initial pieces and is not an isometry. The corrected lift in Theorem 5.1 first restricts to Q_a:

$$W = L_c Q_a + (1 - Q_a), \quad W^* W = I, \quad I - W W^* = Q_a - L_c Q_a L_c^*.$$

A path starting with a and length greater than ℓ has uniquely the full prefix c followed by another path starting with a. The missing range consists exactly of lengths $1, \ldots, \ell$. Its terminal multiset is the original graph's incident-edge multiset at that cycle vertex. These finite defect vectors all lie in terminal compact ideals; their classes give the claimed unsigned incidence entries. The sign convention affects only the overall sign of the boundary map.

Here is the Smith calculation without spectral assumptions. Choose a spanning tree of the connected underlying graph and one root. By changing signs on one bipartition of the tree, its $E \times V$ incidence submatrix becomes the ordinary oriented tree incidence, with $V - 1$ unit pivots. After eliminating these pivots, a remaining column represents the common root coordinate, and each nontree edge evaluates on it as zero when its fundamental cycle is even, or ± 2 when odd. A loop is the latter case. If all are even the graph is bipartite and this column is zero; otherwise its entries have gcd two, giving the final pivot two. Remaining rows are free. Consequently the kernel/cokernel, then the six-term sequence, give the K groups stated in the core. No theorem reconstructing a cyclic embedding from those groups is asserted.

The compact ideal is block diagonal, not $\mathcal{K}(\mathcal{H})$: prefixes never move terminal vertices, so cross-terminal rank-one maps do not belong to the displayed ideal. This is why the arithmetic insertion must choose one terminal sector. A graph/Cuntz boundary quotient kills the vacua and hence this ideal, so that insertion need not survive passage to the quotient.

12.2 Spectral domain and arithmetic trace

A vector belongs to the adjoint domain of the diagonal N precisely when its coefficientwise product by the real sequence $|p|$ is square summable; this proves self-adjointness on the given domain. The truncations by length prove finite supports are a core. Resolvent coefficients $(|p| - i)^{-1}$ tend to zero and each length space is finite dimensional, giving compactness. Prefix and deletion shift the length

by at most one, so preserve the domain and prove the commutator formulas there. Leibniz then handles the dense polynomial *-algebra.

Every terminal sector has one vacuum and two paths at each length $n \geq 1$, even when some graph endpoints coincide: distinct formal paths remain distinct basis vectors. Thus restriction to a sector still has a simple vacuum, but it need not be a faithful representation of the whole concrete algebra. The trace of e^{-tN^2} is $E + 2E\sum_{n\geq 1} e^{-tn^2}$, distinct from the semigroup trace of N in the core. The trace identity with $\chi(N)$ is only on $\operatorname{Re} s > 1$, where the product is trace class. Analytic continuation of the scalar L-function does not make that same operator product trace class on $\operatorname{Re} s > 0$.

Chapter 13

Abel determinant estimates and the arithmetic program

For completeness the inserted family has a full elementary convergence proof. Since complete nonprincipal character periods sum to zero, $|A_n| \le B_0 \le q_\chi$. On $2^k \le n < 2^{k+1}$,

$$\sum |u_n|^2 \le \frac{B_0^2}{(k+1)^2}, \qquad \sum |v_n(s)|^2 \le M^2 (k+1)^2 r^k, \quad r = 2^{-2\sigma_0} < 1.$$

The second follows from $n^{-s} - (n+1)^{-s} = s \int_n^{n+1} x^{-s-1}\, dx$. Define

$$Q_J(r) = r^J \left(\frac{(J+1)^2}{1-r} + \frac{2(J+1)r}{(1-r)^2} + \frac{r(1+r)}{(1-r)^3} \right).$$

Expanding $(J+j+1)^2$ in the three geometric moment sums proves $Q_J = \sum_{k \ge J} (k+1)^2 r^k \to 0$. Locally uniform ℓ^2 convergence of the holomorphic finite-coordinate vectors proves Banach holomorphy. The linear functional $z \mapsto \sum v_n z_n$ has norm $\|v\|$; this is the complex-linear coefficient pairing, not a conjugate-linear dependence on s.

Write $\mathcal{I}(v)(z) = \sum_n v_n z_n$ for the coefficient functional, distinct from the arrow ideal J. For $N_J = 2^J - 1$, $J \ge 1$, let $K_N = u \otimes (\mathcal{I}(P_N v) - e_1^*)$. Then

$$\|K - K_{N_J}\|_1 \le \sqrt{2} B_0 M \sqrt{Q_J(r)}.$$

Finite summation by parts gives

$$T_N = \det_F(I + K_N) = \sum_{n \le N} A_n (n^{-s} - (n+1)^{-s}) = S_N - A_N (N+1)^{-s}.$$

Adding the last endpoint times e_1^* to the functional recovers S_N exactly because $u_1 = 1$. Its trace-norm error is at most

$$\sqrt{2} B_0 (M \sqrt{Q_J(r)} + B_0 (N_J + 1)^{-\sigma_0}).$$

For arbitrary cutoffs use the last whole dyadic block for an enclosing tail. Galerkin compressions also converge: $\|u - P_{N_J} u\| \le B_0 / \sqrt{J}$. With $C_K = 1 + M\sqrt{(1+r)/(1-r)^3}$,

$$\|K - P_{N_J} K P_{N_J}\|_1 \le \frac{B_0 C_K}{\sqrt{J}} + \sqrt{2} B_0 M \sqrt{Q_J(r)}.$$

The raw endpoint adds the same endpoint bound as before. Isometric insertion by U preserves nonzero singular values, every norm estimate, and the Fredholm determinant. Everything is in one compact ideal of the actual infinite dessin Fock algebra.

For primitive parity a use $g(s) = (q_\chi/\pi)^{(s+a)/2}\Gamma((s+a)/2)$, holomorphic nonzero on $\operatorname{Re} s > 0$. The completed functional is $g(s)\mathcal{I}(v) - e_1^*$, and the completed raw functional is $g(s)(\mathcal{I}(P_N v) + A_N(N+1)^{-s}e_1^*) - e_1^*$. Their determinants are exactly gL and gS_N. Scaling the entire old perturbation would instead give $1 + g(L-1)$, the wrong normalization. Trace-class background is standard [26], while these special rank-one identities follow directly from $\det_F(I + u \otimes \varphi) = 1 + \varphi(u)$.

The metric obstruction is likewise elementary. With inner product conjugate-linear in its first variable, symmetry of $T = u \otimes \varphi$ implies

$$\varphi(z)\langle u, u\rangle = \overline{\varphi(u)}\langle u, z\rangle.$$

A nonzero φ cannot have $\varphi(u) = 0$; thus all symmetric maps with that range have proportional functionals. For $n \geq 2$ the coefficient ratios at $\sigma + it$ and σ satisfy

$$\frac{v_n(\sigma + it)}{v_n(\sigma)} = \frac{\sigma + it}{\sigma}\, n^{-it}(1 + O_{\sigma,t}(1/n)).$$

Integers nearest $\exp(2\pi k/|t|)$ and $\exp((2\pi k + \pi)/|t|)$ give opposite nonzero subsequential limits. Hence no common positive metric symmetrizes this family for all critical heights. Restricting any proposed metric on the inserted family to $U\ell^2$ gives the same contradiction. An independently constructed positive comparison is still a possible different route.

13.1 Precisely how classical and recent sources enter

Over a finite field the Grothendieck trace-formula architecture writes a zeta function as alternating determinants on finite-dimensional cohomology; Deligne's purity supplies the weights of eigenvalues [22, 20, 21]. If a weight-w Frobenius is $q^{w/2}$ times a positive-metric unitary, each eigenvalue has modulus $q^{w/2}$. The norm conclusion is immediate; constructing that form or proving purity is the deep step. Symplectic/alternating duality $F^t J F = q^w J$ alone only pairs reciprocal eigenvalues and does not give their individual moduli. An order resolution or Schur-positive character expansion alone cannot replace this step in number-field cohomology.

The recent primary comparisons used here are specific:

- [53, Theorem3.12, pp.15–17] gives finite abelian covers with local orbit length $\log p$ and genuine arithmetic monodromy. It suggests what a descent-equipped order functor should match; it does not identify our u or a face cycle with a prime without such a functor.

- [54, Propositions4.1–4.3, Theorem4.6] builds finite-prime local models with archimedean Fourier factors. Its Hilbertian isomorphisms are not an asserted uniformly controlled common isometry for our arithmetic approximants. [55, Proposition3.1] gives determinacy of a particular local moment problem; positive finite Hankel matrices are not the missing global Weil inequality.

- [56, Theorem3.6] gives discreteness/lower boundedness for the stated restricted Weil form. Theorem5.10, pp.23–24, additionally assumes a finite truncation, $\lambda > 1$, a simple even bottom eigenfunction ξ and $\delta_N \xi = 1$. Its quotient metric is positive and its regularized determinant is $-i\lambda^{-iz}\widehat{\xi}(z)$. Passing to the infinite arithmetic comparison, including the ground-state and convergence conditions in its §8, is not automatic.

- [57, Theorem6.1, p.13] assumes a specified real distribution on an interval, the corresponding trigonometric-core form, essential self-adjointness, and an isolated simple lowest eigenvalue with reflection-even eigenfunction. Its Fourier transform then has only real zeros. Our degree operator's vacuum and length parity are not those data.
- [60, Theorem6.1] gives prime-indexed periodic packets in its dynamical construction, not one proved orbit per prime with the full desired cohomological trace formula. [62] isolates an obstruction to a particular real-coefficient cohomology with full functorial Artin multiplicities; it is not a theorem excluding every Hilbert–Pólya operator. The current survey [58] distinguishes remaining steps. The recent Picard-monoid construction [59] is not already a polarized abelian variety with a number-field purity theorem. [63] uses RH as an input in the cited spectral identification, so cannot be used circularly as its proof.

These are targeted theorem/proof readings, not claims that every page of every cited work has been audited. The preserved review records specify actual page coverage and hashes, including inaccessible original PDFs where the modern primary proof was used instead.

13.2 A concrete sequence of remaining constructions

First choose a number-field model of a Belyi pair, its finite covers and conductor/ramification data. Construct actions or exact correspondences on the labelled order/module or analytic categories that recover the local prime Frobenius polynomials. The elliptic twists show why omitting this datum loses information. Test at more than one arithmetic model of the same complex dessin, not just its topological passport.

Second specify a finite projective or otherwise determinant-admissible complex and its *independent* arithmetic endomorphism. Prove exact-sequence/determinant-line compatibility. Use semi-invariants to encode presentation degeneracies, and Schur/exterior characters to calculate local factors, retaining every central-action and nilpotence distinction above. General derived Schur replacements require their own construction; termwise Schur on the resolution fails the exhibited split sequence test.

Third prove the same local factors, gamma terms and conductor give a globally convergent or correctly regularized determinant. Trace-norm convergence must be established for the chosen analytic perturbations; strong convergence is insufficient in general. The inserted Abel family supplies this topology for an elementary scalar encoding, but has neither the independent arithmetic spectrum nor the positive metric needed for zero location.

Finally produce an independently proved positive form or self-adjoint determinant-preserving realization. One acceptance criterion is the Hurwitz target in the core; another is equality with the precise arithmetic quadratic form in the recent spectral theorems, with all their domain, ground-state, normalization and limit hypotheses verified. The local Cauchy kernel, K-group computation and graph spectral dimension alone meet none of these final comparison hypotheses. No remaining condition is silently asserted true here. Dirichlet, Hecke, Dedekind and Artin scopes must remain individually recorded.

Chapter 14

Graph pullbacks and completions with specified fibers

14.1 The complete Fock graph presentation

For graph C* conventions the edge operator has initial projection at its range vertex and final projection under its source vertex. Put

$$S_h = L_{a_h} Q_{a_{\tau h}}, \qquad T_h = L_{a_h} v_{[\tau h]}.$$

The two right-hand projections are orthogonal subprojections of $L_{a_h}^* L_{a_h} = Q_{a_{\tau h}} + v_{[\tau h]}$. Thus $S_h^* S_h = Q_{a_{\tau h}}$, $T_h^* T_h = v_{[\tau h]}$, and $S_h^* T_h = 0$. For distinct h the prefix ranges are orthogonal, so every other cross-edge CK1 product is zero. Their final projections sum to Q_{a_h}. The vacua have no outgoing edges. The resulting graph has $|H| + |Q_0|$ vertices and $2|H|$ edges, with cycles only among its state vertices and exits to the paired sinks. Gauge covariance and nonzero vertices prove faithfulness as well as surjectivity. This explains precisely how a graph C* algebra is obtained without imposing the incompatible original-quiver CK1 relations.

The graph integer map has one column for each regular vertex h and no sink columns:

$$d : \mathbb{Z}^H \longrightarrow \mathbb{Z}^H \oplus \mathbb{Z}^{Q_0}, \qquad d(e_h) = e_h - e_{\tau h} - b_{[\tau h]}.$$

The state part of $d(x) = 0$ says $x_h = x_{\tau^{-1} h}$; the sink part then says the constant cycle coefficients are in $\ker B_\Gamma$. For the cokernel choose a base state on each cycle and successively use $q_h = q_{\tau h} + b_{[\tau h]}$ to eliminate all other states. Closing the cycle leaves $\sum_{h \in c} b_{[h]} = 0$, while its base state remains free. Hence

$$\operatorname{coker} d \cong \mathbb{Z}^{\#\,\mathrm{cyc}\,\sigma} \oplus \operatorname{coker} B_\Gamma.$$

This gives an independent integral check of the previous compact-extension computation. The full labelled graph, diagonal projections and edge families carry more information than those coarse groups.

14.2 A genuine application of the supplied pullback theorem

For the interval dessin let $E_2 = G_\alpha$ have loop vertices q_x, q_y, common sink b, loops s_x, s_y and exits t_x, t_y. Remove b to obtain F_2, two disjoint loop graphs. Its quotient is exactly $C(\mathbb{T})^2$. Let E_1 have the same three vertices and edges $e_{x,n} : q_x \to b$, $e_{y,n} : q_y \to b$, $n \geq 0$, and no other edges. Remove b to obtain the edgeless F_1, whose algebra is $\mathbb{C}^2$.

The inclusions $F_i \subset E_i$ are admissible in the precise sense of [65, Definition3.1]: the removed sink set is hereditary/saturated; all edges ending at retained vertices are retained; no vertex emits infinitely many edges into the removed set and a positive finite number outside it. For E_1 the latter outside count is zero, which is allowed. Define a path functor identity on objects by

$$e_{x,n} \longmapsto s_x^n t_x, \qquad e_{y,n} \longmapsto s_y^n t_y.$$

Images of distinct positive edges are prefix incomparable: at the first different length one has an exit and the other a loop. Source vertices and identities cause no exception. The image consists of all positive pointed paths in E_2, whose last edge is not a short loop; vertices are included separately. The first graph is acyclic and the removed sink has no short loop in the second. All hypotheses of Theorem3.4 are now verified, giving

$$C^*(E_1) \cong \mathcal{T}_A \underset{C(\mathbb{T})^2}{\times} \mathbb{C}^2.$$

The quotient of $C^*(E_1)$ by its sink ideal is $\mathbb{C}^2$, and its sink ideal is compact operators on the paths ending at b. The graph is acyclic, hence its algebra is AF. There are no regular vertices in E_1: the two sources are infinite emitters and b is a sink. The graph K matrix has no columns and three rows, yielding $K_0 = \mathbb{Z}^3$ and $K_1 = 0$.

This is a positive instance of the new source inside our infinite dessin framework. It is a constants-symbol subalgebra of the Fock model, not the entire Fock algebra or the formal order. For a permitted cycle of length greater than one, its full closed path is itself pointed in this paper's convention, so an object-preserving acyclic path functor cannot cover all such paths. That simple obstruction prevents copying this example to every cycle by analogy. The sink-extension Theorem4.3 also requires a specified original pullback, injective gluing maps, annihilation of glued vertex projections and coverage of all paths ending at the gluing sinks. A two-to-one identification of sink occurrences within one branch does not satisfy these assumptions without another construction.

14.3 Diagonal residue obstruction and the corrected pullback

Suppose one tries to glue full matrix-circle blocks by their diagonal values. In M_2, $E_{12}E_{21} = E_{11}$, although both factors have zero diagonal. Thus the proposed boundary map is not a homomorphism. For $\ell > 1$, more generally any *-homomorphism from the full matrix-circle algebra to a commutative C*-algebra has equal images of $E_{ii} = E_{ij}E_{ji}$ and $E_{jj} = E_{ji}E_{ij}$. These images are both orthogonal and equal, hence zero; summing the diagonal units kills the image of one and then the whole homomorphism. This works for nonunital *-maps as well: a nonzero image of one would give a nonzero unital corner map. The failure concerns full matrix-circle blocks, not the already-restricted fiber algebra D_ℓ.

Evaluation $D_\ell \to \mathbb{C}^\ell$ is a genuine surjective *-homomorphism since at one both factors are diagonal. Its kernel is $C_0(\mathbb{T} \setminus \{1\}, M_\ell)$. Pulling it back along the repeated-coordinate embedding $\Delta_\alpha : \mathbb{C}^{Q_0} \to \mathbb{C}^H$ gives exactly the fiber conditions of $\mathcal{D}_H$. The paired constant diagonal matrices give a unital *-section. Hence this is a split C*-extension, not a claim deduced from a merely commutative diagram of nonmultiplicative maps. Since $\mathbb{T} \setminus \{1\} \cong \mathbb{R}$, each kernel block is a matrix suspension of $\mathbb{C}$, with $K_0 = 0$, $K_1 = \mathbb{Z}$. Split exactness now proves Theorem 5.3.

The pairing is retained in this specified presentation, not recovered from the unmarked isomorphism class. In particular arbitrary constant permutation matrices can relabel a branch's diagonal positions, so the spherical and elliptic theta examples have isomorphic unmarked $\mathcal{D}_H$ despite different ribbon orderings. Keep the labelled arrows, cycle cuts and distinguished fiber as

additional data when using it as a dessin model. Also the central branch unitary zI_ℓ represents ℓ times the primitive corner generator in K_1, not a new independent generator per matrix entry.

14.4 Generation, faithfulness and the exact failure of extension

Let ρ denote the polynomial representation in Theorem 5.3. Its vertex images are orthogonal projections summing to one. Every arrow image lies under the correct source and range vertex projections. A forbidden product vanishes because the intermediate state indices are τh and $\alpha\tau h$, which are distinct. Thus ρ is an actual algebra map, not a fitted characteristic polynomial.

For a full turn starting at h, $\rho(c_h) = w_c E_{hh}$ with $w_c(z) = g(z)^{|c|}z$. It is a normal element of $\mathcal{D}_H$. The function w_c vanishes exactly at one; away from that point $w_c/|w_c| = z$. Therefore it is injective on all of $\mathbb{T}$, and a homeomorphism to its image. Nonunital continuous functional calculus at $\rho(c_h)$, using functions zero at zero, gives every fE_{hh} with $f \in C(\mathbb{T})$, $f(1) = 0$. The zero condition makes the result vanish in all other matrix entries and branches.

To obtain an off-diagonal entry choose the unique shortest permitted path from state h to j within its cycle, of length $r > 0$ and cut-crossing count $\kappa \in \{0, 1\}$. Its image is $g^r z^\kappa E_{hj}$. Given $f(1) = 0$, multiply it on the left by the already available diagonal element

$$\frac{f(z)z^{-\kappa}}{g(z)^r + \varepsilon} E_{hh}, \qquad \varepsilon > 0.$$

The product is $fg^r/(g^r + \varepsilon)E_{hj}$ and converges uniformly to fE_{hj}. Indeed near one $|f|$ is small, and on the complement of that neighborhood g^r has a positive minimum. These two bounds prove uniform convergence without assuming division by a function vanishing at the fiber. Thus the generated algebra contains the whole kernel ideal. Its quotient contains the entire paired diagonal algebra, so it is $\mathcal{D}_H$.

For injectivity let a finite linear combination of path images vanish. At one, all positive paths vanish, forcing the coefficients of the original vertex idempotents to vanish. Fix a state-matrix entry in a branch. Its remaining paths have lengths $r + m\ell$ and images

$$g^r z^\kappa \sum_m a_m (g^\ell z)^m E_{hj},$$

with a possible initial zero coefficient if $r = 0$. On $\mathbb{T} \setminus \{1\}$ the prefactor is nonzero, and w_c has infinite range, so the polynomial has infinitely many distinct zeros. All coefficients are zero. Repeating by entries proves faithfulness.

No boundedness improvement repairs full formal extension. Pick any $z_0 \neq 1$, let $\lambda = w_c(z_0)^{-1}$, and consider $1 - \lambda c_h$. Its inverse $\sum_{m \geq 0} \lambda^m c_h^m$ exists in the arrow-adic order. The matrix image at z_0 has a zero eigenvalue on state h, so cannot be invertible. A unital algebra homomorphism would preserve the inverse identity, an immediate contradiction. This proves nonextension without any continuity assumption. It does not rely on the false claim that a norm greater than one alone forces a geometric-series denominator to be singular. Scaling g to have norm at most one still does not allow arbitrary formal series.

The two valid operator completions are structurally different. In the bipartite case $K_0(\mathcal{T}_A) = \mathbb{Z}^{E+1}$ differs from $K_0(\mathcal{D}_H) = \mathbb{Z}^E$; in the nonbipartite case the former has the extra $\mathbb{Z}/2$ torsion. Neither is automatically the norm completion of the arrow-adic topology. No invariant constructed solely from these complex objects distinguishes the rational twists already exhibited; arithmetic descent remains a separate input.

14.5 Invariant basis number and the Leavitt comparison

If a unital k-algebra R_1 has a unital representation on a nonzero finite-dimensional vector space V, an isomorphism $R_1^m \cong R_1^n$ induces an isomorphism $V^m \cong V^n$ by tensoring. Comparing dimensions gives $m = n$. This proves invariant basis number without requiring the representation to be faithful. Every original dessin algebra and completed order has a vertex-simple representation, and each permitted-state Laurent algebra has evaluation onto a nonzero matrix block. The Leavitt algebra of G_α has the same finite-dimensional representation after killing the vacuum sinks and evaluating one boundary cycle. Thus all of these examples have invariant basis number. Leavitt's free-rank phenomena are important elsewhere but cannot be assumed in this class simply from the name LPA.

For the specified finite graph G_α, the usual complex Leavitt path algebra embeds in $C^*(G_\alpha)$ with norm-dense image, using its canonical generators and involution [67, Theorem7.3]. Those generators are the split S_h, T_h, not the unsplit original gentle arrows equipped with ordinary CK1. This is the exact algebraic/analytic twin relationship applicable here.

Chapter 15

String finiteness and the smooth quotient surface

For a finite quiver, finitely controlled implies finitely generated: choose one finitely generated submodule containing each vertex space and sum the finitely many choices. Conversely the module itself works. Stabilization at each of finitely many vertices likewise gives ordinary artinianity. Crawley–Boevey's polynomial theorem therefore applies directly to our finitely generated and artinian polynomial modules, with its primitive injective bands in the latter case. It does not assert a classification of every possible algebraic module.

The formal branch $M = k[[x]]$, $y = 0$, gives a precise counterexample to that stronger assertion. An A-endomorphism F preserves each $x^n M$. For the polynomial truncation f_n of f, one has $F(f) - f_n F(1) \in x^n M$. Separatedness and all n imply $F(f) = fF(1)$. Hence $\mathrm{End}_A(M) = k[[x]]$ is local, so M is indecomposable. If an ordinary word module has $y = 0$, its word cannot contain y or y^{-1}, since such a letter would create a nonzero basis map. A reduced word using only $x^{\pm 1}$ cannot change direction. Its nonperiodic choices are finite strings, the free ray $k[x]$, and the opposite locally nilpotent ray. The first and last fail injectivity of x, and the free ray fails invertibility of $1 - x$. A two-sided pure-x word is periodic, and in every associated band x acts invertibly. On M, x is injective but not onto. None of these possibilities equals M. This includes arbitrary nonzero indecomposable Laurent coefficient modules, not just finite bands.

For a general nonzero completed lattice, finite generation over A would imply finite generation over $k[u]$. Every nonzero polynomial in u is a power of u times a unit of $k[[u]]$, so it acts injectively. The lattice would be a nonzero finite free $k[u]$-module. Multiplication by $1 - u$ on such a free module has a nonzero cokernel, but the same map is invertible over $k[[u]]$. This proves the contradiction in the core. The strictly descending u-adic chain separately excludes artinianity. These are actual failed hypotheses, not merely unverified applicability.

Lekili–Polishchuk's precise convention takes $D(A)$ to mean perfect dg modules. Their locally gentle definition allows infinite dimension, and their smoothness criterion says no forbidden cycle; properness says no permitted cycle. Both fail for every nonempty closed dessin here. Their Theorem3.2.2 therefore does not apply to the full polynomial dessin algebra, even though it applies to suitable infinite-dimensional quotients. The OPS finite-dimensional geometric-intersection theorem has its own grading, local-system, boundary and exceptional coincident-curve conditions. An unrestricted statement that every gentle/string derived object is an ungraded geodesic crossing is not justified by either source.

For the interval quotient $B = A/(y) = k[x]$, global dimension one follows from its PID structure: every submodule of a finite free module is free, giving finite projective resolutions of all finitely

generated modules. Thus $\mathrm{Perf}(B) = D^b(\mathrm{mod}_{\mathrm{fg}}\, B)$ in grading zero. Its permitted loop gives one unmarked boundary component and its length-one forbidden thread x gives one stop on the other boundary, the annulus of the smooth partially wrapped model. This construction uses the infinite polynomial quotient, not the completed $k[[x]]$ category.

Finally the quotient does not give a derived fully faithful embedding by ordinary inflation. Resolve B over A by the alternating maps $y, x, y, x, \ldots$. Hom into B gives $0, x, 0, x, \ldots$, so $\mathrm{Ext}^{2r}_A(B, B) = B/xB = k$ for every $r \geq 1$, whereas all positive self-Ext over B vanish. Consequently its inflation is not derived fully faithful, and the perfect B-module B becomes an infinite-projective-dimension A-module. A surface functor transferring these categories must account for that difference explicitly.

Chapter 16

Arithmetic character sums and the cellular determinant

16.1 The uniform calculation and its exceptional cases

In the finite-field calculation write $k = \mathbb{F}_q$; the final order-insertion paragraphs specify a separate characteristic-zero coefficient field. The curve is nonsingular because its discriminant is $-432d^2 \neq 0$, and has one point at infinity. Thus counting $1 + \eta(x^3 - d)$ in each affine fiber proves the trace formula in Theorem 6.2 directly. The multiplicative group of k is cyclic. When $3 \mid q - 1$, a nonzero cube has three roots and a noncube none, detected by $1 + \rho(t) + \bar{\rho}(t)$. At zero this expression is one, using the *constant* one. This convention differs from mathlib's zero-extended trivial multiplicative character, whose value there is zero. The proof and formal components never replace the constant by that character.

The substitution is valid including $u = 0, 1$ because the characters are zero extended:

$$\sum_t \rho(t)\eta(t - d) = \sum_u \rho(du)\eta(du - d) = \rho(d)\eta(-d) \sum_u \rho(u)\eta(1 - u).$$

Multiplying by the initial minus sign gives exactly $\pi_d + \bar{\pi}_d$. Changing ρ to ρ^{-1} conjugates both j and λ_d, so does not change the unordered roots or the integer trace. Characteristics two and three and $d = 0$ are excluded throughout. For a residue characteristic $p \equiv 2 \pmod 3$, odd extension degrees have trace zero by cube bijectivity, whereas even degrees lie in the split case; no claim of zero trace for all extensions of that characteristic is made.

For completeness, choose a nonzero $\mathbb{F}_p$-linear functional on k and compose it with $t \mapsto e^{2\pi i t/p}$ to obtain ψ. For a nontrivial multiplicative character χ,

$$\begin{aligned} |G(\chi)|^2 &= \sum_{u \neq 0} \chi(u) \sum_{y \neq 0} \psi((u - 1)y) \\ &= (q - 1) - \sum_{u \neq 0,1} \chi(u) = q. \end{aligned}$$

Both orthogonality identities follow by translating a finite character sum by an element on which its character is not one. The $z = 0$ terms in $G(\chi)G(\theta)$ sum to $\theta(-1) \sum_x (\chi\theta)(x)$, zero if the product is nontrivial. For $z \neq 0$ set $x = zu$, $y = z(1 - u)$ to obtain $J(\chi, \theta)G(\chi\theta)$. Coprime orders three and two ensure our product character is nontrivial. Taking norms proves $|J|^2 = q$, rather than presuming a desired positivity. This is the classical calculation reviewed directly in [68]; no claim of novelty is made.

At seven the five nonzero summands are $-\omega^2, -\omega, \omega, -\omega^2, 1$, so $J = 3 + 2\omega$. Multiplication by $\lambda_2 = \omega^2$ gives $-1 - 3\omega$. Their matrices are

$$\begin{pmatrix} 3 & -2 \\ 2 & 1 \end{pmatrix}, \qquad \begin{pmatrix} -1 & 3 \\ -3 & 2 \end{pmatrix}.$$

Both have determinant seven and the same positive Gram form; their traces are four and one. The arithmetic distinction between the rational models now comes from explicit finite-field characters, while their complex dessin is unchanged. Exact prime-field diagnostics test every nonzero twist for the listed nineteen primes; the theorem itself is uniform in prime powers and rests on the proof, not that finite list.

16.2 Deck descent and the actual cellular basis

The CM cellular construction in this and the next subsection assumes $q \equiv 1 \pmod 3$. The cover $\beta = x^3/d$ has deck transformations $(x, y) \mapsto (\zeta_3 x, y)$ and $(x, y) \mapsto (x, -y)$, generating $C_3 \times C_2 = C_6$. Label their actions on darts by $\sigma = (012)(345)$ and $\alpha = (03)(14)(25)$; they commute. Their product is $(0\,4\,2\,3\,1\,5)$. They preserve the edge pairing and both successor permutations. Consequently they define algebra automorphisms, preserve J and its completion and fix the common u. The projective-resolution maps are compatible since $\theta_g(a_h x) = a_{gh}\theta_g(x)$. Twisting module actions gives an exact functor and the corresponding relabelled resolution, without an assumption of arithmetic Frobenius.

Over an algebraic closure choose $a^3 = d$, $b^2 = d$ and $\varphi(x, y) = (ax, by)$ from E_1 to E_d. For the q-power Frobenius morphisms,

$$\varphi^{-1} \circ F_d \circ \varphi = h_d \circ F_1, \qquad h_d(x, y) = (a^{q-1}x, b^{q-1}y).$$

When $q \equiv 1 \pmod 6$, the first coordinate is the cubic residue $d^{(q-1)/3}$ and the second $\eta(d)$. Choose the embedding of cube roots defining ρ consistently. On the deck eigenspace labelled by dx/y, the *deck* character is their ratio, $\rho(d)/\eta(d) = \lambda_d$. This labels the twist, not a nonzero Frobenius action on differentials in characteristic p: that pullback is zero. An inverse convention for arithmetic Galois Frobenius changes the action convention; our displayed formula uses the geometric curve morphism explicitly. The character-sum polynomial fixes its trace without an unstated convention.

Orient the three geometric edges as in the core. Each occurs twice with opposite orientation in the single face, so $\partial_2 = 0$; ∂_1 has columns $(-1, 1)^t$. The homology basis $b_0 = e_0 - e_1$, $b_1 = e_1 - e_2$ has Gram matrix G inherited from the orthonormal cellular edges. Rotation sends $b_0 \mapsto b_1$, $b_1 \mapsto -b_0 - b_1$, exactly C. The involution swaps vertices and reverses each edge, hence acts as $-I$ on this plane. Thus multiplication by $A + B\omega$ is literally $AI + BC$, not merely an abstract matrix with the correct trace. Expanding its determinant gives the norm $A^2 - AB + B^2$, and multiplication of this norm proves $M^t G M = qG$. The displayed positive quadratic decomposition proves the metric independently of unknown zeros.

16.3 An explicit chain lift and determinant cancellation

Let P be the three-cycle permutation on C_1. Write $\pi_d = A + B\omega$ and $\gamma = A + B$. Define

$$F_1 = AI + BP, \qquad F_0 = \frac{1}{2}\begin{pmatrix} 1+\gamma & 1-\gamma \\ 1-\gamma & 1+\gamma \end{pmatrix}, \qquad F_2 = q.$$

These are rational cellular maps. The sum-zero plane of F_1 has restriction M_d, while the edge-sum line has eigenvalue γ. On C_0 the vertex-sum line has eigenvalue one and the boundary line eigenvalue γ. It follows directly that $F_0\partial_1 = \partial_1 F_1$; the other chain condition is zero. The induced maps on H_0, H_1, H_2 have characteristic factors $1-T$, $P_d(T)$, $1-qT$. In particular

$$\det(I - TF_0) = (1-T)(1-\gamma T), \qquad \det(I - TF_1) = (1-\gamma T)P_d(T).$$

The contractible boundary factor cancels, giving

$$\frac{\det(I - TF_1)}{\det(I - TF_0)\det(I - TF_2)} = \frac{P_d(T)}{(1-T)(1-qT)}.$$

This equality holds as a formal rational identity and by cancellation at apparent intermediate singularities. The alternating determinant with exponent $(-1)^i$ is its inverse; signs are not changed silently. The standard elliptic trace and determinant theorem identifies the right side with the finite-field curve zeta function [22, 20]. The elementary proof above supplies its actual trace and norm here; it does not independently reprove the general cohomological trace formalism.

This lift also has the required descent character on H_1. Write the base Jacobi value as $j = A_1 + B_1\omega$. Then h_d acts there by $\eta(d)C^r$ when $\rho(d) = \omega^r$, so twisting $A_1 I + B_1 C$ multiplies its eigenvalue j by λ_d, giving $M_d = A_d I + B_d C$. Reduction modulo $C^2 + C + I = 0$ gives the coefficients of π_d. If the corresponding full three-edge polynomials differ by $c(I + P + P^2)$, the coefficients are $c = 0, \eta(d)B_1, \eta(d)A_1$ for $r = 0, 1, 2$ respectively; they have the same homology action and differ only on the contractible boundary pair. An explicit homotopy sends $v_- - v_+$ to $c(e_0 + e_1 + e_2)$ and the vertex sum to zero. Its $\partial H + H\partial$ is the edge difference and the boundary scalar difference $3c$. Thus this normalization is compatible up to rational chain homotopy. It is not asserted to arise from a continuous Frobenius map of the complex Riemann surface, or to give a canonical integral Tate-lattice comparison.

16.4 A decisive restriction at inert primes

For $q \equiv 2 \pmod 3$, Frobenius sends ζ_3 to ζ_3^{-1}. Thus its cohomological relation with the cubic deck action requires $MC = C^2 M$ (the pullback convention yields the equivalent reversed relation). On a real two-dimensional carrier with our fixed C, comparing entries forces

$$M = \begin{pmatrix} a & b \\ a+b & -a \end{pmatrix}, \qquad \det M = -a^2 - ab - b^2 = -\tfrac{1}{2}(a^2 + b^2 + (a+b)^2) \le 0.$$

The genuine elliptic Frobenius determinant is $q > 0$, so such a real carrier is impossible. This is an elementary instance of the real-coefficient cohomology issue discussed by Deninger [62]. It does not rule out a complex realization, a larger real space, a different coefficient theory or an infinite-dimensional height operator. In particular it is not a refutation of an unrestricted Hilbert–Pólya approach.

There is a local complex replacement, with its limitations explicit. The conjugation matrix

$$K = \begin{pmatrix} 1 & -1 \\ 0 & -1 \end{pmatrix} \quad \text{satisfies} \quad K^2 = I,\ KC = C^2 K,\ K^t G K = G.$$

Set $F = i\sqrt{q}\,K$ on the complexification. Then $\operatorname{tr} F = 0$, $\det F = q$, $F^2 = -qI$ and $F^* G F = qG$, giving the correct trace-zero local factor. This is a two-dimensional complex, hence four-dimensional

real, local model. This is a one-q model; independently choosing $i\sqrt{q}K$ for each extension field does not give Frobenius-power compatibility. Such a family would need to start from one F_p and prove the comparison for its powers. The inserted $i\sqrt{q}$ is not a canonical rational or integral Tate comparison or a construction of a global positive determinant. The fixed real two-dimensional model has genuinely failed; its surviving split-prime construction and this complex alternative must not be conflated.

16.5 The remaining order and global spectral obligations

The deck automorphism has finite order, so over characteristic zero its eigenvalues are roots of unity. Its permutation action alone cannot supply a weight-one norm $\sqrt{q} > 1$. The independently computed Jacobi coefficient is essential. The above construction closes the specified local cellular benchmark, not a reconstruction of arithmetic from bare complex dessin data.

For these order-insertion tests take $\Lambda = \Lambda_{\mathbb{Q}}$ over $R = \mathbb{Q}[[u]]$, with rational residue simples and $V_d = \mathbb{Q}^2$. This is distinct from reducing the matrix in characteristic p, where its determinant p becomes zero. There are two concrete failures of a naive order lift. First, a residue simple S_e is annihilated by all arrows; $S_e \otimes V_d$ with independent endomorphism $I \otimes M_d$ has the desired underlying finite-$\mathbb{Q}$ determinant. But its projective dimension is infinite by the nonterminating forbidden resolution, and σ permutes all three vertex labels, so choosing S_e does not preserve full deck symmetry. This is not a perfect-order determinant realization.

Second, $\Lambda \otimes V_d$ is free of rank two over Λ, but the underlying R-rank of Λ is eighteen. In an R-basis, the Λ-*linear* $I \otimes M_d$ is eighteen copies of M_d, hence its determinant is P_d^{18}. On $e\Lambda$ the rank is six: each of its two hereditary branches contributes a row of length three, with finite conductor gluing changing no rank. This gives P_d^6. Restoring all three residue vertices similarly gives P_d^3 for this untwisted tensor action. A genuine deck-semilinear map permutes factors as well and need not have these powers; it requires its own comparison. No justified determinant-line root or multiplicity-removing functor is supplied here.

Finally fix a nonzero rational d when forming its global Euler product from good reductions. Local purity cannot be multiplied into a global zero-location proof. Since each reciprocal root has modulus $\sqrt{p}$, the product of good elliptic factors is absolutely convergent on $\operatorname{Re} s > 3/2$. The local zeros at $\operatorname{Re} s = 1/2$ are outside that initial product argument. The finite-field curve zeta's center is $1/2$; the global elliptic L-function's unshifted center is one, or $1/2$ only after a half-unit shift. Neither is an identification with the primitive Dirichlet function targeted earlier. Ramified factors, archimedean normalization, adelic descent and the global determinant-preserving positive spectral limit remain separate obligations.

Chapter 17

The theta order interface beyond tensor insertion

We prove the multiplicity theorem and construct its actual functorial escape, retaining the completed infinite order as the primary algebra. The finite algebra used below is its specified special fiber, not a replacement primary example.

17.1 Normal forms and the special fiber

Index the three quiver vertices by $i \in \mathbb{Z}/3$. Let $a_i^{\pm} : i \to i-1$ be the two permitted branches. Only paths of constant sign survive. Write $p_{i,r}^{\pm}$ for their length-r paths and $c_i^{\pm} = p_{i,3}^{\pm}$. The split order model gives

$$ue_i = c_i^{+} + c_i^{-}, \qquad p_{i,r+3}^{\pm} = up_{i,r}^{\pm} \quad (r \geq 1).$$

The second identity follows because multiplying by the opposite full cycle switches branch and vanishes. An R-basis of $e_i\Lambda$ is

$$e_i,\ a_i^{+}, a_i^{-},\ p_{i,2}^{+}, p_{i,2}^{-},\ c_i^{+}.$$

These span by the two identities and completion. They are independent in the two hereditary branch orders: off-diagonal entries distinguish the four shorter paths; the remaining combination $xe_i + yc_i^{+}$ has paired diagonal entries $(x + uy, x)$, forcing $x = y = 0$. This is an actual power-series basis, not just a generic rank calculation.

In $B = \Lambda/u\Lambda$, the same eighteen basis vectors remain. A product of nonempty paths is zero if branches switch or the total length exceeds three; the length-three minus path is the negative of the plus path. A length-four path is u times a length-one path, hence zero. The radical $\mathfrak{r}$ is the positive-length ideal, with $\mathfrak{r}^4 = 0$ and $B/\mathfrak{r} = \mathbb{Q}^3$. Each e_iB has radical-layer dimensions $(1, 2, 2, 1)$ and has two composition factors at each vertex, so the Cartan matrix is

$$C_B = \begin{pmatrix} 2 & 2 & 2 \\ 2 & 2 & 2 \\ 2 & 2 & 2 \end{pmatrix}.$$

The Cartan matrix alone does not establish the endomorphism formula; the path filtration supplies that step.

17.2 All projectives and all linear endomorphisms

Let J be the completed arrow ideal. Every $1-x$, $x \in J$, has its convergent formal inverse; the same applies to $1-yx$ for arbitrary $y \in \Lambda$. Thus J is contained in the Jacobson radical, and $\Lambda/J = \mathbb{Q}^3$ implies equality. If P is finitely generated projective, choose a basis of each multiplicity space in P/PJ and lift its vectors into Pe_i. This gives

$$Q = \bigoplus_i (e_i\Lambda)^{n_i} \longrightarrow P$$

inducing an isomorphism on tops. Nakayama makes this surjective. Its split kernel N is finitely generated and satisfies $N/NJ = 0$, hence vanishes by Nakayama. This proves the decomposition for every such projective without an unproved freeness assumption over Λ.

An entry $e_i\Lambda \to e_j\Lambda$ is left multiplication by an element of $e_j\Lambda e_i$. Its degree-zero part is possible only when $i = j$ and is a scalar multiple of e_i. For f on P, write its top blocks A_i and set $A = \operatorname{diag}_i A_i$. The remainder has positive path length. On P/uP it strictly raises the path filtration. At each of the two length-one and two length-two slots, and the length-zero and length-three slots, the degree-preserving map is exactly A. In an increasing-degree basis the matrix is lower block triangular. For the formal upper-block convention order the six slots instead as

$$3+,\ 2+,\ 2-,\ 1+,\ 1-,\ 0.$$

There are six diagonal copies of A, allowing arbitrary positive-length off-diagonal entries. Hence

$$\det_{\mathbb{Q}}(1 - T\overline{f}) = \det_{\mathbb{Q}}(1 - TA)^6.$$

Specializing the finite R-matrix determinant coefficientwise gives Theorem 6.3. No simultaneous diagonalizability, tensor form of f, arithmetic purity or deck equivariance is required.

For a bounded complex $P^\bullet$ with a linear chain map $f^\bullet$, fix either parity convention and write

$$D_R(P^\bullet, f^\bullet; T) = \prod_j \det_R(1 - Tf^j)^{(-1)^j}.$$

The inverse convention gives the reciprocal and the same obstruction. Each factor has constant coefficient one, so its reduction is nonzero. Specialization of this displayed rational expression is defined and equals

$$\left[\prod_{j,i} \det_{\mathbb{Q}}(1 - TA_i^j)^{(-1)^j}\right]^6.$$

Cancellations cannot change the divisibility of divisor multiplicities. In particular adding contractible projective pairs cannot repair it. Here a perfect endomorphism means a bounded-projective chain representative; no statement about an arbitrary infinite resolution or nonperfect lattice is inferred.

Generically put $N = \sum_i n_i$. As a right module over $M_3(K) \times M_3(K)$, P_K is the sum of N copies of each branch's three-dimensional row simple. Any commuting map has matrices $F_+, F_- \in M_N(K)$ on their multiplicity spaces, so

$$\det_K(1 - Tf_K) = \left(\det(1 - TF_+)\det(1 - TF_-)\right)^3.$$

For example, left multiplication by c_i^+ on $e_i\Lambda$ has ordinary characteristic determinant $(1-uT)^3$, so a sixth power is forced only at the special fiber. This also shows why generic reduced determinants

and ordinary determinants must not be conflated. We do not need an integral choice of either generic matrix to derive the cube obstruction.

For the local numerator, suppose a sixth power $A(T)^6/B(T)^6$ equals a nonzero quadratic $P(T)$ with $B \neq 0$. Polynomial degrees would give $6 \deg A = 2 + 6 \deg B$, impossible. This includes repeated roots. For the complete curve zeta, $q > 1$ and $|a_q| \leq 2\sqrt{q}$ imply $P(1) > 0$ and $P(1/q) = P(1)/q > 0$. Its two denominator roots are distinct uncancelled simple poles, again impossible for a sixth power. These arguments do not assume global RH or Artin holomorphy.

17.3 An actual integral determinant functor of multiplicity one

Project to one hereditary cycle order $B_+ \subset M_3(R)$. Its standard column $W_+ = R^3$ is a left Λ-module with $e_i W_+ = R\epsilon_i$. More explicitly a_i^+ acts by $u^{w_i} E_{i,i-1}$, with $(w_0, w_1, w_2) = (1, 0, 0)$ for one cut, while every a_i^- acts by zero. These matrices respect travel-order multiplication as written: $E_{ij}E_{jk} = E_{ik}$. Each plus full cycle is uE_{ii}, so the common central u acts by uI. Nonempty paths tend to zero u-adically, giving the action of the *full completed order*, not just its polynomial subalgebra.

This (Λ, R)-bimodule defines $- \otimes^{\mathbf{L}}_{\Lambda} W_+$. On a bounded complex of finite projectives it is computed termwise, and its terms are finite free over R. An acyclic bounded projective complex is split contractible, and tensor preserves that contraction. Thus it descends to a genuine triangulated functor $\Phi_+ : \mathrm{Perf}(\Lambda) \to \mathrm{Perf}(R)$. No assertion of exactness on all un-derived modules or flatness of W_+ is needed. In fact W_+ has infinite projective dimension as a left Λ-module: Λ is noetherian, being finite over the central noetherian ring R, and W_+ is finitely generated. A finite projective resolution could thus be chosen with finite terms. Every left projective term has R-rank divisible by six by the same top argument, whereas W_+ has rank three, contradicting the alternating rank identity. This does not prevent tensoring bounded right-projective complexes.

The isomorphism $e_i\Lambda \otimes_\Lambda W_+ \to e_i W_+$ sends $e_i\lambda \otimes w \mapsto e_i\lambda w$. Its inverse is $e_i w \mapsto e_i \otimes e_i w$, proving rank one directly. Consequently $e_i\Lambda \otimes_{\mathbb{Q}} V$, with action $1 \otimes M$, maps to $R \otimes_{\mathbb{Q}} V$ with action M. Its characteristic determinant is $\det(1 - TM)$ *once*. On the regular input, the three vertex lines survive and the factor appears three times. A full symmetric vertex orbit therefore is not silently replaced by a chosen vertex.

Different cycle cuts give different integral column lattices. Their generic modules are isomorphic, and branch endomorphism characteristic determinants agree by generic conjugation; equality over K implies equality over R. Nevertheless an equivariant lattice is stronger data. Branch selection also changes the functor: for the unit $f = e_i + c_i^+$ on $e_i\Lambda$, the plus determinant is $1 + u$ and its characteristic factor is $1 - (1 + u)T$. Under α, c_i^+ becomes c_i^-, whose plus action is zero; the determinant becomes one and the characteristic factor $1 - T$. Since these scalar automorphisms differ, no natural isomorphism of determinant lines can make this chosen branch functor α-invariant for all endomorphisms. Pairing both branch functors makes this branch-swap data available but doubles the local factor on the basic-projective tensor input.

17.4 The exact rotation obstruction and its allowed escape

Consider any left Λ-lattice supported on the plus branch, free of rank one at every e_i, and with central u acting as multiplication by u. Write its plus arrow scalars as $a_0, a_1, a_2 \in R$. Their product is u, so they are nonzero and their nonnegative valuations sum to one. A strict R-linear lift of σ that intertwines the order rotation permutes the three vertex lines by units. Its arrow-intertwining

identities make the three arrow valuations equal. Thus $3m = 1$ for an integer $m \geq 0$, impossible. The argument already excludes an invertible intertwiner; imposing that its cube be one cannot help.

There is an elementary coefficient proof of the same obstruction. Intertwining by units makes vanishing of the three arrow constant coefficients equivalent. Their product has constant coefficient zero, so all three vanish. Factoring u from two of them would make $a_0a_1a_2$ divisible by u^2, contradicting that it equals u. This is the literal power-series component checked in Lean.

After the explicit ramified change $R' = \mathbb{Q}[[v]]$, $u = v^3$, *choose a new column lattice* with all three plus arrow maps $vE_{i,i-1}$. Cyclic permutation of the basis is an order-three intertwiner fixing v. The full-cycle relation is now $v^3 = u$ as required. Merely base-changing the old cut lattice gives arrow valuations $(3, 0, 0)$ and does not produce that linearization. This construction changes the base and lattice and does not establish an unramified descent comparison. It also still needs both branches to address α.

The rotation obstruction is about actual central-fixed rank-one branch lattices, not every possible determinant-level natural transformation. The branch characteristic determinants may be rotation-invariant through their generic reduced interpretation even without a lattice lift. Neither this theorem nor the ordinary-determinant obstruction excludes semilinear correspondences, other determinant categories or infinite analytic models.

The next arithmetic task is therefore narrower and more concrete: choose a branch functor, a ramified lattice or a different equivariant replacement only after proving the required descent comparison, and construct its prime action from arithmetic data rather than feed in the already known M_d. Bad primes, integral/Tate comparison, local-to-global conductor and gamma factors, and a positive spectral limit for the same completed function remain unresolved.

Chapter 18

Explicit tame descent and the full-deck coefficient test

The preceding ramified column is a precisely specified object of the original infinite theta-order framework. We compute its descent instead of treating a semilinear group action as automatically effective integral descent. All constructions below are over the new *mathematical* coefficient field $k = \mathbb{Q}(\omega)$; they do not replace the original rational model or its source history.

18.1 The action on genuine formal series

Write $R = k[[u]]$, $S = k[[v]]$ with $u = v^3$. Every $f(v) \in S$ has a unique expansion $f_0(u) + vf_1(u) + v^2 f_2(u)$. Hence S is finite free of rank three over R. Let $g(f)(v) = f(\omega v)$. On $W = S^3$ with basis $\epsilon_0, \epsilon_1, \epsilon_2$, put

$$\Theta(f_i\epsilon_i) = \omega^{-i} f_i(\omega v)\epsilon_i, \qquad C\epsilon_i = \epsilon_{i+1} \quad (i \bmod 3).$$

The plus arrow a_i^+ acts by $vE_{i,i-1}$, the minus arrow by zero. The cyclic product is uE_{ii} and arbitrarily long permitted paths tend to zero v-adically, so this is a module over the full completed order. For $i = 0$, the coefficient in the target under Θa_i^+ is ωv and in the source under $a_i^+\Theta$ it is $v\omega^{-2}$; these agree. The same calculation for $i = 1, 2$ proves compatibility with every arrow and therefore the completed algebra. Also $\Theta^3 = 1$.

For all coefficients, not just a truncated diagnostic,

$$[v^n](\Theta w)_i = \omega^{n-i}[v^n]w_i.$$

Thus invariance is equivalent to support in $n \equiv i \pmod 3$, and

$$L = W^\Theta = \bigoplus_{i=0}^{2} Rb_i, \qquad b_i = v^i\epsilon_i.$$

The arrows on this basis have scalar valuations $(1, 0, 0)$, since $a_0^+ b_2 = ub_0$, $a_1^+ b_0 = b_1$ and $a_2^+ b_1 = b_2$. Tame invariants return the old cut column rather than a new strict unramified rotation carrier.

The extension $R \subset S$ is faithfully flat. Its group quotient has a fixed point at $v = 0$, however, and is not a μ_3-torsor there: $S \otimes_R S$ is not the product of the three group graphs at that point. Semilinear μ_3 data are therefore not by themselves the full effective fpqc descent data for a vector bundle on $\operatorname{Spec} R$. The adjunction comparison is explicitly

$$S \otimes_R L \longrightarrow W, \qquad \operatorname{diag}(1, v, v^2),$$

with determinant $v^3 = u$, cokernel $S/(v) \oplus S/(v^2)$ and closed-fiber rank one. This proves its failure to be an isomorphism directly. The quotient stack $[\operatorname{Spec} S/\mu_3]$ retains the stabilizer at zero; its coarse space is $\operatorname{Spec} R$. This local root construction is described in [71]. Alper's vector-bundle criterion [70, Theorem10.3, pp.2386–2387] distinguishes precisely this nontrivial closed-fiber action from effective coarse descent. We use these sources for interpretation; all displayed module comparisons have already been calculated explicitly.

18.2 A projective symmetry survives on the endomorphism order

Directly on the three basis vectors,

$$\Theta C = \omega^{-1} C\Theta = \omega^2 C\Theta.$$

Thus the rotation does not preserve L. The corrected vC does preserve it, since the factor ω from $g(v)$ cancels the commutator. In the invariant basis it has matrix A from the core, with $A^3 = uI$ and $\det A = u$. This is generically invertible and gives an order-three projective rotation after passing to the fraction field; it is not an invertible integral lattice intertwiner.

There is nevertheless a true algebra-level escape. Let $F = (f_{ij})$ be an arbitrary S-linear endomorphism of the *carrier* W, not assumed Λ-linear. Its tame conjugate has entries $\omega^{j-i} g(f_{ij})$. Invariance therefore means

$$f_{ij} \in v^{[i-j]_3} R,$$

where $[r]_3 \in \{0, 1, 2\}$ is the least residue. After replacing ϵ_i by $b_i = v^i \epsilon_i$, the entry becomes

$$v^{[i-j]_3 + j - i} R = \begin{cases} R & i \geq j, \\ uR & i < j. \end{cases}$$

Consequently $\operatorname{End}_S(W)^\Theta$ is exactly the hereditary three-cycle order B_+ in the previously fixed cut convention. Multiplication is preserved by conjugation in the generic matrix algebra; no entrywise projection is being misused as a ring map.

The scalar commutator cancels under conjugation. Hence $\operatorname{Ad}(C)$ commutes with tame conjugation, descends to B_+, and is represented there by $\operatorname{Ad}(A)$ because v is central. Since $A^3 = uI$, this algebra automorphism has order three and preserves the hereditary order despite A^{-1} not lying in its unit group. This explains the surviving rotation of the branch order; it does not construct a rotated integral column or a Frobenius endomorphism. The other branch is identical and the paired-diagonal residue gluing still defines the original Λ_R inside $B_+ \times B_-$. Its label automorphisms remain the already constructed dessin symmetries.

18.3 Determinants, Schur weights and the boundary correction

On $\det_S W = \bigwedge_S^3 W$, both $\det D = 1$ and $\det C = 1$. Thus this line has the scalar action g and genuinely descends to the trivial R-line with trivial rotation. This is a real positive determinant-level result. On the other hand,

$$b_0 \wedge b_1 \wedge b_2 = v^3(\epsilon_0 \wedge \epsilon_1 \wedge \epsilon_2),$$

so the natural comparison

$$\det_R(W^\Theta) \longrightarrow (\det_S W)^\Theta$$

is multiplication by u, with cokernel $R/(u)$. Both lines are abstractly free of rank one, but their given comparison is not an isomorphism. It records the divisor of the ramified lattice change and

cannot be discarded when claiming a determinant-preserving integral comparison. For the *same* operator vC, the scalar on both determinant lines is u; naturality is respected. Comparing $\det C = 1$ with $\det A = u$ without the extra factor v would instead compare different operators.

For an induced nonzero homogeneous tensor/Schur carrier of degree d, the same functorial calculation gives commutator ω^{-d}. The two induced invertible operators commute exactly when $3 \mid d$. This is a group-action statement, not sufficient for coarse vector-bundle descent: $\mathrm{Sym}^3(S^3)$ has closed-fiber inertia multiplicities $(4, 3, 3)$ for the three characters, so its inertia is still nontrivial even though the commutator vanishes. A mixed tensor $W^{\otimes a} \otimes (W^*)^{\otimes b}$ has net exponent $a - b$; $\mathrm{End}_S(W)$ is the concrete net-zero escape computed above.

These carrier functors are not automatically functors on unital Λ-modules. Applying global $\bigwedge^3$ to the rank-one vertex idempotents kills all three, although their sum is the identity before applying this nonadditive functor. Vertexwise $\bigwedge^3$ of the rank-one column is zero. Vertexwise Sym^d for $d > 0$ does retain the monomial zero relations but changes the full-cycle scalar from u to u^d. It leaves the fixed central-action problem unless that change of base is included. None of the Schur cancellations supplies the missing two-dimensional arithmetic carrier.

18.4 Why the local Jacobi plane does not descend by this repair

On constant vectors Θ acts by D. It sends the sum-zero cycle $\epsilon_0 - \epsilon_1$ to $\epsilon_0 - \omega^2\epsilon_1$, whose coordinate sum is nonzero. It also does not preserve the edge-sum line. Thus this order column is not already a tame descent of the earlier cellular chain complex or its H_1.

For the split local coefficient $\pi = a + b\omega$, the natural rotation polynomial on the balanced carrier is $F = aI + bC$. Since g fixes $a, b \in k$, its commutator is

$$\Theta F - F\Theta = b(\omega^2 - 1)C\Theta.$$

It is nonzero when $b \neq 0$. In particular both prime-seven twists $(a, b) = (3, 2), (-1, -3)$ fail this test. Prime powers with $b = 0$ are not excluded by this particular calculation; that exceptional commutation does not prove a global comparison.

Replacing C by vC gives an invariant carrier map, whose matrix on L is $aI + bA$. Direct determinant expansion gives

$$\det(1 - T(aI + bA)) = (1 - aT)^3 - ub^3T^3.$$

At $u = 0$ this is $(1 - aT)^3$, not $P_d(T)$, which has degree two. The weighted replacement therefore does not recover the established local factor. Neither C nor vC is generally Λ-linear: they intertwine its rotation. Their sums with the identity have not thereby become endomorphisms in the ordinary order-module category. They are legitimate carrier tests, with exactly the failures just exhibited.

18.5 The full orbit is the normalization, and arbitrary kernels still fail

Twist the tame action by the three characters: $\Theta_r = \omega^r\Theta$, $r \in \mathbb{Z}/3$. Its fixed lattice is

$$L_r = \bigoplus_i Rv^{[i-r]_3}\epsilon_i,$$

with the cut at arrow r. Rotation gives an actual R-linear isomorphism $C : L_r \to L_{r+1}$, taking each displayed basis vector to the corresponding next one. Thus the direct sum of all three cuts has

a strict rotation, at the cost of vertex rank three. Before invariants the orbit is S-rank nine and underlying R-rank twenty-seven; its invariant lattice is R-rank nine. The original orbit S-lattice still has nontrivial inertia at zero and need not be the base change of these invariants.

There is an explicit integral identification of the invariant orbit. In the standard hereditary order B_+, the j-th column has entries $u^{t_{ij}}R$ with $t_{ij} = 1$ for $i < j$ and zero otherwise. In its natural column basis the arrow coefficient has exponent

$$\delta_{i0} + t_{i-1,j} - t_{ij} = \delta_{ij}.$$

Its cut is therefore exactly j, identifying B_+e_j with L_j. Consequently $\bigoplus_r L_r \simeq B_+$ as left Λ_R-modules. Adding the exchanged minus-branch orbit gives the hereditary normalization $\widetilde{\Lambda} = B_+ \times B_-$, with full deck action and vertex rank six. On a basic projective with external arithmetic matrix this yields P_d^6.

In fact this normalization replacement restores the *ordinary determinant for every linear endomorphism*. For a finite projective P, the inclusion $P \to P \otimes_\Lambda \widetilde{\Lambda}$ becomes an isomorphism over $K = k((u))$, since $\widetilde{\Lambda}_K = \Lambda_K$. Its extended endomorphism and the original f therefore have the same generic characteristic determinant. Both polynomials have coefficients in R, which injects into K, so they agree already over R. The bounded-complex ratios also agree. This symmetry repair returns the previously refuted ordinary determinant, not merely its external tensor exponent.

For completeness we now prove Theorem 6.4 for arbitrary finite free kernels, without assuming they split integrally into branches or are projective over the order. Let

$$p_\pm = u^{-1} \sum_i c_i^\pm \in \Lambda_K.$$

These are central generic branch idempotents. Put $W_\pm = W_0 \cap p_\pm(W_0 \otimes_R K)$. They are finite free R-lattices spanning those generic summands, stable under Λ_R. Rotation fixes them, branch swap exchanges them. Their vertex ranks are constant on each branch: the arrow maps become invertible over K because a full cycle acts by u. Write these ranks r_+, r_-. No equality $W_0 = W_+ \oplus W_-$ over R has been assumed.

On the plus branch choose integral vertex bases. The three square arrow matrices X_0, X_1, X_2 have product uI_{r_+} around every cycle. Thus

$$\sum_i \operatorname{val}_u(\det X_i) = r_+.$$

The rotation lift is an invertible matrix between each pair of vertex lattices, hence its determinant is an R-unit. Its intertwining identities make all three arrow determinant valuations equal. Therefore $r_+ = 3m_+$; similarly $r_- = 3m_-$. The branch-swap lift gives $r_+ = r_-$, so generic rank additivity gives $\operatorname{rank}_R e_i W_0 = 6m$. Only invertible intertwiners were used; imposing the full deck group relations is stronger than necessary.

For the general determinant statement reduce modulo u and filter each vertex by $e_i J^a(W_0/uW_0)$. Since $J^4 \subset u\Lambda$, this is a finite filtration. In $\Phi_{W_0}(P) = \bigoplus_i (e_i W_0)^{n_i}$, positive-length entries of f raise the filtration, while the degree-zero matrix A_i repeats on its i-th graded pieces. Their total dimension is $6m$. Multiplying the graded characteristic determinants proves the asserted power for *every* f. Flatness of the kernel as a left Λ-module was not used; bounded right projectives compute the derived functor and give finite free R-outputs. Denominator constant terms one permit specialization of every bounded alternating ratio. The regular left kernel Λ_R attains $m = 1$, making the rank bound sharp.

The remaining arithmetic problem therefore requires a different hypothesis: a genuinely twisted determinant target, a specified non-linear/semilinear correspondence with its own comparison, or an infinite analytic realization. One cannot continue searching inside this closed finite-free, fixed-center, full-deck kernel class for a missing small estimate. The local purity and descent computations supply no global positive spectral limit. Moreover $k/\mathbb{Q}$ has degree two: ordinary restriction of scalars sends a rational local factor P_d to P_d^2. A relative k-determinant must not silently be identified with the former rational determinant. Finite-field curve zeta, global elliptic L-functions and Dirichlet GRH keep their distinct scopes and centers.

Chapter 19

Literal formal and computational coverage

`DessinNodeOrder.lean` defines the fiber-product subring of two genuine power-series rings. Its twenty printed declarations establish diagonal injectivity, nonzero branches, branch zero products, factorizations and annihilator identities, regular u, augmentation surjectivity and exactness, bijective power-series coordinates, square-zero alternating maps, three kernel-equals-range statements and projectivity of the actual regular/product terms. The coordinate statement is a bijection; the file does not yet package it as an R-linear equivalence or declare the general order freeness theorem.

`DessinInvariantComponents.lean` has ten printed declarations: three CK1-collapse components, determinant change of bases with an actual inverse, symmetric-square zero, identity, composition and relation preservation, the degree-two Schur character identity, and the nodal central-parameter obstruction. Its symmetric-square matrix is

$$\operatorname{Sym}^2\begin{pmatrix} a & b \\ c & d \end{pmatrix} = \begin{pmatrix} a^2 & ab & b^2 \\ 2ac & ad+bc & 2bd \\ c^2 & cd & d^2 \end{pmatrix}.$$

This is the ordinary monomial basis convention; composition is proved universally over any commutative ring. The full project suite prints types and foundational dependencies; no theorem about general LR decompositions, abstract derived equivalences, C* K-theory or GRH is inferred from these component proofs.

The exact Python checks use fixed rational matrices and exhaustive finite sets, with no random seeds. They construct the displayed permutations and paths, evaluate rational Cartan identities, check formal Euler products through a stated degree, verify unrestricted-output Fock basis identities on bounded input sets, compute incidence Smith data, construct all nine string representations and signed maps, and test Young/LR/Cauchy examples and the seven-element-field counts. Those diagnostics are reproducible finite corroboration. The universal infinite claims rest on the proofs or explicitly cited theorems, and the arithmetic comparison remains open.

The ten additional `DessinPullbackComponents.lean` statements use literal two-state matrix units and a subring of a product of arbitrary rings, with equal images under two ring homomorphisms. The section hypothesis of its projection-surjectivity theorem is explicit. It does not construct a topological C*-pullback or assert that a commutative diagram is automatically a pullback. The no-residue theorem is universal algebra, with a nontrivial commutative target; it is not an assumed arithmetic positivity premise.

`ThetaJacobi.lean` adds sixteen declarations, with actual finite fields and MulChar rather than

a hypothetical arithmetic kernel. Its norm theorem invokes mathlib's proved Jacobi inverse-product identity and derives characteristic inequality from finiteness; all three character nontrivialities are explicit. Other declarations verify the rescaling substitution, CM multiplication/polynomial/Gram positivity components, exact prime-seven matrix polynomials, and actual 2-by-3 boundary/3-by-2 cellular inclusion, homology restriction, three-edge determinant and rational vertex chain-map equations. It does not formalize elliptic point counting or an order-to-Frobenius functor.

The final three `ThetaJacobi.lean` declarations verify the actual conjugation matrix identities and the real two-dimensional inert-prime determinant obstruction. They do not assume any unknown zero location or exclude complex/enlarged spectral realizations.

`ThetaProjectiveMultiplicity.lean` adds five declarations. Arbitrary upper-block matrices with six identical diagonal blocks have the sixth-power determinant and characteristic-factor formulas. A literal power-series constant-coefficient map specializes matrix determinants, for each scalar parameter; the coefficientwise polynomial/order interface remains prose. A nonzero quadratic cannot be a rational-function sixth power, proved by polynomial degrees. Three genuine power series with product X cannot have the cyclically equivalent residue-vanishing conditions forced by invertible rank-one arrow intertwiners. The derivation of the path blocks from the full theta order, projective decomposition, derived branch functor and deck intertwiners is not fully formalized. The complete analytic suite now has208printed declarations; no hidden arithmetic positivity assumption occurs.

`ThetaTameDescent.lean` adds twelve printed component declarations (and one auxiliary polynomial lemma). They use genuine formal-power-series rescaling and its coefficient operator, proving the coefficient action, invariant-coefficient vanishing and the tame operator's cube. Actual three-state matrices satisfy the clock/rotation commutator, comparison determinant and closed-fiber equations, weighted-cycle cube and full characteristic polynomial, noncommutation of a nonconstant rotation polynomial, trivial whole-column determinant actions, and the same-operator comparison for vC. The classification of all fixed series by residues, full order action and hereditary invariant algebra, stack interpretation, valuation/rank theorem, general coefficient-functor filtration and GRH are not thereby fully formalized. No unknown positivity premise is assumed.

Chapter 20

Explicit covering calculations and derived transport

Here are the linear details behind Proposition 7.2. Write $x_i = F_{ii}$ for diagonal entries of an endomorphism F preserving the three vertex spaces. Solving $FX = XF$, $FY = YF$ leaves a common diagonal c and six free entries. More explicitly its only off-diagonal entries are

$$F_{02} = F_{13} = t_0, \quad F_{15} = t_1, \quad F_{35} = F_{46} = t_2, \quad F_{48} = t_3, \quad F_{68} = F_{70} = t_4, \quad F_{72} = t_5.$$

All others vanish. This can be checked column by column using the displayed partial permutation matrices X, Y: their nonzero columns specify all constraints, and the zero columns force the remaining entries to zero. If $r_0, \dots, r_5$ are the six radical generators in the displayed order, the only nonzero products are

$$r_0 r_2 = r_1, \qquad r_2 r_4 = r_3, \qquad r_4 r_0 = r_5.$$

Every product of three radical generators vanishes. Thus the listed space is a local algebra, not merely a seven-dimensional vector-space count.

For clarity, the Euler equations used in the obstruction are

$$\begin{aligned} &d_1 - d_0 = 1, && d_3 - d_2 = 1, && d_4 - d_3 = 1, \\ &d_6 - d_5 = 1, && d_7 - d_6 = 1, && d_0 - d_8 = 1, \\ &d_1 - d_2 = 1, && d_4 - d_5 = 1, && d_7 - d_8 = 1, \end{aligned}$$

where $d_i = D_{ii}$. Subtract the last row from the first two rows after summing: the left side is zero and the right side is three. This is an identity in every characteristic-zero coefficient field. In characteristic three this particular obstruction disappears; no conclusion there is asserted. The closed walk

$$a_0 b_0^{-1} a_0 \, a_1 b_1^{-1} a_1 \, a_2 b_2^{-1} a_2$$

explains the computation: six forward steps and three backward steps have net lift displacement three. Arbitrary changes of basis cannot remove the Euler contradiction. This is an algebraic lift obstruction, not an assertion about the topological fundamental group of the compact torus.

The Hom formula also has a constructive perfect-complex version. Each cover representable is finite support and pushes to its corresponding vertex projective. For a bounded complex $P^\bullet$ of finite sums of such representables and a bounded finite-support complex $N^\bullet$, decompose the entire Hom complex by deck displacement. The differential preserves that displacement, and only finitely

many displacements occur across the finitely many degrees. Taking cohomology gives

$$\mathrm{Hom}_{D(B_L)}(FP^\bullet, FN^\bullet[q]) \simeq \bigoplus_g \mathrm{Hom}_{D(\widetilde{B}_L)}(P^\bullet, T_g N^\bullet[q]).$$

Bounded projective complexes compute these derived Hom spaces, which justifies the equality rather than assuming every module perfect. This elementary proof does not assert an equivalence of entire bounded derived categories or cover objects of infinite projective dimension. The broader precovering statement has its exact hypotheses in [74, Theorem 4.28].

Finally the branch involution $\sigma(a_i) = b_i$, $\sigma(b_i) = a_i$ preserves B_L. In characteristic different from two one can study $B_L \# k[\mathbb{Z}/2]$ through equivariant representations. For the nine-dimensional module above, $\operatorname{rank} X = 6$ and $\operatorname{rank} Y = 3$, so $M \not\simeq \sigma M$. The sum $M \oplus \sigma M$ has canonical swap equivariance. This is an explicit orbit construction; it neither chooses a nonexistent fixed-object linearization nor identifies this skew group algebra with the particular GLS algebras in [75].

`CoveringBand.lean` verifies the actual matrices' mixed products, cubic nilpotence and absence of an arbitrary Euler matrix. The general covering, module endomorphism classification, derived functors and the surface/Fukaya equivalence are not fully formalized by these three matrix statements. Exact rational code additionally checks thirty-six Hom pairs, the full radical multiplication and the two Hom-control defects.

Part III

AR components, derived categories and Hom-control

Chapter 21

Almost-split geometry, Hom control and determinant relations

The infinite locally gentle algebra and its completed order are the primary objects in this chapter. Finite algebras enter through specified quotients or special fibers, with their categories named. This distinction makes the pictures more useful: a mesh is an exact sequence in an actual category, a tube is an entire component, and a diagram of limiting words records additional objects rather than suggesting an unproved equivalence.

We work over an algebraically closed field k of characteristic zero unless stated otherwise. Paths are written in travel order, so ab traverses a first. We use right projectives e_iA with this convention. Column representations may instead be read over the opposite algebra; all displayed arrow-intertwining identities use the usual column convention. This change of handedness does not identify different categories or change any of the explicit Hom dimensions.

21.1 What an AR picture specifies

For an infinite polynomial dessin algebra $A = kQ/I$, write $\operatorname{mod}_{\mathrm{fd}} A$ for finite-dimensional modules. Its arrow-adic completion is $\Lambda = \varprojlim A/J^n$. A continuous finite-dimensional Λ-module kills some J^n. A lattice over an order finite over $R = k[[u]]$ is instead a finitely generated module free over R. A nonzero lattice has infinite dimension over k and is not a continuous finite-length module. These distinctions persist even when the same quiver describes the arrow actions.

The polynomial string classification has its own finiteness boundary. For a finite quiver, finitely controlled modules are precisely finitely generated modules: sum the finitely many finitely generated submodules containing their vertex spaces. Pointwise artinian is likewise equivalent to artinian, by taking a common stabilization index over the finitely many vertices. Thus [83, Theorems 1.2–1.4] genuinely applies to finitely generated and artinian modules over our infinite polynomial string algebras, with its specified strings, finite-dimensional bands and primitive injective bands. It does not say that every module, or every infinite word, satisfies these finiteness conditions.

A nonzero completed node lattice is outside both classes after restriction to $A = k[x, y]/xy$. Indeed A is finite over $k[u]$, $u = x + y$. If the lattice were finitely generated over A, it would be finite torsion free over the PID $k[u]$, hence finite free of positive rank. Multiplication by $1 - u$ would then have a nonzero cokernel. On a $k[[u]]$-lattice it is a unit, a contradiction. The strictly decreasing chain $M \supset uM \supset u^2M \supset \cdots$ also excludes artinianness. The lattice classification below is proved independently.

An almost-split sequence in an exact category is a nonsplit conflation

$$0 \longrightarrow L \xrightarrow{i} E \xrightarrow{p} M \longrightarrow 0$$

whose first map factors every nonsection out of L and whose second map factors every nonretraction into M. Minimality removes redundant summands. The Auslander–Reiten translate is $\tau M = L$. Irreducible maps give the arrows of the AR quiver; projective vertices do not have nonsplit almost-split sequences ending there. The mesh relation says that the two component composites, with the appropriate signs, sum to zero. It is an identity between module maps. It is not, without a further construction, a polynomial relation between functions on a representation variety.

A quotient $A \twoheadrightarrow B$ gives a fully faithful exact inflation $\operatorname{mod} B \to \operatorname{mod} A$. Its image is the full subcategory annihilated by the kernel. That image need not be extension closed. Even when it is closed under the relevant exact operations, an almost-split factorization in the smaller category need not factor every map from the larger category. Throughout, the word *intrinsic* records this boundary.

21.2 The corrected elliptic-theta face example

21.2.1 Permutations, genus and the actual algebra

The six-dart data

$$\sigma = (1\,2\,3)(4\,5\,6), \qquad \alpha = (1\,4)(2\,5)(3\,6)$$

are transitive and commute. With composition acting rightmost first,

$$\sigma\alpha = (1\,5\,3\,4\,2\,6), \qquad \phi = (\sigma\alpha)^{-1} = (1\,6\,2\,4\,3\,5).$$

The uncontracted dessin has two black vertices, three white vertices, six edges and one face. Thus $2 - 2g = 2 + 3 - 6 + 1 = 0$, and $g = 1$. After contracting the clean white vertices the counts are two vertices, three edges and one face. A global Belyi realization is $y^2 = x^3 - 1$, $\beta = x^3$.

Its algebra has vertices x_1, x_2, x_3 and arrows

$$a_1 : x_1 \to x_2, \quad a_2 : x_2 \to x_3, \quad a_3 : x_3 \to x_1,$$
$$b_3 : x_1 \to x_2, \quad b_1 : x_2 \to x_3, \quad b_2 : x_3 \to x_1.$$

The forbidden travel pairs are

$$a_1b_1,\ b_1a_3,\ a_3b_3,\ b_3a_2,\ a_2b_2,\ b_2a_1.$$

This is the same presentation as the corresponding right-to-left list after reversing each product. For an explicit dart identification use

$$1 \mapsto a_1, \quad 2 \mapsto a_3, \quad 3 \mapsto a_2, \quad 4 \mapsto b_3, \quad 5 \mapsto b_2, \quad 6 \mapsto b_1.$$

Then the face is the forbidden indexing cycle $a_1b_1a_3b_3a_2b_2$. It repeats vertices; it does not create six distinct primitive idempotents. The algebra is infinite dimensional because its two permitted three-cycles can be traversed indefinitely. Its completed order has R-rank 18.

Proposition 21.1 (A face unfolds to a complex, not a six-vertex quotient)**.** *Put $B_6 = k\widetilde{A}_5/J^2$, with the cyclic orientation. Every representation of the theta algebra gives a B_6-representation by unfolding the face:*

$$(V_1, V_2, V_3, V_1, V_2, V_3), \qquad (D_1, \ldots, D_6) = (a_1, b_1, a_3, b_3, a_2, b_2).$$

The resulting functor is additive, exact and faithful, but is not full. There is no unital surjection from this theta algebra to B_6.

Proof. Every consecutive product $D_{i+1}D_i$ is zero by the displayed relation list. A theta morphism repeats its three vertex maps at the two corresponding face positions, so it gives a complex morphism and is recovered from those repeated maps. Kernels and cokernels are computed vertexwise, proving exactness and faithfulness. On the vertex simple at x_1 its image is the sum of the two B_6-simples at positions $1, 4$. The original endomorphism space has dimension one and the image endomorphism space dimension two. Thus fullness fails.

If a unital surjection to B_6 existed, compose it with $B_6 \twoheadrightarrow k^6$. Every theta arrow joins distinct idempotents. Its image in this commutative algebra is therefore zero, since those two idempotent images are orthogonal. The image is generated by the three vertex idempotents and has dimension at most three, whereas a surjection to k^6 would have image dimension six. This is impossible. □

For a fixed theta dimension vector (d_1, d_2, d_3), the arrays of six matrices satisfying the equations are indeed the same affine variety as the circular complex variety of dimensions $(d_1, d_2, d_3, d_1, d_2, d_3)$, after identifying the repeated vector spaces and bases. The acting groups differ:

$$G_A = \prod_{i=1}^{3} \mathrm{GL}_{d_i}, \qquad G_{B_6} = \prod_{i=1}^{6} \mathrm{GL}_{d_i}, \qquad d_{i+3} = d_i.$$

The first embeds diagonally in the second. Its orbits, Hom spaces and orbit closures cannot be read by forgetting that diagonal restriction. For example $a_1 = 1, b_3 = \lambda$, all other arrows zero, gives pairwise nonisomorphic theta representations when $\lambda \in k^\times$ varies. Their unfolded complexes are all isomorphic: the two disjoint nonzero edge maps may be rescaled independently by G_{B_6}.

21.2.2 Correct rank strata and the incidence resolution

The standalone circular-complex model remains valuable. Let $V_i = k^{d_i}$, with indices modulo six, and set

$$\mathcal{X}(\mathbf{d}) = \{(D_i) : D_i : V_i \to V_{i+1}, \quad D_{i+1}D_i = 0\}.$$

A feasible rank vector satisfies

$$0 \le r_i, \qquad r_{i-1} + r_i \le d_i.$$

At each vertex choose a decomposition into the incoming image, a homology space and an outgoing complement. Each outgoing complement maps isomorphically onto the next incoming image. This proves that ranks classify orbits for the *six-factor* group. It also decomposes a complex into simple vertex modules and length-two arrow modules. Thus B_6 is representation finite, with six simples and six indecomposable length-two projectives; it is not the category of the HLS diamond or of band tubes.

Theorem 21.2 (Kernel incidence resolution)**.** *For a feasible* $\mathbf{r}$*, use subspaces of dimension* $d_i - r_i$ *and define*

$$Z(\mathbf{r}) = \{(D, R) : \dim R_i = d_i - r_i, \quad \operatorname{im} D_{i-1} \subset R_i \subset \ker D_i\}.$$

This is smooth, and its projection is proper and birational onto the closure of the rank-$\mathbf{r}$ *stratum. The maximal feasible vectors, in componentwise order, give the irreducible components of* $\mathcal{X}(\mathbf{d})$*.*

Proof. Over $\prod_i \mathbf{Gr}(d_i - r_i, V_i)$ the allowed maps are precisely

$$D_i \in \operatorname{Hom}(V_i/R_i, R_{i+1}).$$

They form a vector bundle, so its total space is smooth and irreducible. The incidence conditions are closed, and projection from the Grassmannian product is proper. Its image is the locus of

complexes with ranks at most $\mathbf{r}$. To see surjectivity onto that locus, extend the incoming image to a $(d_i - r_i)$-dimensional subspace inside the kernel; feasibility gives both inequalities needed for that extension. Full ranks r_i occur in the bundle, since $r_i \le d_{i+1} - r_{i+1}$. On this open locus $R_i = \ker D_i$ is forced and varies regularly, giving an inverse there. The image is consequently the closure of this stratum. Every complex has a feasible rank vector and these closures are ordered by componentwise rank, proving the last assertion. □

The index in $d_i - r_i$ matters. With all $d_i = 1$ and the maximal vector $(1, 0, 1, 0, 1, 0)$, the alternative $d_i - r_{i-1}$ asks for a one-dimensional R_1 inside the zero kernel of D_1. It has no fiber on the intended dense stratum. This elementary counterexample rules out that incorrect resolution. The theorem above describes the circular variety and its own group, not a nonexistent quotient map of Proposition 21.1.

21.3 Zwara's Hom control and covering Hom decompositions

For finite-dimensional algebras B, A, an exact k-linear functor $F : \operatorname{mod} B \to \operatorname{mod} A$ is *Hom-controlled* in the sense used by Zwara if there is a bilinear form $\beta : K_0(B) \times K_0(B) \to \mathbb{Z}$ with

$$\dim_k \operatorname{Hom}_A(FU, FV) - \dim_k \operatorname{Hom}_B(U, V) = \beta([U], [V]) \tag{21.1}$$

for all U, V. This is the explicit definition in [80, Section 6, p.71], following [79]. A fully faithful exact functor satisfies it with $\beta = 0$, when these finite-dimensional hypotheses apply. The definition does not consist of a sum indexed by deck translates.

The face unfolding fails even this weaker numerical test on a finite radical-square-zero quotient. Take dimension $(1, 1, 0)$, let M_0 have $a_1 = 1, b_3 = 0$, and M_1 have $a_1 = b_3 = 1$, all other maps zero. Their original endomorphism dimensions are both one. After unfolding, M_0 has one edge summand and two isolated simples, so its endomorphism dimension is three; M_1 has two disjoint edge summands, so its endomorphism dimension is two. The defects are two and one at the same K_0 class. No bilinear form in (21.1) can describe them.

Under suitable covering and support hypotheses a Galois push-down has the separate identity

$$\operatorname{Hom}(F_\lambda M, F_\lambda N) \cong \bigoplus_{g \in G} \operatorname{Hom}(M, {}^g N).$$

The support hypotheses ensure the asserted direct sum is the appropriate space; an arbitrary infinite-support category may require other finiteness conditions. Even when this formula holds, the nonidentity summands can depend on the module structures rather than only on their K_0 classes.

Proposition 21.3 (A finite covering counterexample)**.** *A Galois push-down need not be Hom-controlled in Zwara's sense, even for a finite group and finite-dimensional algebras.*

Proof. Let B be the two-cycle $1 \rightleftarrows 2$ with all length-two paths zero, and let $A = k[t]/t^2$. The deck involution exchanges vertices and arrows; the quotient sends either arrow to t. Its push-down adds the vertex spaces. Put $M = S_1 \oplus S_2$ and $N = P_1$. Both have dimension vector $(1, 1)$. The first pushes to k^2 with $t = 0$, the second to a single nilpotent Jordan block of size two. Thus

	$\dim \operatorname{End}_B$	$\dim \operatorname{End}_A(F(-))$	defect
M	2	4	2
N	1	2	1

while $[M] = [N]$ in $K_0(B)$. This contradicts any proposed bilinear defect. The example satisfies the ordinary finite covering Hom decomposition; it is the inference from that decomposition to (21.1) which fails. □

The geometric conclusion must also be stated accurately. The module-scheme result of [79] gives the specified smooth-then-immersion factorization on connected components. The smoothness results for orbit closures and rank schemes in [80, Theorems 6.1–6.2] have their own source, image and singularity-type hypotheses. They do not say that every induced map into the entire module variety is smooth. For example the quotient $k[t]/t^2 \twoheadrightarrow k$ has fully faithful inflation, but the point with $t = 0$ is not a smooth immersion into the positive-dimensional square-zero matrix variety in dimension two.

21.4 Homogeneous tubes in a polynomial dessin algebra

Let $A = k[x, y]/xy$ be the actual polynomial interval dessin algebra. At the smooth support $(x-\lambda, y)$, with $\lambda \neq 0$, put

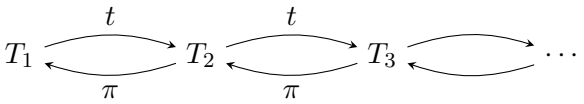
$$T_r = k[t]/t^r, \qquad x = \lambda + t, \quad y = 0, \qquad r \geq 1.$$

All finite-dimensional modules at that support decompose into these modules. This is the elementary divisor theorem for the nilpotent operator t. Modules at different maximal supports have no nonzero maps between them.

Proposition 21.4 (A full tube with explicit meshes)**.** *In* $\mathrm{mod}_{\mathrm{fd}}\, A$ *the* T_r *form a homogeneous tube,* $\tau T_r = T_r$*. With* $T_0 = 0$ *its sequences are*

$$0 \to T_r \xrightarrow{(\pi,t)} T_{r-1} \oplus T_{r+1} \xrightarrow{(t,-\pi)} T_r \to 0.$$

Proof. The two composites cancel. The first map is injective, including when $r = 1$, and the second is surjective; their dimensions then prove exactness. A nonsurjective map into T_r lands in tT_r and factors through the first summand of the right map. A surjective nonretraction from an indecomposable T_s has $s > r$ and factors through T_{r+1}. Decomposing a general source proves right almost splitness; finite-dimensional duality proves the left property. The sequence is nonsplit and its endpoints have local endomorphism rings. Maps from other support blocks are zero, so this is an ambient finite-dimensional-module assertion, not merely an intrinsic quotient claim. □



$$\tau T_r = T_r$$

Figure 21.1: The entire rank-one tube, continued to every quasi-length. The translation fixes quasi-length; the irreducible arrows change it.

The same construction is available on each branch of the polynomial theta algebra at a nonzero central value. After inverting u, its algebra is the product of two matrix algebras over $k[u, u^{-1}]$. Morita equivalence at a fixed nonzero λ gives one such tube on each branch, with dimension vector (r, r, r). Two arrows may be identities and the third $J_r(\lambda)$. These are two *families* indexed by λ, not two tubes in total. They do not extend to continuous finite-dimensional completed-order modules: there the central u must be nilpotent.

There are also mixed inverse bands. Already $a_1 b_3^{-1}$ is a primitive cyclic string. More generally

$$(a_1 a_2 a_3)^j a_1 b_3^{-1}, \qquad j \geq 0,$$

has no forbidden subword in any power and is primitive because its cyclic period contains exactly one inverse letter. Their different lengths give infinitely many band classes. This disproves a classification by just the two pure permitted triangles. Existence of these modules does not by itself identify their entire ambient completed-order AR components.

At the singular central support, a pure-branch tube has only an intrinsic meaning. The augmentation of the y-Jordan block T_2^y onto k cannot factor through $T_2^x \to k$: any map $T_2^y \to T_2^x$ lands in $\operatorname{ann}(x) = \operatorname{soc}(T_2^x)$, which that projection kills. Thus even the first right almost-split factorization would fail in the larger node finite-length category.

21.5 The whole nodal lattice AR category and its stable triangles

Let $\Lambda = k[[x,y]]/xy$, $u = x + y$, $X = (x)$ and $Y = (y)$. Normalize inside a lattice's total quotient module using $B = k[[x]] \times k[[y]]$. If $\overline{M} = BM$, the conductor gives $u\overline{M} \subset M \subset \overline{M}$. The subspace $M/u\overline{M}$ projects onto both branch residue spaces. Separate its pure branch intersections and its common quotient; bases put it into pure lines and diagonal pairs. Lifting those bases over the two DVRs gives

$$M \cong \Lambda^r \oplus X^a \oplus Y^b. \tag{21.2}$$

This is a classification of the finite R-free lattice category. It is not a claim that ordinary polynomial strings classify every completed lattice.

Theorem 21.5 (The full nodal lattice meshes)**.** *The two almost-split sequences in this category are*

$$0 \to Y \hookrightarrow \Lambda \xrightarrow{x} X \to 0, \qquad 0 \to X \hookrightarrow \Lambda \xrightarrow{y} Y \to 0.$$

Its indecomposable lattice vertices are Λ, X, Y, with $\tau X = Y$, $\tau Y = X$.

Proof. The first kernel is $\operatorname{ann} x = Y$. It does not split: a map $X \to \Lambda$ lands in X, and its composite with multiplication by x is a nonunit of $\operatorname{End} X = k[[x]]$. Every map $\Lambda \to X$ factors through the displayed epimorphism, $\operatorname{Hom}(Y, X) = 0$, and each nonunit of $\operatorname{End} X$ is divisible by x, so factors through $X \hookrightarrow \Lambda \to X$. Classification (21.2) now proves right almost splitness for every finite sum: a unit component from X would already be a retraction. For the left property, maps $Y \to X$ are zero, maps $Y \to \Lambda$ extend to multiplication on Λ, and each nonunit endomorphism of Y factors through $Y \hookrightarrow \Lambda \xrightarrow{y} Y$. A unit component into a Y summand would be a section. This proves the factorization for every target in the classification. Endomorphisms of Λ fixing a structure map have constant term one, hence are units; this proves minimality. Interchange x, y for the other sequence. There are no other nonprojective indecomposables in this category. □

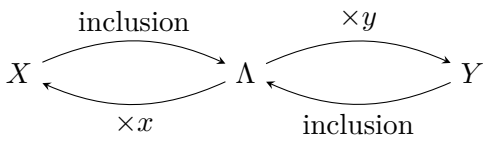


Figure 21.2: The full lattice AR quiver. Its two opposite-branch mesh composites are zero. The central vertex is projective.

Reduction by the regular parameter gives a genuine exact diagram:

$$\begin{array}{ccccccccc}
0 & \longrightarrow & Y & \hookrightarrow & \Lambda & \xrightarrow{x} & X & \longrightarrow & 0 \\
 & & \downarrow{\scriptstyle u} & & \downarrow{\scriptstyle u} & & \downarrow{\scriptstyle u} & & \\
0 & \longrightarrow & Y & \hookrightarrow & \Lambda & \xrightarrow{x} & X & \longrightarrow & 0 \\
 & & \downarrow & & \downarrow & & \downarrow & & \\
0 & \longrightarrow & k & \xrightarrow{-\epsilon} & k[\epsilon]/\epsilon^2 & \longrightarrow & k & \longrightarrow & 0.
\end{array}$$

All three terms are R-free, so the columns are exact and the free quotient X prevents a Tor term in the bottom row. With generators y, x on the two ends, $y = -\epsilon$ explains the sign. This diagram is stronger than a drawing of four mesh diamonds with arbitrary rows and columns.

The lattice category is Frobenius. Alternating annihilator resolutions give

$$\operatorname{Ext}^1(X, \Lambda) = \operatorname{Ext}^1(Y, \Lambda) = 0.$$

The embeddings above give enough injectives, and Λ gives enough projectives. The projective-injectives are its finite free sums. Modulo them,

$$\underline{\operatorname{End}}(X) = \underline{\operatorname{End}}(Y) = k, \qquad \underline{\operatorname{Hom}}(X, Y) = \underline{\operatorname{Hom}}(Y, X) = 0.$$

For example maps through Λ give precisely $xk[[x]]$ inside $\operatorname{End} X = k[[x]]$. The two rows identify $\Sigma X = Y$ and $\Sigma Y = X$. Hence its stable AR triangles are

$$Y \to 0 \to X \xrightarrow{\sim} \Sigma Y, \qquad X \to 0 \to Y \xrightarrow{\sim} \Sigma X.$$

The four ordinary arrows through Λ disappear in the stable category. A tube or a stable three-vertex

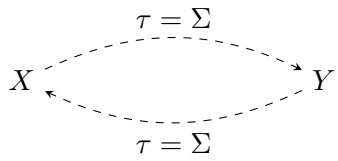


Figure 21.3: Stable node data: two isolated ordinary-quiver vertices. The dashed arrows indicate translation and suspension, not irreducible morphisms.

grid would describe a different category.

21.6 A complete HLS component inside a genuine dessin quotient

For $n \geq 3$ take the infinite closed algebra with chain vertices $1, \ldots, n$, endpoint loops l, r, opposite chain arrows a_i, b_i, and relations

$$l^2 = r^2 = a_i b_i = b_i a_i = 0.$$

Its single permitted cycle goes along the chain and back, with the endpoint loops. It is an actual sphere dessin: the Chebyshev map

$$\beta_n(z) = \frac{1}{1 - T_n(z)^2}$$

has passport $(2n), (2^n), (2^{n-1}, 1, 1)$; the dual contracted graph is a path. Its completed order has rank $(2n)^2$. Kill one arrow at every chain edge, retaining an orientation Ω. The quotient is exactly $H_n(\Omega)$ of [77]. A surviving path has length at most $n+1$, so the killed-arrow ideal contains J^{n+2}. The polynomial algebra and its completion have the same finite quotient. Denote its intrinsic module category by $\mathcal{C}_\Omega$.

For H_3 retain $1 \xrightarrow{a} 2 \xrightarrow{b} 3$. The following table fixes all nine modules, not merely nine unlabelled points. A superscript $-$ denotes a formal inverse, not an inverse in the algebra. Each word is read in travel order; its basis consists of its successive visited positions.

Label	word	start	dimension vector
A	$l^-abrb^-a^-l^-$	1	$(4,2,2)$
B	l^-ab	1	$(2,1,1)$
C	l^-abr^-	1	$(2,1,2)$
D	$abrb^-a^-l^-$	1	$(3,2,2)$
E	ab	1	$(1,1,1)$
F	abr^-	1	$(1,1,2)$
G	$r^-b^-a^-labrb^-a^-l^-$	3	$(4,3,4)$
H	$r^-b^-a^-lab$	3	$(2,2,3)$
I	$r^-b^-a^-labr^-$	3	$(2,2,4)$

A direct letter sends the preceding position vector to the following one; an inverse letter sends the following vector to the preceding one. These rules satisfy $l^2 = r^2 = 0$ and specify every arrow matrix.

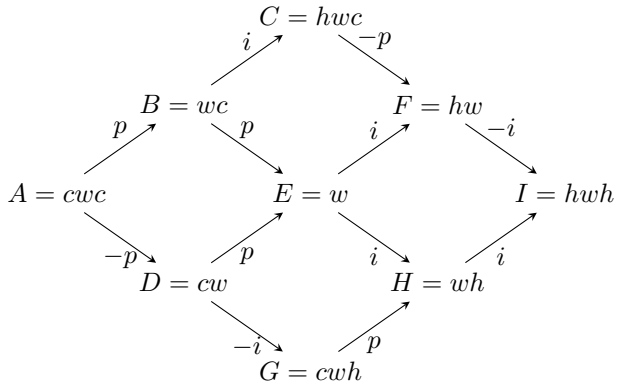


Figure 21.4: The actual signed HLS nine-object window, rotated to show four meshes. Hook/cohook labels are operation labels from the source convention; the preceding table fixes their travel words.

The maps p are projections and i embeddings. In the rectangular version, horizontal projections retain initial positions and horizontal embeddings append the final r^- position. Upper vertical projections delete the initial l^-, and lower vertical embeddings prepend $r^-b^-a^-l$. The signs in Figure 21.4 give

$$0 \to A \to B \oplus D \to E \to 0, \qquad 0 \to B \to C \oplus E \to F \to 0,$$
$$0 \to D \to E \oplus G \to H \to 0, \qquad 0 \to E \to F \oplus H \to I \to 0.$$

Their module intertwinings and exactness follow by these basis rules; for example the two maps out of A have joint kernel zero, the maps to E are onto, and the displayed dimension vectors give the kernel dimension at each vertex. The almost-split property is the additional content of the source's hook/cohook theorem, after identifying $w = ab$ as minimal of type $(2,2)$.

21.6.1 Extending the window to the whole component

Let X_p mean p left hooks when $p \geq 0$ and $-p$ left cohooks when $p < 0$; let Y_q mean q right hooks when $q \geq 0$ and $-q$ right cohooks when $q < 0$. These are operations on strings, not paths. Lemma 3.10 and the four-region induction in [77, Proposition 3.12] give every vertex of this component as

$$M(p,q) = M(X_p w Y_q), \qquad (p,q) \in \mathbb{Z}^2.$$

There are arrows increasing one coordinate at a time and

$$\tau M(p,q) = M(p-1, q-1).$$

The meshes are

$$0 \to M(p-1,q-1) \to M(p,q-1) \oplus M(p-1,q) \to M(p,q) \to 0.$$

The nine labels are $A = M(-1,-1), E = M(0,0), I = M(1,1)$ and the six intermediate positions. This is the whole $\mathbb{Z}A_\infty^\infty$ component, continued in all four directions. Its two outgoing axis rays are monomorphisms and its two incoming axis corays epimorphisms. It is not an exact array of rows and columns.

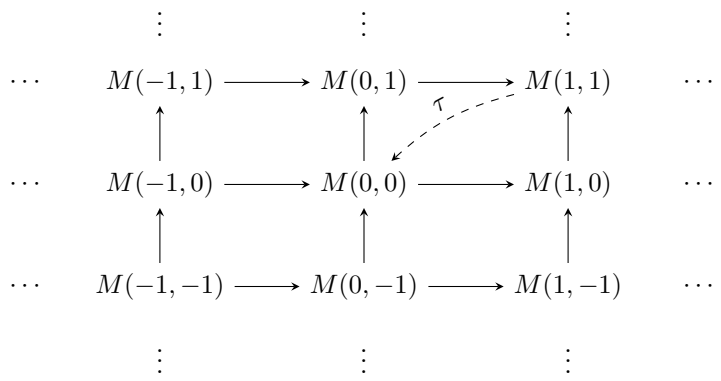


Figure 21.5: The coordinate model of the entire HLS component: continue this lattice for all $(p,q) \in \mathbb{Z}^2$. Ordinary arrows increase one coordinate; translation decreases both. Only mesh diagonals give the exact sequences above. The dashed arrow is translation, not an additional irreducible map.

The complete H_n module AR inventory is also known [77, Theorem 3.14]: one component containing the projective and injective modules; one exceptional tube of rank $n-1$; homogeneous band tubes; and the components above indexed by minimal type $(2,2)$ strings. In H_3 the exceptional tube has rank two. The source's Figures 3–4 show the projective-injective component, not another view of this stable grid. None of these classifications is silently asserted for every module of the completed infinite order.

21.7 Derived AR triangles and an explicit infinite-pd endpoint

A short exact sequence gives a triangle in a bounded derived category, but almost splitness is tested against all objects there, including shifts. For a finite-dimensional algebra an indecomposable endpoint admits a derived AR triangle precisely when it is perfect; its translate is the derived

Nakayama functor followed by $[-1]$. In the finite-dimensional gentle setting, [81, Section 5.1] applies this criterion to finite homotopy strings and bands, and describes the graded endpoint rotations. The finite-dimensional, grading, boundary and local-system hypotheses remain part of that theorem.

For our $E = M(ab)$ the failure of perfectness can be seen directly. In the travel/right convention the cover $P_1 \to E$ has the eight-element basis

$$e_1, l, a, la, ab, lab, abr, labr,$$

and retains e_1, a, ab. Its kernel is $lP_1 \oplus abrP_3$. Multiplication by $l : P_1 \to P_1$ has kernel lP_1, since exactly the paths already starting with l are killed. The terminal simple repeats through $r : P_3 \to P_3$ in the same way. Thus its minimal projective resolution has nonzero terms in arbitrarily high degrees. In particular the module AR sequence ending in E in Figure 21.4 does not become an AR triangle ending in E in $D^b(\operatorname{mod} H_3)$.

By contrast, the smooth-support block of Proposition 21.4 is hereditary. The operator $(x - \lambda)^r$ is a nonzerodivisor of the polynomial node algebra and its cokernel is T_r, so it gives a length-one free resolution. The indecomposables in that block's bounded derived category are $T_r[n]$; each fixed shift has the same homogeneous tube and the AR triangles are induced by its exact sequences. The nodal stable lattice triangles above give another actual triangulated model. These two positive descriptions do not identify all of D^b, Perf and the stable lattice category.

21.8 Ringel's wind-wheel tile: limits are part of the object

Ringel's $\mathcal{T}(45)$ is the tile of the five-bar wind wheel specified in [78, Section 12.5, p.30], drawn in Section 13.3, pp.38–40. It is part of an AR *quilt*, not just a finite-module mesh. Four finite-word quarters surround a nonperiodic two-sided word. The diagonals consist of one-sided limiting words; contracting ones use ordinary strings, expanding ones need completion, and the central mixed word needs partial completion. The five-tile example glues to a torus with five holes. Deleting the infinite-word points recovers its finite string components. This is a category construction, not an assertion that its quilt genus equals the genus of our original Belyi pair.

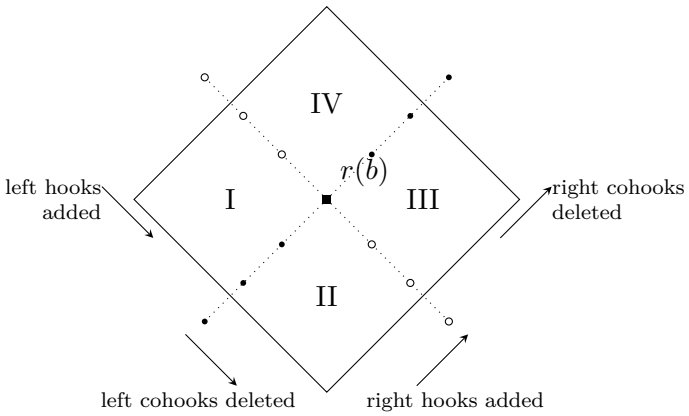


Figure 21.6: A scoped schematic of Ringel's tile, with the directions and limit symbols of the primary picture: circles are contracting one-sided words, bullets expanding ones, and the black square the two-sided limit. The border and neighboring-tile identifications are additional data.

This picture is a finite-dimensional special-biserial wind-wheel category. It is not already our completed node or the H_n quotient. No map identifying its algebra with a present dessin quotient

has been supplied here. Such an import would have to specify that map and the category in which almost-split or pure-injective statements are retained. In particular deleting diagonals from one tile leaves four quarters, not one HLS grid. Ringel also lists indecomposable algebraically compact modules omitted from the quilt; it is not itself a classification of all ordinary modules or all pure-injectives.

21.9 What meshes can and cannot say about semi-invariants

21.9.1 Presentation determinants and their exact scope

Let $P_0 = \bigoplus_i (e_i A)^{a_i}$, $P_1 = \bigoplus_i (e_i A)^{b_i}$ and $f : P_1 \to P_0$. When the evaluated Hom map is square, chosen bases give

$$c_f(N) = \det \mathrm{Hom}_A(f, N), \qquad c_f(gN) = \prod_i \det(g_i)^{b_i - a_i} c_f(N).$$

The action on coordinate functions has inverse weight. Put $M = \operatorname{coker} f$. The kernel is $\mathrm{Hom}(M, N)$. If f is the injective first map of a short projective resolution, its cokernel after Hom is Ext^1 and higher Ext vanish. A general presentation keeps its other cohomology. These distinctions are part of the semi-invariant construction, not consequences of an AR drawing [82].

Proposition 21.6 (A mesh whose determinant is a free generator)**.** *An actual dessin quotient has an almost-split mesh whose presentation semi-invariant generates a polynomial algebra with no relation.*

Proof. In the Chebyshev order above kill vertex3, both loops and the reverse arrow, retaining only $a : 1 \to 2$. This is the genuine continuous quotient $B = k(1 \to 2)$. Its three indecomposables are S_1, S_2 and $E = (k \xrightarrow{1} k)$, by rank decomposition of a linear map. Its sequence

$$0 \to S_2 \xrightarrow{i} E \xrightarrow{p} S_1 \to 0$$

is almost split: a nonretraction to S_1 from these indecomposables factors through p, and dually for i. At dimension $(1, 1, 0)$ the single arrow coordinate z transforms by g_2/g_1. The short projective resolution $P_2 \to P_1 \to S_1$ gives the determinant z. Its semi-invariant algebra is $k[z]$, with no algebraic relation among this one generator, although $pi = 0$ is a mesh relation. □

This quotient AR sequence is not ambient. A module supported at vertex1 with $l = J_2$ has a nonsplit epimorphism to S_1. Every map from it to E has zero vertex1 component, because $E(a) = 1$ while its own arrow a is zero. That epimorphism cannot factor through $E \to S_1$ in the larger category.

Proposition 21.7 (The product theorem with its hypotheses)**.** *Suppose $0 \to M' \to M \to M'' \to 0$ is exact and M', M'' have resolutions of length one by finite projectives. At a dimension vector where both end presentation Hom matrices are individually square, compatible bases give $c_M = c_{M'} c_{M''}$, up to the fixed basis sign.*

Proof. Lift the projective cover of M'' through M. Its horseshoe kernel sits in a split sequence of the two projective kernels, so the middle resolution has differential $\left(\begin{smallmatrix} d' & \xi \\ 0 & d'' \end{smallmatrix}\right)$. Applying contravariant Hom gives a triangular square matrix, with diagonal the two prescribed Hom maps. Its determinant is their product. □

Total squareness alone does not give the two factors. In the A_2 mesh at $(1,1)$, S_1 has balanced Euler value, but P_1 and P_2 do not. The three individual Schofield factors are therefore not all defined there. Nor is the infinite-pd H_3 object $E = ab$ eligible for the short-resolution hypothesis just because it sits in a diamond.

21.9.2 Relative nodal determinants really are multiplicative

The lattice AR sequence of Theorem 21.5 is split after forgetting to R. It therefore has the canonical line comparison

$$\det_R \Lambda \cong \det_R Y \otimes_R \det_R X.$$

For any compatible endomorphisms it gives

$$\det_R(1 - TF_\Lambda) = \det_R(1 - TF_Y) \det_R(1 - TF_X). \tag{21.3}$$

Indeed an R-splitting gives a triangular matrix. For every Λ-linear endomorphism of Λ, multiplication by $f = a(u) + b(u)x$, the basis $(y, 1)$ makes that matrix and its two scalar actions explicit:

$$F_\Lambda = \begin{pmatrix} a & -b \\ 0 & a + ub \end{pmatrix}, \qquad F_Y = a, \quad F_X = a + ub.$$

Both sides of (21.3) are $(1 - Ta)(1 - T(a + ub))$. This is a determinant-functor statement after the specified passage to R-perfect modules, not a nonexistent finite Λ-projective resolution of X, Y. On applying $\mathrm{Hom}_\Lambda(-, Y)$ the map between the first two R-lines is multiplication by u; its cokernel is $\mathrm{Ext}^1(X, Y) = k$. Thus underived Hom is not exact on this AR row.

On the reduced continuous finite-length representation variety of the one-vertex node, the ordinary semi-invariant ring is in fact just k. Scalar matrices act trivially by conjugation, forcing a nonzero character $\det^m$ to have $m = 0$. The nilpotent filtration gives a one-parameter conjugation degeneration to zero arrows, so all invariant regular functions are constant. This coexists with the nontrivial lattice AR sequences and the relative determinant identity; the target data differ.

21.9.3 An inherited relation and a framed Plücker relation

A genuine unframed relation is obtained by retaining both arrows $a : 1 \to 2$, $b : 2 \to 1$ in a Chebyshev quotient and killing all other vertices and arrows. At dimension $(1, 1, 0)$ its coordinate ring is $k[a, b]/(ab)$. The two coordinates have opposite torus weights and are semi-invariants. Their relation $ab = 0$ comes from the inherited zero relation; it was not deduced merely from a mesh picture.

For a richer example use the actual elliptic-square order. Its dessin has $\sigma = (1\,2\,3\,4)$ and $\alpha = (1\,3)(2\,4)$; there is one face, so $1 + 2 - 4 + 1 = 0$ gives genus one. The two quiver vertices are the two clean edges. The four arrows form one permitted alternating four-cycle, with two arrows in each direction; every mixed consecutive choice is forbidden. Put $V_0 = k^2, V_1 = k^2$, set both reverse arrows to zero, and let the two forward arrows have arbitrary matrices X, Y. These modules kill J^2 and therefore extend continuously to the completed order. *Frame the source* V_0: the acting group in the next statement is only $\mathrm{GL}(V_1)$. Write $Z = [x_1, x_2, y_1, y_2]$ and $p_{ij} = \det(z_i, z_j)$.

Theorem 21.8 (A genuine framed generator relation)**.** *All six p_{ij} have target character* det*, and their subalgebra is*

$$k[p_{12}, p_{13}, p_{14}, p_{23}, p_{24}, p_{34}]/(p_{12}p_{34} - p_{13}p_{24} + p_{14}p_{23}).$$

The stated quadratic is the full ideal of relations of these six functions.

Proof. On the dense open $p_{12} \neq 0$, express the other columns in the basis x_1, x_2 and expand determinants. This gives the quadratic; polynomial continuation proves it everywhere. Conversely any nonzero six-tuple satisfying it is realized by four columns. After choosing a nonzero p_{12} by permutation, one may take

$$z_1 = (1,0)^t, \qquad z_2 = (0, p_{12})^t,$$
$$z_3 = (-p_{23}/p_{12}, p_{13})^t, \quad z_4 = (-p_{24}/p_{12}, p_{14})^t.$$

Their last minor has the required value by the equation. The zero tuple is also realized. The quadratic is irreducible: its nondegenerate quadratic form has rank six, whereas a product of linear forms has rank at most two. Its ideal is prime, so this image description proves it is the kernel of the polynomial map on coordinate rings. □

The exact diagram behind decomposability is $0 \to \ker Z \to k^4 \xrightarrow{Z} k^2 \to 0$ on the rank-two locus. Its exterior comparison relates complementary kernel minors. The relation extends to the lower-rank boundary; the diagram is not an almost-split sequence.

The framing is indispensable. With $X = I$ and $Y = \left(\begin{smallmatrix} 0 & 0 \\ 1 & 0 \end{smallmatrix}\right)$, $p_{13} = 1$, but swapping the source basis makes it zero. Thus p_{13} is not a semi-invariant for the full original source-and-target group. The coefficients of

$$\det(X + tY) = p_{12} + t(p_{14} - p_{23}) + t^2 p_{34}$$

are full-group semi-invariants. The framed six-generator relation must not be misidentified as their unframed presentation.

21.10 What this supplies to the physical and arithmetic program

The chapter supplies exact categories, explicit maps, parameter spaces and some genuine determinant identities. These can organize an effective model's states and couplings, but τ is a categorical translation, not by itself a unitary time evolution or a self-adjoint Hamiltonian. A physical realization must supply its Hilbert space, domain, positive metric and dynamics independently. Likewise an arbitrary tube parameter is not a prime norm or a Frobenius eigenvalue.

A valid route from a mesh to polynomial functions requires a universal family, the correct group action, determinant-admissible complexes and rank hypotheses. A route from those functions to an arithmetic L-function additionally requires descent, local factors, conductor and archimedean data, convergence or regularization, and a determinant-preserving positive comparison. The relative nodal identity and framed Plücker model give usable intermediate constructions. They do not discharge those arithmetic or analytic obligations, and no full GRH conclusion is asserted.

Chapter 22

Coverings, equivariant AR theory and graded surfaces

The new covering calculation in Section 7 is useful for computations, precisely because its limitations are explicit. The parent algebra is our infinite elliptic theta algebra, its truncation is a finite quotient, and its cover has vertices indexed by integers. They are three different categories. The nine-dimensional module in Proposition 7.2 and the higher-genus family in Corollary 7.3 show that even this genuine covering does not reach every module or bounded derived object.

22.1 Reading the three new sources together

The lecture [74] supplies the covering language: locally bounded categories, exact push-down/pull-up and graded adjunctions. Its Theorems 4.23 and 4.28 give precoverings on the specified module and bounded derived categories. The AR-quiver covering theorem 4.18 assumes an actual Galois covering between Hom-finite Krull–Schmidt categories; density cannot be dropped. A universal bound-quiver cover also depends on its relations and presentation. It is not automatically a covering of the compact surface.

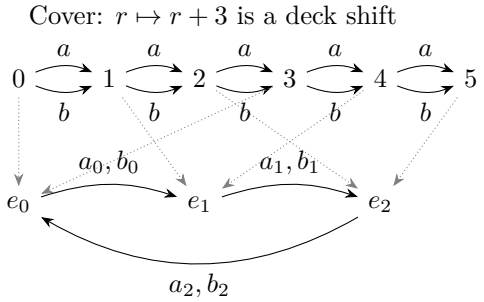
The GLS paper [75] works over an algebraically closed characteristic-zero field. Its finite-dimensional GLS algebras are not automatically our surface orders. Equivariantization transports string/band information to certain non-string algebras. Theorem B includes $\mathbb{Z}D_\infty$, $\mathbb{Z}A_\infty^\infty$ and tube components, with its stated indexing. Theorem C corrects a naive root-only prediction by the additional rank vectors $\beta_i + (2r+1)\delta$ in the named tame affine family. Here vertexwise local freeness and its rank vector concern the GLS vertex algebras; they are not freeness over the central DVR of a surface order.

Schroll's survey [76], particularly Sections 4–5, organizes graded gentle categories through curves, local systems, crossings and twisted complexes, with grading encoded by a line field. Its framework distinguishes perfect objects from finite-total-cohomology objects, and properness from homological smoothness. These distinctions matter for our infinite algebras and for any completed or operator-algebra version.

22.2 A cover that really preserves the relations

At finite L the cover is locally bounded: only lengths below L survive. Without the cutoff, each vertex has infinitely many nonzero outgoing Hom spaces, so the locally bounded hypothesis used above fails. Locally finite dimensional Hom spaces alone do not repair that failure. Conversely the

truncated algebra has extra length relations and is not entitled to all quadratic-gentle surface-model theorems.



Mixed paths and all paths of length L vanish on both rows.

The dotted lines are the residue map, not irreducible module maps. The top row is a finite window of the quiver of the cover, not an AR component. The proof of the Hom decomposition, the unequal Hom-control defects and the full nine-dimensional calculation appear in Sections 7 and 20; they are not inferred from the drawing. The genus-four instance has dimension 27, and the genus-nine instance dimension 57, with the same exact lift obstruction.

22.3 Equivariant modules, not an unproved folding rule

The branch swap preserves J^L. Thus it acts on B_L and defines the skew-group algebra $B_L \# k[\mathbb{Z}/2]$. An equivariant object needs an isomorphism $c : M \to \sigma M$ satisfying the cocycle equation $\sigma(c)c = 1$. A fixed isomorphism class alone does not choose c. For the nine-dimensional witness, the unequal branch ranks rule out even $M \simeq \sigma M$; the induced pair $M \oplus \sigma M$ instead has a canonical swapping action. This is a concrete use of equivariant machinery without asserting a GLS identification or drawing a folded AR component whose almost-split property has not been checked.

22.4 Derived and symplectic data needed for a determinant

Let $\mathrm{Perf}(A)$ be the compact objects of $D(A)$ and write $D_{\mathrm{fd}}(A)$ for finite-dimensional total cohomology. For our infinite zero-graded A, its regular projective is perfect but has infinite k-dimension; hence it does not belong to $D_{\mathrm{fd}}(A)$. This is already a counterexample to silently embedding the former category in the latter. In proper smooth finite-dimensional situations the familiar equivalences have stronger conclusions; these hypotheses cannot be omitted here.

A curve model must retain the dissection, stops or punctures, line field, intersection grading, local systems and dg or A_∞ enhancement. For perfect complexes, the finite deck decomposition of the Hom complex is proved in Section 20. For a matrix of paths representing a differential one must still check $d^2 = 0$; intersections without degrees and composition data do not prove this identity. A marked surface used for the Fukaya model is not automatically the original compact Riemann surface with its analytic metric and all boundary data forgotten.

Completion to Λ requires continuous modules and the chosen topology. The finite quotient modules extend continuously because a power of J annihilates them. They are not central-DVR-free lattices: in the nine-dimensional example the central length-three parameter already acts by zero. The cutoffs A/J^L must not be renamed Λ/u^L.

For a useful arithmetic determinant, these categorical constructions still need a prime correspondence, a perfect or analytically controlled target, a determinant-line comparison, conductor and

archimedean normalization, and positivity for that same arithmetic function. A derived equivalence or an AR translation does not supply a Hilbert-space adjoint, self-adjoint domain or trace-class convergence. These remain the spectral program's independent load-bearing obligations.

The executable calculation is `examples/coverings/theta_cover.py`. The exact matrix obstruction is checked in `lean-analytic/CoveringBand.lean`. The latter verifies selected actual matrices, not every category or equivalence discussed here.

22.5 From a Hom decomposition to a semi-invariant calculation

For a finite-dimensional quotient, a specified projective presentation $P_1 \xrightarrow{d} P_0$ can define a determinant semi-invariant by applying $\mathrm{Hom}(-, V)$, provided the resulting vector spaces have equal dimension and bases or determinant lines have been fixed. The covering Hom decomposition makes this matrix computable by deck components; it does not make it block diagonal. Cross-degree terms matter. Already

$$\det \begin{pmatrix} 1 & a \\ b & 1 \end{pmatrix} = 1 - ab$$

differs from the product of its two diagonal entries. A product over deck blocks is justified only after proving that the relevant differential or endomorphism preserves those blocks. AR mesh relations concern morphism compositions; a presentation of a ring of semi-invariants requires additional generation and relation proofs, as in the earlier determinant chapters. Neither the GLS root lattice nor a surface intersection form provides the missing positive Hilbert inner product or arithmetic factor identity by itself.

Part IV

Quantum models, trained renormalization and physical realizations

Chapter 23

Positive evolution, clocks and holonomy

A surface order specifies algebraic paths, relations and gluing. A quantum model additionally specifies a complex Hilbert space, an adjoint, a state, a Hamiltonian and a measurement rule. Calling a path a clock does not make those choices. Keeping them explicit turns the physical interpretation into a sequence of calculable questions.

This part retains finite clocks, infinite observable algebras, quantum field theory, learned graph networks and molecular hardware. We distinguish a finite identity, a theorem with hypotheses, a reported experiment and a proposed design. A conditional arithmetic spectral criterion belongs to the second category: its missing determinant comparison cannot be supplied by a hardware analogy.

23.1 States and two different logarithms

A finite system has Hilbert space $\mathcal{H} = \mathbb{C}^d$, inner product v^*w, and density matrix $\rho = \rho^* \succeq 0$, $\operatorname{Tr}\rho = 1$. An observable is Hermitian and has expectation $\operatorname{Tr}(\rho O)$. A pure state is vv^* for a unit vector v; its global phase is invisible. A self-adjoint energy operator gives two different evolutions:

$$U(t) = e^{-itH/\hbar}, \qquad T(a) = e^{-aH/\hbar}, \qquad H = H^*. \tag{23.1}$$

Here t and a have time units. The first is unitary, the second positive. The formal substitution $t = -ia$ is Wick rotation. A continuum field theory needs additional reconstruction and convergence hypotheses.

Proposition 23.1 (Unitary and positive logarithms)**.** *For any unitary U and $\Delta t > 0$ there is a bounded self-adjoint H with $U = e^{-iH\Delta t/\hbar}$. For a positive-definite finite matrix T, the matrix $H = -(\hbar/a)\log T$ is Hermitian. These are distinct operations.*

Proof. Choose a bounded real Borel argument ϑ on the circle. Functional calculus gives $H = -(\hbar/\Delta t)\vartheta(U)$. For the second assertion, unitary diagonalization of T gives positive eigenvalues λ_j and $\log T = V \operatorname{diag}(\log \lambda_j) V^*$. □

The Borel argument need not be continuous or lie in the original C^*-algebra. The coordinate z on S^1 has no continuous logarithm because of its winding. A singular nilpotent matrix has no matrix logarithm: every exponential is invertible. One must model such a truncation by a contraction or channel instead. For a unitary clock, different phase branches change energies by integer multiples of $2\pi\hbar/\Delta t$; identical ticks can thus conceal different dynamics between ticks.

For positive T, top-eigenvalue normalization has the precise sign

$$\widetilde{T} = T/\lambda_{\max}, \qquad \widetilde{H} = H + (\hbar/a)\log\lambda_{\max} I \succeq 0.$$

In contrast, $i\log T$ is skew-Hermitian, not this Hamiltonian. If $\rho > 0$ is a density matrix, $K = -\log\rho$ is a dimensionless, nonnegative modular Hamiltonian. Its interpretation as physical energy requires a chosen scale and a state-dynamics comparison.

23.2 Total positivity, wiring and physical positivity

A real matrix is totally nonnegative when all minors are nonnegative, and totally positive when all are positive. These are ordered-basis conditions. Hilbert-space positivity means Hermitian positive semidefiniteness. The matrix $\left(\begin{smallmatrix}1 & 2\\ 3 & 7\end{smallmatrix}\right)$ is totally positive but not Hermitian. Its positive real eigenvalues do not make its logarithm self-adjoint in the given inner product.

The oscillation theorem and planar-network parametrizations remain valuable [84]. In an acyclic directed planar network, the Lindström–Gessel–Viennot cancellation expresses a minor of its path matrix as a sum over disjoint path families. With compatible boundary orderings the surviving signs are positive. Missing families can give zero minors, so positive edge weights do not guarantee strict positivity. Concatenation preserves total nonnegativity, but not symmetry. Positive Grassmannians organize these minors; they do not alone identify a physical scattering amplitude or reconstruct a quantum field theory.

Proposition 23.2 (A finite positive-energy certificate)**.** *A real symmetric positive-definite transfer matrix has a self-adjoint logarithmic Hamiltonian. Normalization by its largest eigenvalue makes the energy nonnegative. Symmetric total positivity is sufficient, not necessary.*

Proof. Use Proposition 23.1 and the normalization formula. Only positivity of the eigenvalues is needed. □

An exactly soluble transfer. For $T = \left(\begin{smallmatrix}2 & 1\\ 1 & 2\end{smallmatrix}\right)$, let $|\pm\rangle = (1, \pm 1)^T/\sqrt{2}$. Its eigenvalues are $3, 1$, hence

$$T/3 = |+\rangle\langle+| + \tfrac{1}{3}|-\rangle\langle-|, \qquad \widetilde{H} = (\hbar\log 3/a)|-\rangle\langle-|.$$

The normalized state $T/4$ instead has eigenvalues $3/4, 1/4$, with modular energies $\log(4/3), \log 4$. Transfer normalization, Gibbs normalization and energy extraction answer different questions.

For a row-stochastic P, detailed balance $\pi_i P_{ij} = \pi_j P_{ji}$ makes $S = D^{1/2} P D^{-1/2}$ symmetric, where $D = \operatorname{diag}\pi$. Its positivity is extra: $P = \left(\begin{smallmatrix}1/4 & 3/4\\ 3/4 & 1/4\end{smallmatrix}\right)$ is reversible but has eigenvalues $1, -1/2$. The lazy transition $(I + P)/2$ has eigenvalues $1, 1/4$ and a positive logarithmic interpretation. Laziness changes the process; it is not a proof of positivity for the original process.

In Euclidean QFT, reflection positivity is a condition on reflected correlations. Together with the other reconstruction axioms it produces a Hilbert space and positive-energy translations [85]. Entrywise nonnegative Boltzmann weights are not that condition. A family of finite positive matrices does not alone prove locality, covariance or a continuum limit.

23.3 Weyl clocks, Fourier phases and twisted successors

On $\mathbb{C}^d$ with $\omega = e^{2\pi i/d}$ set

$$X_d|j\rangle = |j + 1 \bmod d\rangle, \qquad Z_d|j\rangle = \omega^j|j\rangle.$$

Acting on basis vectors proves $Z_dX_d = \omega X_dZ_d$ and $X_d^d = Z_d^d = I$. The unitary Fourier matrix $(F_d)_{jk} = d^{-1/2}\omega^{jk}$ satisfies $F_dX_dF_d^* = Z_d$. For a qutrit the explicit matrices are

$$X_3 = \begin{pmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{pmatrix}, \qquad Z_3 = \operatorname{diag}(1, \omega, \omega^2).$$

Their commutator $Z_3X_3Z_3^*X_3^* = \omega I$ gives an interference calibration. The differently ordered product $Z_3X_3X_3^*Z_3^* = I$ contains no Weyl phase. Discrete X_d, Z_d generate a finite Pauli group, not all of $\mathrm{SU}(d)$. Universality needs suitable continuous controls and entangling gates.

A successor with wrap-around entry $r \in \mathbb{C}$ obeys

$$S_d(r)^d = rI, \qquad S_d(r)^*S_d(r) = \operatorname{diag}(1, \ldots, 1, |r|^2).$$

It is unitary precisely when $|r| = 1$. For a phase $r = e^{i\phi}$ this is an explicit twisted clock. A formal uniformizer in a surface order is not already a complex phase. A representation, specialization and norm must be supplied before interpreting it physically. A weight with $|r| \leq 1$ is a contraction, for which ancilla implementation is available below.

23.4 Infinite clocks: rotations and odometers are distinct limits

On $\ell^2(\mathbb{Z})$ let $U|n\rangle = |n+1\rangle$ and $V|n\rangle = e^{2\pi i\theta n}|n\rangle$. Then $VU = e^{2\pi i\theta}UV$. Under Fourier transformation, U becomes multiplication by e^{it} and V becomes angle translation. They are not two names for the same multiplication operator. The universal algebra on this relation is the rotation algebra A_θ.

Proposition 23.3 (Finite dimension detects a rational phase)**.** *If q-dimensional unitaries satisfy this relation, then $q\theta \in \mathbb{Z}$. For $\theta = p/q$, the pair X_q, Z_q^p gives a finite representation, but not a faithful representation of the entire universal $A_{p/q}$.*

Proof. Take determinants and cancel the nonzero determinants of U, V. At a rational phase U^q, V^q are central; the finite model fixes them to I, whereas other central circle values are possible. □

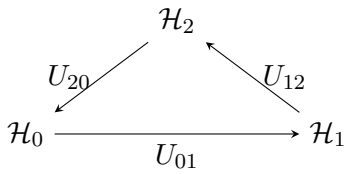
The irrational algebra has a canonical trace selecting the coefficient of U^0V^0 in a Fourier polynomial. This is not the ordinary trace of the identity on an infinite Hilbert space. Abstract algebraic states, normal states in a representation and trace-class densities are distinct notions.

Continued fractions provide rational approximations $|\theta - p_j/q_j| < q_j^{-2}$, but the denominators need not divide one another: successive Fibonacci numbers already refute that assertion. Exact isometries intertwining both Weyl generators at different phases would imply $(e^{2\pi i\theta_j} - e^{2\pi i\theta_{j+1}})J = 0$, which is impossible for a nonzero J. A limiting construction must specify approximate embeddings or a continuous field. Elliott–Evans approximations use matrix algebras over the circle, not a naive AF limit of finite matrix algebras [86].

A separate infinite clock chooses $n_j \mid n_{j+1}$ and forms $X = \varprojlim \mathbb{Z}/n_j\mathbb{Z}$. Addition by one is an odometer with Haar invariant measure. Compatible local extensions $\mathbb{C}((u^{1/n_j}))/\mathbb{C}((u))$ have these finite cyclic groups. Their inverse Galois limit and the direct union of fields point in opposite directions. The full set of denominators gives $\widehat{\mathbb{Z}}$; a selected tower gives the corresponding procyclic quotient. This preserves the ramification interpretation without identifying it with one irrational rotation algebra.

23.5 Holonomy, DG compatibility and persistence

Assign Hilbert fibers to graph vertices and unitary transports to oriented edges, with inverse on reversed edges. Their based loop product is holonomy; changing vertex bases conjugates it. For a triangular loop of complex lines with phases $e^{i\phi_j}$ it is $e^{i(\phi_1+\phi_2+\phi_3)}$. This is a gauge-invariant phase defect, not yet a gravitational redshift.



Operators on different tensor-product registers instead define many-body observables. Their product is not automatically transport on one fiber. The distinction must be fixed before designing a holonomy measurement.

A DG enhancement has $d^2 = 0$ and a graded Leibniz rule. A clock must commute with d to act on cohomology. On $\mathbb{C} \xrightarrow{0} \mathbb{C}$, phases on the two terms can be chosen independently. On $\mathbb{C} \xrightarrow{1} \mathbb{C}$, compatible scalar actions agree but cohomology vanishes. Thus the differential alone creates no physical time. Higher A_∞ operations have their own degrees and signs; periodic syzygies are not themselves a Hamiltonian or a time period.

A state-derived bifiltration can use increasing spatial radius and a decreasing correlation threshold. The order must give actual inclusions. For example, four edges around a square give an H_1 class; adding two triangles fills it. Multiparameter rank invariants, Betti tables and jump loci are meaningful diagnostics, but there is generally no canonical complete two-parameter barcode. Stability needs bounds on the actual filtration functions, including the statistical error of a nonlinear correlation-to-distance conversion. It does not identify a spectral or arithmetic determinant.

Chapter 24

Entanglement and trained quantum renormalization

A renormalization channel needs an objective: which observables, states or phase labels should its coarse variables retain? Training selects parameters from data. This chapter supplies a solvable optimization task and a graph extension, while distinguishing an analytic optimum from an executed training run.

24.1 An entangled pair and the finite first law

Entropy is $S(\rho) = -\operatorname{Tr}(\rho \log \rho)$, with $0 \log 0 = 0$. Reduced states use partial trace; mutual information is $I(A:B) = S(\rho_A) + S(\rho_B) - S(\rho_{AB})$. We use natural logarithms. For

$$|\psi_p\rangle = \sqrt{p}|00\rangle + \sqrt{1-p}|11\rangle, \quad 0 < p < 1,$$

the reductions are $\operatorname{diag}(p, 1-p)$, so $S(\rho_A) = h(p) = -p \log p - (1-p)\log(1-p)$ and $I(A:B) = 2h(p)$. The modular Hamiltonian is $K_A = \operatorname{diag}(-\log p, -\log(1-p))$. At $p = 1/2$ its flow is trivial although the pair is maximally entangled. Entanglement amount therefore does not determine a nontrivial local modular clock rate.

Proposition 24.1 (Finite first law)**.** *If $\sigma > 0$ is a density matrix and $\rho(\epsilon)$ is differentiable, $\rho(0) = \sigma$, then*

$$\frac{d}{d\epsilon} S(\rho(\epsilon))\Big|_0 = \operatorname{Tr}(\rho'(0)(-\log \sigma)).$$

Proof. In a diagonal basis, the derivative under the trace is $d \operatorname{Tr}(X \log X)[Y] = \operatorname{Tr}(Y(\log X + I))$; off-diagonal divided-difference terms have zero trace. Use $\operatorname{Tr} \rho'(0) = 0$. □

For the pair above, $h'(p) = \log((1-p)/p)$ is exactly $\operatorname{Tr}(\operatorname{diag}(1,-1)K_A)$. The reference modular Hamiltonian stays fixed in this variation. A rank-deficient reference needs support control or a declared regularization such as $(1-\eta)\rho + \eta I/d$.

Relative entropy $D(\rho\|\sigma) = \operatorname{Tr} \rho(\log \rho - \log \sigma)$ is nonnegative, with equality only at equal states; incompatible supports give infinity. This is a strict matrix-convexity result, not the assertion that equality in a limit forces equality in an intermediate Hölder inequality. For the Gibbs state $\sigma = e^{-\beta H}/Z$ one obtains $K = \beta H + \log Z$. A general reduced state has no automatically identified physical energy or physical time parameter.

24.2 Correlation geometry and the gravitational research question

On a finite graph assign symmetric regularized lengths

$$\ell_e = \ell_* \max\{\varepsilon, -\log(\max\{I_e, \delta\}/I_*)\}, \qquad \ell_*, \varepsilon, \delta, I_* > 0.$$

They define a shortest-path metric. Zero lengths would yield a pseudometric; infinite lengths could disconnect the graph. If two length assignments differ by at most η on a graph with N vertices, shortest simple paths have at most $N-1$ edges, hence their distances differ by at most $(N-1)\eta$. A patchwise bound can be smaller. The logarithm is ill conditioned near zero correlation, which makes the regularization and its error propagation part of the model.

A regular tree and $\mathbb{H}^2$ are both hyperbolic metric spaces but have different boundaries, Cantor and circle, and are not quasi-isometric. Negative curvature does not alone give logarithmic diameter for every finite quotient. Expansion or a controlled family is needed. Arbitrary edge lengths need not form valid triangles for angle-defect curvature; metric-only curvature surrogates also require their neighborhood measures.

One can postulate a rate ratio $r_v/r_* = e^{-\alpha\varepsilon_v}$ for an entanglement proxy. This is a falsifiable ansatz. A bounded cyclic number expectation cannot increase forever, so an operational rate should use an unwrapped phase, counted transitions or calibrated oscillations. Matched noise-only controls must distinguish frequency changes from lost contrast. The entanglement-equilibrium route to Einstein equations has extra geometric, area, energy and continuum assumptions [87]. The finite first law does not supply them. Persistence can test stability across thresholds; it cannot convert this ansatz into Einstein dynamics.

24.3 Pooling is a channel: an exact trainable example

A channel has Kraus form $\mathcal{E}(\rho) = \sum_m K_m \rho K_m^*$, $\sum_m K_m^* K_m = I$. For unitary evolution followed by discarding one qudit,

$$\mathcal{P}_\theta(\rho) = \mathrm{Tr}_2(U_\theta \rho U_\theta^*), \qquad K_m = (I \otimes \langle m|) U_\theta.$$

The completeness identity follows by summing the environment projectors. An isometry $\mathbb{C}^{d^2} \to \mathbb{C}^d$ would be injective and is impossible for $d > 1$. Stinespring instead adds an environment [88]. Outcome-dependent unitaries replace K_m by $V_m K_m$. Ignoring the outcome remains linear; normalization of a selected branch by its success probability is nonlinear. Partial trace alone is not a nonlinear activation.

A solvable training objective. Encode inputs as $R_y(x)|0\rangle \otimes R_y(y)|0\rangle$, where $R_y(t) = e^{-itY/2}$, and choose

$$U_\theta = e^{-i\theta X \otimes X/2}, \qquad f_\theta(x,y) = \mathrm{Tr}((Z \otimes I) U_\theta \rho_{x,y} U_\theta^*).$$

This encoding varies with its input; $Z^j|0\rangle$ and $e^{ix\hat{n}}|0\rangle$ do not. Anticommutation gives

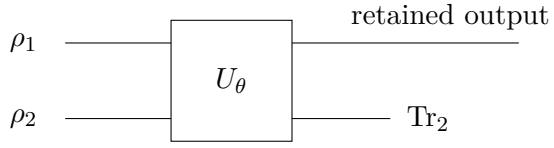
$$U_\theta^*(Z \otimes I)U_\theta = \cos\theta (Z \otimes I) + \sin\theta (Y \otimes X), \qquad f_\theta(x,y) = \cos\theta \cos x.$$

The second term averages to zero for these real product states. Before training, this proves a limitation: that readout cannot retain y. Richer objectives need a changed ansatz or another observable.

For targets $t_j = c\cos x_j$, $-1 < c < 1$, put $A = M^{-1}\sum_j \cos^2 x_j > 0$. The exact squared loss and gradient are

$$L(\theta) = A(\cos\theta - c)^2, \qquad L'(\theta) = -2A\sin\theta(\cos\theta - c).$$

The minima satisfy $\cos\theta = c$ and have zero loss. The stationary points $0, \pi$ must not be mistaken for a trained solution. A numerical optimizer can be tested against this answer; no run is asserted here. For comparison, CNOT followed by target-Z readout gives $f(x, y) = \cos x \cos y$, a true two-input feature.



24.4 Trained QGCN and QGAT on ribbon graphs

Give the graph its cyclic half-edge ordering and a proper edge-coloring. Shared two-site gates can run in parallel on each color class. A simple bounded-degree graph requires at most $\Delta + 1$ colors, not universally Δ. Matching-based pooling specifies disjoint retained/discarded subsystems and a ribbon contraction. Independently applying channels to overlapping marginals does not automatically produce a consistent global coarse state.

Quantum convolutional and graph-network proposals offer concrete ansatz families [89, 90, 91]. A trained RG task may target phase classification, selected correlators, a reconstructed state, or an effective Hamiltonian. Success on one does not prove success on the others. A possible loss combines observable errors and symmetry penalties:

$$\frac{1}{M}\sum_{j,a} w_a (f_{\theta,a}(\rho_j) - t_{j,a})^2 + \lambda \sum_e |\theta_e - \theta_{\mathrm{rot}(e)}|^2.$$

The penalty encourages cyclic weight sharing; exact equivariance requires intertwining by the architecture. Geodesic CNN coordinates and ribbon orderings are related design inputs, not identical constructions. Random quanvolutional circuits provide a useful untrained baseline. They must be compared with trained circuits and a classical model on the same data rather than credited with an automatic quantum advantage.

For a gate $e^{-i\theta P/2}$ with $P^2 = I$, parameter-shift differentiation gives

$$\partial_\theta f = \tfrac{1}{2}(f(\theta + \pi/2) - f(\theta - \pi/2)).$$

Indeed the surrounding expectation has only constant, cosine and sine terms in this parameter. Other generators require a different rule. Finite-shot gradients and exact-state-vector gradients have different noise and optimization costs. Clifford layers can create entangled graph and GHZ states; non-Clifford gates are needed for universality, not for entanglement itself.

QGAT scores can be obtained from edge-register query/key circuits followed by classical softmax. For $d > 2$, Z_d is not Hermitian, so choose real quadratures

$$C_m = (Z_d^m + Z_d^{-m})/2, \qquad S_m = (Z_d^m - Z_d^{-m})/(2i).$$

For example, $s_{vu} = -(\langle C_m\rangle_{q_v} - \langle C_m\rangle_{k_u})^2$ gives real weights $\alpha_{vu} = e^{\beta s_{vu}} / \sum_w e^{\beta s_{vw}}$. The normalization is classical. The mixture $\sum_u \alpha_{vu}\rho_u$ is a state; a coherent sum of arbitrary unknown value states needs a separate physical protocol.

Generic amplitude loading of a D-component classical vector can cost order D input operations; $\log D$ qubits do not remove that cost. Several commuting heads can reuse one measurement basis, but their expectations still require repeated preparation. Quantum circuits are linear in ρ; nonlinear input encoding and classical processing are where these feature maps can become nonlinear in classical data.

24.5 Associative memory with a correct measurement model

A normalized readout uses positive effects E_a summing to I, with $p_a = \mathrm{Tr}(E_a U \rho U^*)$. An incomplete readout has sum at most I and a failure effect. For orthogonal rank-one projectors and pure input,

$$E_a = |\xi_a\rangle\langle\xi_a|, \qquad p_a = |\langle U^*\xi_a, \psi\rangle|^2.$$

This is exact. A common unitary preserves the templates' inner products; it does not independently prescribe arbitrary nonorthogonal templates. Arbitrary projectors are not a POVM: two identical $|0\rangle\langle 0|$ projectors give probability sum two at $|0\rangle$. A general Hermitian observable also need not be positive and cannot be interpreted as one rank-one pattern effect through a generalized square root.

For a two-pattern test, measure the computational basis on $R_y(x)|0\rangle$. Then $p_0 = \cos^2(x/2)$ and $p_1 = \sin^2(x/2)$. A larger-probability rule retrieves zero for $|x| < \pi/2$ modulo 2π. Sampling makes queries near the boundary difficult. This is a concrete retrieval task, not an exponential-capacity theorem.

Classical modern Hopfield networks have capacity and retrieval guarantees under specified separation, normalization and temperature assumptions [92]. Substituting d^{n_q} for their associative dimension does not prove a QGAT memory theorem. A neighborhood with δ keys contains δ explicit candidates per head. Capacity requires a precision criterion, a dataset and a resource model. A D-dimensional system perfectly distinguishes at most D orthogonal states in one use; nonorthogonal packing, classical parameter storage and repeated queries are different resources.

Holographic reduced representations bind classical vectors by circular convolution. For the unnormalized discrete transform, $\widehat{a * b} = \widehat{a}\,\widehat{b}$. Correlation is an exact inverse only under additional nonvanishing/normalization conditions. A qudit's $\langle Z_d^m\rangle = \sum_j \rho_{jj}\omega^{mj}$ determines diagonal populations when all modes are known, but no off-diagonal phases. Additional bases distinguish $|+\rangle$ and $|-\rangle$, which have equal populations.

A channel hierarchy can be drawn as a tensor network with coarse bulk variables and fine boundary variables. Holographic-code models use a specified isometric encoding and reconstruction property [93]. Arbitrary discarding does not have it: a channel sending both $|0\rangle$ and $|1\rangle$ to one state destroys the label. A pooling tree alone proves neither boundary-area scaling of an actual state, error-correcting reconstruction nor AdS/CFT. These remain separate questions that a trained graph model can investigate.

24.6 Evidence for trained quantum renormalization

An executable example must record graph, encoding, task, ansatz, loss, train/validation split, seed, optimizer and stopping rule. It should save parameters, predictions and loss history, compare classical and untrained baselines, and distinguish exact, finite-shot and hardware expectations. A noise study must name the channel and where it is applied. A successful finite toy task does not certify a thermodynamic RG limit, quantum speedup, Einstein dynamics or an arithmetic determinant.

The solvable pooling objective supplies a precise acceptance test for the separate executable examples. A training result is reported only when its checkpoint and verification outputs exist. The mathematical model and its analytic optimum remain meaningful independently of such a run.

Chapter 25

Noise, molecular spins and hyperbolic-code architectures

An architecture needs controls, interactions, preparation, measurement and a noise model. A coupling graph supplies none of these calibrations by itself. Exact small operations provide useful acceptance tests before considering larger processors or molecular assemblies.

25.1 Ancillas, noise and measurement resources

For a contraction A, the map

$$v \longmapsto Av \otimes |0\rangle + (I - A^*A)^{1/2} v \otimes |1\rangle$$

is an isometry: its squared norm is $\|v\|^2$. It extends to a unitary with an environment. Postselecting zero realizes A with probability $\mathrm{Tr}(A\rho A^*)$. Tracing out the ancilla includes both branches and does not produce just $A\rho A^*$ as a trace-preserving map. An operator of norm above one must first be scaled, with its success probability and amplitude normalization accounted for.

Amplitude damping has Kraus matrices

$$K_0 = \begin{pmatrix} 1 & 0 \\ 0 & \sqrt{1-\gamma} \end{pmatrix}, \qquad K_1 = \begin{pmatrix} 0 & \sqrt{\gamma} \\ 0 & 0 \end{pmatrix}, \quad 0 \le \gamma \le 1.$$

Direct multiplication gives $K_0^*K_0 + K_1^*K_1 = I$ and multiplies the excited population by $1-\gamma$. Pure dephasing instead can preserve populations while reducing coherences. Under a Markovian qubit model, $1/T_2 = 1/(2T_1) + 1/T_\phi$; these relaxation times alone do not determine pulse calibration, leakage or gate fidelity.

A Hadamard test estimates $\mathrm{Re}\,\mathrm{Tr}(\rho U)$; a changed ancilla basis estimates the imaginary part. Repeated preparations are required. Local randomized measurements and classical shadows can reuse settings for many observables, with locality-dependent costs [94]. Entropy is nonlinear, so a linear-observable sample bound does not automatically estimate mutual information. Reconstructing small marginals or using dedicated purity/entropy estimators is necessary.

At $N = 10^4$ sites and maximum degree Δ, at most $\Delta N/2$ edge marginals are needed. This defines a proposed workload, not an available processor specification. Routing, state preparation, correlated readout errors and reset must enter a resource estimate. Likewise 10^5 qudits of dimension eight require at least 3×10^5 binary qubits, before ancillas and coding. A generic encoded qudit entangler does not cost merely $\log_2 d$ two-qubit gates. Its actual compilation and unused basis states

must be checked. An infinite rotor needs a cutoff and a bound on leakage/aliasing; a cyclic quotient and a hard truncation have different boundary behavior.

25.2 The peptide-spin evidence and its limits

Lanthanide-binding peptides offer chemically programmable spin organization. Rosaleny and collaborators reported ensemble pulsed X-band EPR in frozen solutions at 4 K, with representative relaxation values [95]:

Complex	T_1	T_2
GdLBT	approximately $10\,\mu$s	approximately $2\,\mu$s
NdLBT	approximately $3\,\mu$s	approximately $1\,\mu$s

Driven nutations motivate a molecular-spin platform, but are not randomized-benchmark gate fidelities or individual single-shot register readout. Proposed multispin organization and transport interfaces must be distinguished from a demonstrated controlled multi-qubit processor. Binding-affinity improvements are not measurements of increased coherence.

PDB 1TJB is a 2.00 Å Tb-bound structure [96]. Its motif has seventeen amino acids, YIDTNNDGWYEGDELLA, with a terminal amide entry in the deposited entity. Eight nearby oxygen donors come from Asp3, Asn5, Asp7, the Trp9 backbone carbonyl, and bidentate Glu11 and Glu14. Deposited residue identifiers and sequence indices need reconciliation in atomic constraints. Replacing the metal or scaffold requires new coordination and magnetic validation; a Tb structural model is not already a calibrated Gd or Nd control model.

An effective spin Hamiltonian is

$$H_{\mathrm{spin}} = H_{\mathrm{CF}} + \mu_B \mathbf{B}\cdot\mathbf{g}\mathbf{J} + \mathbf{I}\cdot\mathbf{A}\mathbf{J} + H_{\mathrm{drive}}.$$

Selecting a qubit requires a named transition, selection rules and leakage bounds. In a positive-field spin-7/2 Zeeman model, the lowest states have $m = -7/2, -5/2$; the central $-1/2 \leftrightarrow 1/2$ transition is not the lowest doublet. Spectroscopic selection can still use it, but must not be justified by a false energy ordering.

For an ideal two-level splitting hf, the thermal polarization is

$$p_0 - p_1 = \tanh\!\left(\frac{hf}{2k_BT}\right).$$

At $f = 9.3$ GHz, $hf/k_B \simeq 0.446$ K; at 4 K the polarization is only about 0.056. Initialization near a pure state needs a separate protocol. A peptide's melting temperature and a spin's operational temperature therefore answer different questions.

25.3 Control calculations: rotations and exchange

Choose the laboratory drive convention

$$H(t) = \frac{\hbar\omega_0}{2}Z + \hbar\Omega\cos(\omega_0 t + \phi)X.$$

The resonant rotating-wave approximation gives $H_{\mathrm{rot}} = \hbar\Omega(\cos\phi X + \sin\phi Y)/2$ with the stated rotating-frame convention. For $\phi = 0$, a pulse of length π/Ω gives $R_x(\pi)$, equal to X up to phase.

A virtual Z changes the microwave reference phase; avoiding a separate physical pulse does not mean infinite fidelity of all subsequent gates. Spectral separation and drive strength control the approximation and leakage.

For dimensionless $\mathbf{S} = \boldsymbol{\sigma}/2$, isotropic exchange is $H_{12} = J\mathbf{S}_1 \cdot \mathbf{S}_2$, with J in energy units. If P swaps qubits, then $P = (I + XX + YY + ZZ)/2$, so

$$e^{-itH_{12}/\hbar} = e^{iJt/(4\hbar)} \left[\cos\left(\frac{Jt}{2\hbar}\right) I - i \sin\left(\frac{Jt}{2\hbar}\right) P\right]. \tag{25.1}$$

Thus $Jt/\hbar = \pi$ gives SWAP up to phase and $\pi/2$ gives $\sqrt{\text{SWAP}}$. The latter sends $|01\rangle$ to a Bell-type superposition proportional to $|01\rangle - i|10\rangle$. It is not iSWAP: an XY Hamiltonian omits the ZZ term and has different evolution. Whether a peptide pair realizes any of these models needs measurement; dipolar anisotropy and g-tensor orientation can be essential.

Echo decay, nutations, addressability, gate fidelity and single-shot measurement are distinct experimental tasks. EPR generally measures an ensemble signal. CISS and spin-dependent transport are possible future interfaces, requiring a calibrated relation to a qubit observable. A discussed SOMO–lanthanide exchange energy is not a measured Ln–Ln exchange across two nanometres. A fitted law $J_0 e^{-\kappa r}$ also needs consistent units: with $J_0 = 25$ cm^{-1}, $\kappa = 1$ Å^{-1} and $r = 20$ Å, it predicts about 5.2×10^{-8} cm^{-1}, not an order-one coupling. Geometry and chemistry must determine an actual model.

25.4 A proposed molecular-design program

RFdiffusion3 supports atomically conditioned protein design and structural benchmarks [97]. Proteína-style backbone generation and BoltzDesign-style inverse-predictor optimization are complementary candidate methods. None guarantees a target spin coherence time, interaction graph or gate. A reproducible calculation pins a release, checkpoint, input schema, atomic identifiers and constraints; generated structures become preserved inputs for later evaluation. Current RFD3 documentation uses structured input specifications and its model-specific design interface, not an assumed legacy RFdiffusion invocation.

A conservative motif workflow fixes a resolved coordination cage, specifies metal/donor atoms and the intended state, generates surrounding scaffolds, assigns sequences with fixed motif positions, and tests structure and metal-aware energetics before synthesis. Hydrogen-bond conditioning is not a lanthanide coordination validator. RMSD and prediction confidence are filters, not physical yields or gate measurements. A small-motif benchmark does not guarantee a pass rate for a seven-motif, 1000-residue assembly.

ThermoGFN-IF can be formulated as a proposed multi-fidelity optimization [98]. SPURS supplies mutation stability predictions, BioEmu-v1 protein backbone ensembles, and UMA learned atomistic energies and forces in evaluated domains [99, 100, 101]. Model domains matter: BioEmu-v1 does not explicitly represent metal ions and their donor-oxygen geometry, so it cannot directly certify an intact Ln cage. All-atom reconstruction and a separately validated metal-aware calculation would be extra stages. Atom-count throughput is not a certificate of spin-vibronic accuracy.

For declared dimensionless scores define $R(x) = \exp(\sum_j w_j s_j(x))$. If positive $\Delta\Delta G$ denotes destabilization, its stability reward contribution should have a negative sign. A GFlowNet aims to sample terminal candidates proportionally to this specified reward. Trajectory balance compares full forward/backward path probabilities to the terminal reward; subtrajectory balance additionally uses weighted subpath constraints. A unique construction path restricts the legal backward policy. These definitions can be implemented and tested without claiming a synthesized or operating processor.

Increasing melting temperature need not monotonically raise a Debye scale, suppress a particular spin-phonon channel or improve T_2. At a fixed transition frequency, Bose occupation depends on that frequency and the temperature, not directly on bulk rigidity. A dephasing power law needs its spectral density, coupling operator and mechanism, including direct or Raman processes when relevant. One nanosecond of dynamics does not alone estimate a converged folded/unfolded partition-function ratio. Proposed tables mapping ΔT_m to numerical coherence, gate fidelity and fault tolerance are therefore hypotheses, not measured outcomes.

25.5 Hyperbolic kagome geometry and the assembly problem

A regular $\{p, q\}$ tiling has p sides per face and q faces at a vertex. At curvature -1, a face has angle $2\pi/q$ and area $\pi(p - 2 - 2p/q)$. In a closed genus-g quotient, $pF = 2E$, $qV = 2E$ and $V - E + F = 2 - 2g$. For $\{7, 3\}$,

$$E = 42(g - 1), \quad V = 28(g - 1), \quad F = 12(g - 1), \quad A_f = \pi/3. \tag{25.2}$$

Then $FA_f = 4\pi(g - 1)$ as Gauss–Bonnet requires. The orientation-preserving $(2, 3, 7)$ triangle group has coarea $\pi/21$; a torsion-free subgroup has index $84(g - 1)$. Its reflection triangle has half this area. The Klein quartic, with usual model $x^3y + y^3z + z^3x = 0$, has genus three and $(V, E, F) = (56, 84, 24)$. It is not a construction of every proposed quotient.

For a simple cubic tiling graph, its line graph places a qubit at each edge and joins pairs that share a vertex. Each qubit has four neighbors; the line graph has E vertices and $2E$ edges. Quotient degeneracies such as loops or repeated incidences need separate counting. Equal hyperbolic edge lengths do not automatically become equal Euclidean separations or physical couplings in an assembly.

A finite graph can be embedded and routed in three-dimensional space, but this does not create a complete smooth isometric $\mathbb{H}^2$ in $\mathbb{R}^3$. A gyroid has varying intrinsic curvature. Schwarz–Pick concerns holomorphic metric contraction, not an SE(3) protein-placement algorithm. A finite assembly needs target Euclidean coordinates, distortion and steric bounds, mechanical constraints and explicit chemical ownership of shared sites. Kabsch alignment is a least-squares rigid fit; it does not remove inconsistent global distances. Two independently generated chunks cannot both contain one supposedly shared physical peptide without an actual chemical construction realizing that sharing.

25.6 CSS codes from boundary maps

Let a finite connected closed orientable cellulation have

$$C_2(\Sigma; \mathbb{F}_2) \xrightarrow{\partial_2} C_1(\Sigma; \mathbb{F}_2) \xrightarrow{\partial_1} C_0(\Sigma; \mathbb{F}_2), \qquad \partial_1\partial_2 = 0.$$

This is a chain complex, not a short exact sequence. Its first homology is precisely the code's logical information. Put qubits on edges and set $H_X = \partial_1$, $H_Z = \partial_2^T$. A binary row x defines an X product and a row z a Z product. Their commutation sign is $(-1)^{x\cdot z}$, so the boundary identity proves commutation. The code is the common $+1$ eigenspace in $(\mathbb{C}^2)^{\otimes E}$, not a binary kernel intersection regarded as that Hilbert space.

Proposition 25.1 (Surface-code dimension and distance)**.** *For this genus-g cellulation,*

$$k = E - \operatorname{rank} H_X - \operatorname{rank} H_Z = 2g.$$

Distance is the smaller weight of a nontrivial primal homology class and a nontrivial dual homology class.

Proof. Connectedness gives $\operatorname{rank}\partial_1 = V - 1$, and the one-dimensional top homology gives $\operatorname{rank}\partial_2 = F - 1$. Thus $k = E - V - F + 2 = 2g$. A commuting Z error is a cycle; multiplication by Z stabilizers changes it by a boundary. A nontrivial logical Z is therefore a nonzero homology class. The dual argument treats X errors. □

A noncontractible separating curve can be null homologous. Searching one geometric systole therefore does not certify a code's distance. In a $\{7, 3\}$ code, vertex X checks have weight three and face Z checks weight seven; a global swap of X, Z conventions also swaps the logical descriptions.

At genus thirteen, (25.2) gives 504 data qubits and 26 logical qubits. If $E = 1008$, the genus is 25 and $k = 50$. The asymptotic rate is $1/21$. A $\{5, 4\}$ cellulation gives $E = 20(g - 1)$ and rate $1/10$. Syndrome ancillas add physical resources. Suitable hyperbolic quotient families have logarithmic distance at constant rate [102, 103]; topology alone gives no such lower bound for every quotient, or distance eight for an unspecified finite surface. A finite certificate needs its incidence matrices and primal/dual logical-operator search. Logical qubit count grows linearly, not exponentially, with the number of physical qubits.

A finite acceptance test. An $L \times L$ periodic square cellulation, $L \geq 3$, has $V = F = L^2$, $E = 2L^2$. Thus $k = 2$ and $d = L$: a straight winding row or column gives an L-edge logical cycle in either lattice, and a lifted winding cycle must traverse at least L edges. This yields $[[2L^2, 2, L]]$ with weight-four checks. The ranks $L^2 - 1, L^2 - 1$ verify k directly. This is a useful exact decoder test before claiming a larger unprovided hyperbolic instance.

A storage code is not yet a fault-tolerant computer. Generic exchange and Zeeman terms do not commute with every stabilizer. An XX interaction meeting a Z check on exactly one site anticommutes with it, giving a commutator of order the coupling. Syndrome extraction needs ancillas, controlled operations, reset and measured error models; three exchange pulses alone do not prove nondemolition parity measurement. A threshold is conditional on code family, decoder and noise model. Code-capacity noise with ideal checks and circuit-level noisy checks are different. A gate estimate from another model cannot be compared to one threshold and declared fault tolerance.

25.7 The arithmetic comparison is still an additional task

A physical graph, a complex Belyi dessin and a rational arithmetic model are different data. A passport need not determine a Galois orbit, and one finite dessin cannot support a faithful action of the entire absolute Galois group on its finite sheet set. Frobenius requires descent and good reduction. Measured couplings do not automatically recover an Euler factor. The earlier finite-field theta benchmark is a concrete comparison test with its already specified scope, not a global GRH proof.

There is a useful circuit sanity check. XY exchange annihilates $|00\rangle$, so its exponential fixes it. A network of these gates, on an all-zero register preceded only by diagonal Z rotations, has a vacuum overlap that is just a phase. It cannot generate L-function zeros. Also e^{-izG} for Hermitian G is not generally unitary at complex z. Finite products $\prod_p (1 - \chi(p)p^{-s})$ are exponential polynomials in s, not degree-N polynomials; their unramified factor zeros have $\operatorname{Re} s = 0$, not the critical line.

A Hilbert–Pólya proposal still has to identify its determinant with the same completed arithmetic function, including gamma and conductor factors. The transformation $\rho \mapsto 1 - \overline{\rho}$ fixes a zero on the critical line; it does not reverse the ordinate. Conjugate twists matter for complex characters. Chiral

symmetry is extra structure. An unbounded $\pm$-paired self-adjoint spectrum has no Gibbs trace $e^{-\beta H}$; $e^{-\beta H^2}$ or $e^{-\beta|H|}$ is a different diagnostic. Similarly $\sqrt{\Delta - 1/4}$ needs a spectral restriction away from negative shifted modes and a treatment of continuous spectrum.

These chapters give finite calculations, trainable channels and concrete experiments. They preserve the physical research program while leaving its arithmetic determinant, purity, convergence and positive-realization obligations explicit. They establish no complete proof of GRH.

Part V

Computational and learned constructions

Chapter 26

Higher-genus dessins and reproducible finite carriers

The companion separates three objects: an exact combinatorial dessin, a chosen visualization of its surface, and a finite operator model. Agreement of their dimensions is not an equivalence of categories. We give constructions for which these distinctions can be checked directly.

26.1 A hyperelliptic family beyond the torus

For odd $m \geq 3$, let C_m be the smooth compactification of

$$y^2 = x^m - 1, \qquad \beta = x^m : C_m \longrightarrow \mathbb{P}^1.$$

The affine curve is smooth: simultaneous vanishing of $2y$ and mx^{m-1} would force $x = y = 0$, which does not lie on the curve in characteristic zero. Its double cover of the x-line branches over the m roots of unity and infinity. Riemann–Hurwitz gives $g = (m-1)/2$.

Proposition 26.1 (Clean passports and infinite orders)**.** *The degree-$2m$ Belyi map β has passport*

$$(m, m), \quad (2, \ldots, 2)\ (m \text{ entries}), \quad (2m).$$

Its contracted dessin has two vertices, m edges and one face. With darts $0, \ldots, 2m-1$, one may take

$$\sigma = (0\ 1\ \cdots\ m-1)(m\ m+1\ \cdots\ 2m-1), \quad \alpha = \prod_{j=0}^{m-1} (j\ j+m).$$

The associated completed order has $k[[u]]$-rank $2m^2$, and the polynomial locally gentle algebra is infinite dimensional.

Proof. Over zero, x is a local coordinate at the two points $y = \pm i$, and x^m has index m. At a root ξ of unity, y is a local coordinate and $\beta - 1 = y^2$, giving index two. At the unique point at infinity, x has pole order two, so β has pole order $2m$. Away from these points, differentiating x^m introduces no further branching. This proves the passport and the Belyi property. The two cycles describe the two lifts of the radial segments in the x-plane. The involution exchanges their lifts at the clean branch points. For odd m, $\phi = \alpha\sigma^{-1}$ is a single $2m$-cycle: its square advances two steps in each m-cycle, and two is invertible modulo m. Hence $2 - 2g = 2 - m + 1$. Each permitted m-cycle repeats without meeting a forbidden pair. This proves infinite dimension; the earlier split-order theorem gives rank $m^2 + m^2$. □

For $m = 5$ this gives genus two and order rank fifty; for $m = 7$ genus three and rank ninety-eight. These are actual cellular dessins, not finite-dimensional stand-ins. The all-degree resolutions are obtained by iterating their forbidden successor permutation, as in the general resolution theorem. The number of vertices of the algebra is m, while its coefficient-order rank is $2m^2$; they should not be interchanged in a quantum register count.

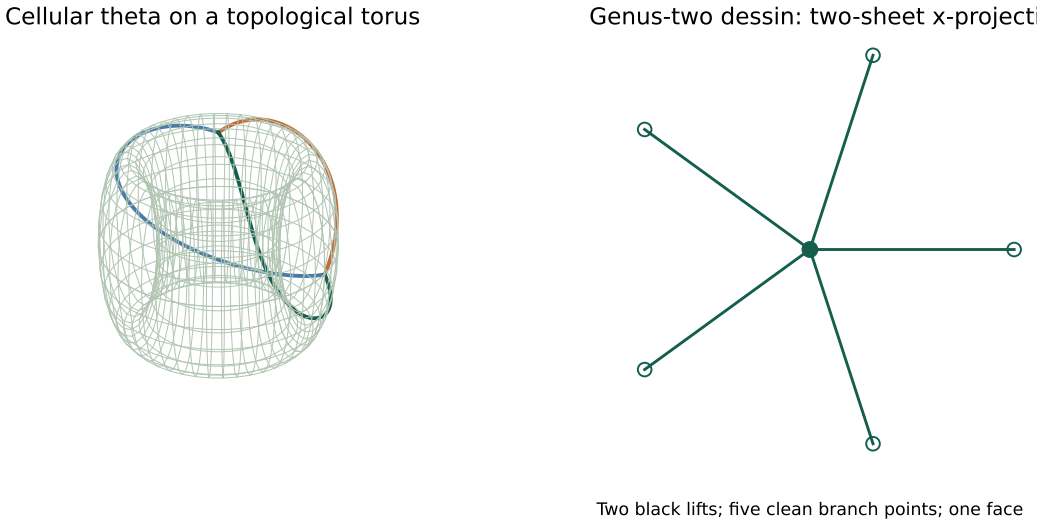


Figure 26.1: Two distinct kinds of faithful data display. Left: a topological torus embedding with the three periodic lifts of the theta graph. Right: the genus-two dessin's x-projection; every radial segment has two lifts. The right-hand disk is not the genus-two surface itself.

26.2 Larger genera and infinite hyperbolic windows

The same exact family gives genus five ($m = 11$, rank242), genus nine ($m = 19$, rank722) and genus sixteen ($m = 33$, rank2178). These ranks refer to the infinite completed orders over $k[[u]]$, not finite vector-space dimensions. The browser includes these examples and a four-dimensional rotation control: write $(x, y) \in \mathbb{C}^2$ in real coordinates (x_1, x_2, y_1, y_2), rotate the (x_1, y_1) plane, then project to three coordinates. The omitted coordinate and projected intersections are explicitly identified. This is a view of the affine double-cover sheets, not a claimed injective embedding of a compact higher-genus surface.

A separate example is the infinite regular hyperbolic $\{7, 3\}$ tiling. Its heptagons have angles $2\pi/3$, so the defect formula gives area $5\pi - 7(2\pi/3) = \pi/3$ at curvature -1. For a regular $\{p, q\}$ tiling, its central polygon's Poincaré circumradius is

$$r = \sqrt{\frac{\cos(\pi/p + \pi/q)}{\cos(\pi/p - \pi/q)}}.$$

To derive this, bisect the polygon into right triangles with acute angles π/p and π/q. The hyperbolic right-triangle identity is $\cosh R = \cot(\pi/p)\cot(\pi/q)$, while the disk radius is $r = \tanh(R/2)$. Substituting and applying elementary sum-to-product identities gives the formula. The condition $1/p + 1/q < 1/2$ ensures positive real radius.

A side is an arc of a circle with center c and squared radius $|c|^2 - 1$, orthogonal to the unit disk. Its endpoints a, b determine $\mathrm{Re}(a\bar{c}) = (|a|^2 + 1)/2$ and the analogous equation for b. Circle inversion gives its hyperbolic reflection

$$z \longmapsto c + \frac{|c|^2 - 1}{\bar{z} - \bar{c}}.$$

If a side is a diameter, the circle formula is replaced by reflection in the line through the origin; the implementation handles that case separately. Reflecting adjacent polygons constructs finite windows of the infinite tiling. The implementation caps both depth and tile count and reports those bounds. It never calls a truncated window the whole infinite lattice. An optional Lorentz-hyperboloid view maps z to $(2\,\mathrm{Re}\, z, 2\,\mathrm{Im}\, z, 1 + |z|^2)/(1 - |z|^2)$, satisfying $T^2 - X^2 - Y^2 = 1$. Its metric is Lorentzian ambient data, not the Euclidean metric of the rendered mesh. A finite compact surface quotient, its side pairings and an arithmetic correspondence would require separate certificates. The current lattice patch supplies none of them.

26.3 What a three-dimensional view can claim

The genus-one examples use a standard torus parameterization

$$((R + r\cos v)\cos u, (R + r\cos v)\sin u, r\sin v), \quad R > r > 0.$$

The theta paths lift from $(0,0)$ to $(1/3, 1/3)$ with translation vectors $(0,0), (-1,0), (0,-1)$ in the universal covering plane. Their interiors are disjoint on the quotient, and their complement is a disk; the induced rotations at the two black vertices give the one-face theta dessin. The elliptic square uses the two fundamental loops. These are topological realizations; the induced Euclidean torus metric is not asserted to be the constant-curvature or uniformizing metric of the algebraic curve.

For $m \geq 5$, the companion instead plots a three-real-coordinate projection of the two sheets of $y^2 = x^m - 1$. Branch cuts, changes of sheet and the compact point at infinity remain part of the mathematical construction. Projected sheets can meet. No injective embedding or numerical uniformization follows from a screen image. The Sage source gives exact passports, quivers, forbidden pairs and determinant calculations. The installed `sage` launcher on this machine points to a missing application, so those Sage scripts are supplied but have not been executed in Sage in this milestone. Exact stdlib Python checks independently verify the displayed permutations, genus and ramification counts.

26.4 Finite clocks and cross-backend agreement

For $d \geq 2$, let $Ue_j = \zeta^j e_j$, $Ve_j = e_{j+1}$, with $\zeta = e^{2\pi i/d}$. Then

$$UV = \zeta VU, \quad U^d = V^d = I.$$

Applying the two sides to e_j proves the first identity, including the wraparound because $\zeta^d = 1$. The discrete Fourier matrix prepares the uniform superposition. The executable comparison applies F, then U, then V, on a four-level carrier; its relative phases are nonconstant. Cirq, Qiskit and PennyLane statevectors are compared to the reference vector. Strawberry Fields realizes the same $U(4)$ matrices on the single-photon subspace of four optical modes and compares the entire density block, including its off-diagonal phases. Preparing a Fock state need not leave the simulator's internal representation marked pure, so reading only a ket would be incorrect.

The photonic cutoff two is exact for this passive, single-photon computation: photon number is preserved and no mode ever contains more than one photon. This justification does not extend to arbitrary squeezing or multiphoton simulations. Separate NVIDIA cuQuantum and CUDA-Q adapters and pinned upstream forks are provided. GPU execution is not performed on this Mac; source availability is distinct from a successful accelerator run.

Chapter 27

Trained quantum coarse-graining architectures

These experiments implement small, nontrivial versions of the original graph-quantum renormalization proposal. They use specified finite physical models, exact labels where available, held-out samples and recorded seeds. A learned approximation is tested against its target; it is not promoted to a new exact renormalization theorem. The classical optimizer, feature readout and data generation are explicit parts of each hybrid architecture.

27.1 Exact teacher data from Ising decimation

For spins $s_1, s_2, s_3 \in \{-1, 1\}$, summing the middle spin gives

$$\sum_{s_2} e^{K_1 s_1 s_2 + K_2 s_2 s_3} = 2\cosh(K_1 s_1 + K_2 s_3) = A e^{K' s_1 s_3},$$

where

$$K' = \tfrac{1}{2}\log\frac{\cosh(K_1+K_2)}{\cosh(K_1-K_2)} = \operatorname{arctanh}(\tanh K_1 \tanh K_2), \quad A = 2\sqrt{\cosh(K_1+K_2)\cosh(K_1-K_2)}.$$

Checking $s_1 s_3 = \pm 1$ proves the identities. Thus the normalization and coupling have separate roles. The network below learns K', not the full partition-function normalization A.

For a three-valent star, let $s = (s_1, s_2, s_3)$ be boundary spins and sum its center. The exact boundary logarithmic weight is

$$f_K(s) = \log\bigl(2\cosh(K_1 s_1 + K_2 s_2 + K_3 s_3)\bigr).$$

Fourier expansion on the Boolean cube is orthogonal for $2^{-3}\sum_s$. The symmetry $f_K(-s) = f_K(s)$ removes every odd monomial. Consequently

$$f_K(s) = c(K) + \sum_{i<j} J_{ij}(K) s_i s_j, \quad J_{ij} = 2^{-3}\sum_s f_K(s) s_i s_j.$$

This proves the zero-field star-to-triangle reduction and supplies exact labels. Such trivalent patches occur at the black vertices of the theta dessin. The associated Ising Hamiltonian is a finite model on that graph, not a faithful representation of its completed surface order.

27.2 Shared quantum graph convolution

The input patch (K_1, K_2) is encoded by $R_y(K_1) \otimes R_y(K_2)|00\rangle$. Two layers of local R_y, R_z rotations, a CZ entangler after each layer, and a final CNOT define an eight-parameter circuit Q_θ. The three features are

$$z_\theta(K) = \big(\langle Z_0\rangle, \langle Z_1\rangle, \langle Z_0 Z_1\rangle\big).$$

A shared affine head predicts $\widehat{K'} = w \cdot z_\theta + b$. Twelve trainable parameters include the four readout coefficients. Independent samples are drawn uniformly from $[0.05, 1.7]^2$: ninety are used for training and thirty are held out. The loss is the mean-square error against exact decimation. The independent constant predictor is fitted only on training targets.

For a gate $e^{-itP/2}$, $P^2 = I$, and a fixed state and observable, the expectation has the form $a + b\cos t + c\sin t$. Hence

$$\partial_t z = \tfrac{1}{2}\big(z(t + \pi/2) - z(t - \pi/2)\big).$$

Each gate parameter occurs once. The code differentiates the expectation features with this formula and then applies the classical chain rule to the squared loss and affine head. Applying parameter shift to the squared loss itself would generally give a wrong gradient. A central finite-difference check of the implemented full loss gradient is recorded.

27.3 QCNN compression with a genuine pooling channel

The four-site periodic transverse-field Ising Hamiltonian is

$$H(h) = -\sum_{(i,j)\in C_4} Z_i Z_j - h\sum_{i=0}^{3} X_i.$$

We diagonalize its 16×16 Hermitian matrix at thirty-six fields in $[0.15, 2.4]$, using the normalized lowest eigenvectors as input states. Every fourth field, starting at index zero, is held out. The circuit has two shared-layout local rotation layers with CZ edges on C_4, followed by CNOTs $(0, 2)$ and $(1, 3)$. Its sixteen parameters are optimized to minimize

$$\mathcal{L}(\theta) = 1 - \frac{1}{N_{\text{train}}}\sum_j \langle\psi_j|Q_\theta^*(I_{01} \otimes |00\rangle\langle 00|_{23})Q_\theta|\psi_j\rangle.$$

The output on retained qubits is

$$\mathcal{E}_\theta(\rho) = \mathrm{Tr}_{23}(Q_\theta \rho Q_\theta^*).$$

Its Kraus maps are $K_b = (I \otimes \langle b|)Q_\theta$, so $\sum_b K_b^* K_b = I$. Thus it is completely positive and trace preserving. For $\rho \geq 0$ and any v, each summand $v^* K_b \rho K_b^* v$ is nonnegative; cyclicity of the finite trace proves trace preservation. The Lean interface verifies the finite-matrix positivity and trace statements with explicit completeness hypotheses. It does not formalize the optimizer or claim convergence of an infinite renormalization flow.

The compression loss is the weight discarded by the chosen vacuum projection. Its successful decrease measures improved concentration for this finite input family. It does not imply exact recovery of arbitrary four-qubit states, an isometry from dimension sixteen to four, or a critical fixed-point construction. The held-out loss remains appreciable, and is reported without calling the model lossless.

27.4 A larger genus-four nine-qubit example

The $m = 9$ hyperelliptic dessin has genus four and completed-order rank162. Its nine edge labels index a permitted quiver cycle. A separate finite transverse-field Ising model is placed on this nine-cycle, with Hilbert-space dimension512. This uses the dessin's combinatorics, without claiming a faithful representation of its infinite order.

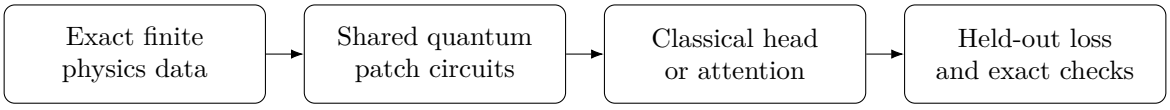
The larger experiment uses twelve ground states over fields $h \in [0.2, 1.9]$, with nine training and three held-out states. These are numerical lowest-eigenstate approximations. The residual alone is not a rigorous lowest-energy certificate. A sparse Hermitian eigensolver has maximum recorded eigenpair residual 1.32×10^{-13}. The circuit has two layers of local R_y, R_z gates and nearest-neighbor CZ around the cycle, with thirty-six parameters. Four CNOTs connect retained to discarded qubits, and tracing out four leaves a five-qubit output. Actual PennyLane gate simulation and parameter-shift gradients train for fifty updates. Final training and held-out discard losses are approximately 0.548 and 0.409. These substantial residuals are reported: this architecture is not an accurate lossless coarse graining. It is a reproducible larger testbed, with saved parameters, held-out examples and numerical ground-state residuals. The local page permits comparison with the smaller four-qubit model, while retaining their different data families and capacities.

27.5 Quantum-feature graph attention

Three two-qubit patch encoders use the same eight circuit parameters on the three pairs of the Ising star. Let their features be $z_e \in \mathbb{R}^3$. A learned query q defines

$$\alpha_e = \frac{e^{q \cdot z_e}}{\sum_f e^{q \cdot z_f}}, \quad c = \sum_e \alpha_e z_e, \qquad \widehat{J}_e = w_{\text{local}} \cdot z_e + w_{\text{context}} \cdot c + b.$$

There are eighteen trainable parameters. One hundred star couplings are sampled from $[0.1, 1.5]^3$, with seventy-five training and twenty-five held-out examples. The objective compares the three predictions with the exact Boolean-Fourier couplings above. The nonlinear attention head is differentiated by central differences of the complete training objective. This experiment is quantum-feature attention with a classical softmax head; it is not a fully coherent attention algorithm. The scalar labels do not assert permutation equivariance under conventions the implementation has not checked.



27.6 Photonic Gaussian pooling

The Strawberry Fields experiment prepares two modes with opposite real squeezing r and $-r$, then applies a fixed unknown-to-the-optimizer beam splitter angle $\alpha = 0.47$. A learned shared beam splitter θ is optimized on twelve of sixteen squeezing values in $[0.1, 0.9]$, holding out every fourth value. With $\hbar = 2$, the covariance convention is (x_0, x_1, p_0, p_1). The objective is

$$\operatorname{Var}(x_0) + \operatorname{Var}(p_1) = 2\left[e^{-2r}\cos^2(\alpha + \theta) + e^{2r}\sin^2(\alpha + \theta)\right].$$

This follows by applying the real beam-splitter rotation to the diagonal covariance of the two squeezed modes. Since $r > 0$, its minimum occurs at $\theta = -\alpha$ modulo π, with value $2e^{-2r}$. The code

trains the angle with central-difference gradients and compares its result with this known optimum. Pooling subsequently discards a mode; the experiment itself optimizes principal axes before that discard. It is a Gaussian finite-mode surrogate, not a demonstration of an interacting field-theory RG flow.

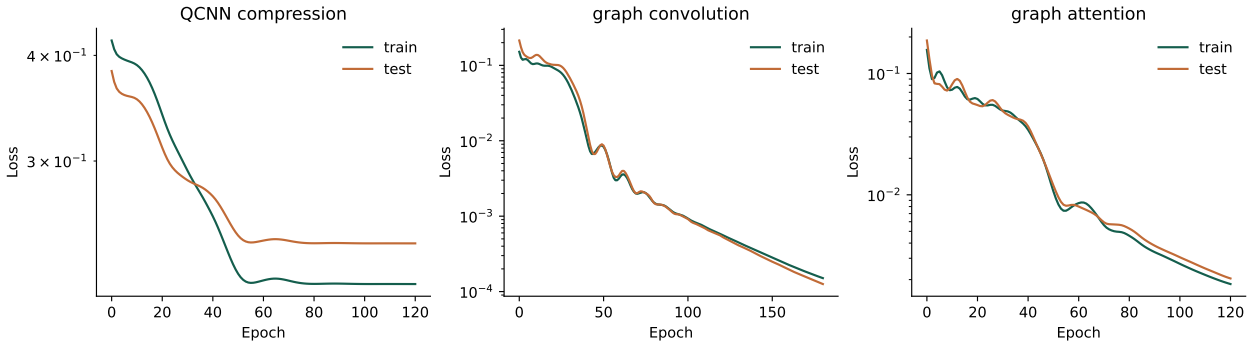


Figure 27.1: Actual PennyLane CPU training runs, seed 20261103. Green is training loss; orange is held-out loss. Architectures have different objectives, so their vertical scales are not cross-model performance rankings.

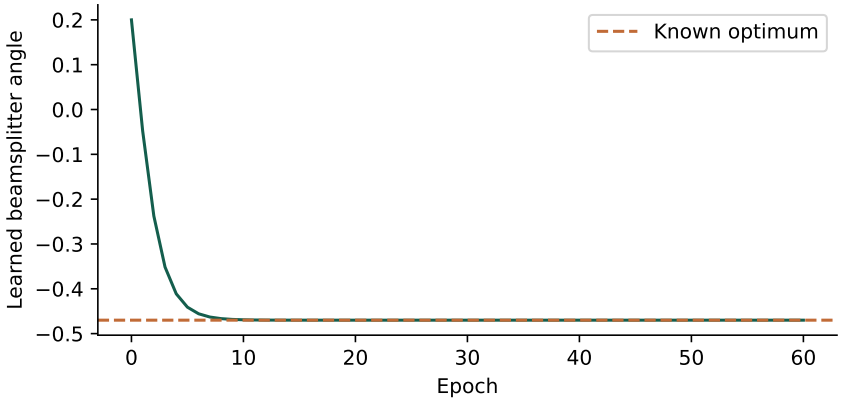


Figure 27.2: Actual Strawberry Fields Gaussian simulation: the learned principal-axis angle approaches its independently known optimum.

Model	Initial held-out loss	Final held-out loss
Four-qubit QCNN compression	0.38332	0.239725
Shared quantum graph convolution	0.213652	0.000125897
Quantum-feature graph attention	0.187065	0.00204477

27.7 Reproduction and limits of the evidence

The scripts in `examples/quantum` specify each circuit, data generator, random seed and optimizer. The approved isolated environment is recorded in `requirements-cpu.lock`; existing environments and all pinned Lean toolchains remain unchanged. NumPy reference and actual PennyLane runs are saved separately. Cross-framework tests exercise a nontrivial Fourier/clock/shift state, not just

a vacuum on which every diagonal clock is invisible. The Strawberry Fields run uses its Gaussian backend for training and its Fock backend for the separate one-photon comparison.

The companion's checkpoint slider selects a saved loss checkpoint. Its held-out examples use final weights; no intermediate predictions are implied. There are no hardware shots, cloud jobs or accelerator benchmarks. The training/test splits are deterministic interpolation tests on small synthetic families, not out-of-distribution or broad physical-validity tests. A model can help rank a proposed construction; it cannot discharge arithmetic holomorphy, determinant convergence, positivity or zero-location obligations merely by achieving low loss.

Chapter 28

Reading an architecture from its tensors

The trained models of Chapter 26 and the ToricGT references have a common computational interface: a finite directed acyclic graph whose vertices contain actual tensors and whose edges record the inputs to each operation. The local inspector at `http://127.0.0.1:8765/architecture.html` makes this graph navigable. Its scene is a computation graph, not a hardware layout or a physical embedding of the infinite surface order. The source code and numerical records are part of the project.

28.1 Three distinct kinds of evidence

The twenty-two saved cases cover six trained CPU model families and six fixed-weight reference variants. The QCNN4 and QCNN9 weights, shared two-qubit convolution, quantum-feature attention, Gaussian photonic pooling and finite Harper hybrid are replayed without further optimizer steps. The classical TokenGT-style, tropical, block-ring, attention-hybrid and two expert-routing variants use explicit seeded weights. They are executable miniatures, not upstream trained language-model checkpoints. The adjustable forward clock is a separately computed finite carrier. GPU and hardware execution remain unavailable, and optional upstream heads remain proposals unless their own numerical evidence is supplied.

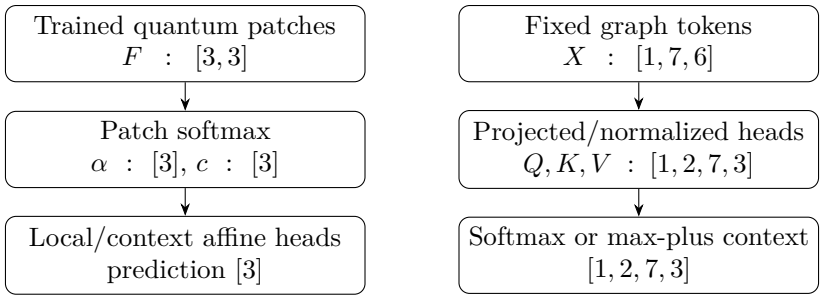
The record contains the source/checkpoint hashes, seed, conventions, formulas, shapes, axes, parameter values and checked errors. Each saved record is immutable and content addressed; its SHA-256 is verified before the interface accepts it. The inspector never changes the trained weights. Selecting a forward-operation step is not selecting a training epoch.

28.2 Shapes explain the operation

Two different attention mechanisms illustrate why the axes matter. The trained quantum-feature model produces three feature vectors,

$$F \in \mathbb{R}^{3\times 3}, \quad \ell = Fq \in \mathbb{R}^3, \quad \alpha = \operatorname{softmax}(\ell), \quad c = \alpha^{\mathsf{T}} F \in \mathbb{R}^3.$$

Softmax here runs over three patches. The classical reference has seven token positions, including one padded node, model width six and two heads. After projection and splitting, its Q, K, V have shape $[1, 2, 7, 3]$. These heads are different from the three observables of the quantum circuit. Its Q, K also undergo the source's maximum shift and Euclidean normalization before dot scores.



The shape labels are literal exported dimensions. The two columns are different finite models; no equivalence of their learned functions is asserted.

For tropical aggregation, a scalar output is

$$C_{bhic} = \max_{j \text{ legal}} (S_{bhij} + V_{bhjc}), \qquad S = QK^{\mathsf{T}}/\sqrt{d_{\text{head}}}.$$

The inspector shows every maximizing key and the gap between the largest two legal values. A tie means equality of computed floating-point values, not a symbolic equality certificate. A query with no legal key has zero context by an explicit guard. Invalid block entries are stored as zero *with a validity mask*; including these zeros in a maximum would be incorrect. Softmax weights displayed in tropical records are comparison diagnostics rather than the tropical aggregation rule.

Soft-MoE dispatch normalizes across valid tokens for each expert-slot; combination normalizes across expert-slots for each valid token. These are different axes. Strict-prefix slots exclude the current and later token indices, while equal-rank attention can permit different tokens with the same reveal rank. The two rules must not be conflated into a stronger causality assertion. Likewise padding can acquire projection bias in a residual, so its mask must remain in later routing and pooling.

28.3 Quantum and Gaussian inspection

The stored StatePrep input precedes the displayed circuit. The variational gates have their exact wire order and angle metadata. Qubit zero is most significant. For QCNN9, the final CNOT targets start at wire five; the retained density matrix has shape $[32, 32]$, not a five-qubit pure state obtained by discarding amplitudes. The partial trace follows the unitary gates and is distinct from their dimension-preserving action.

The photonic model uses covariance order (x_0, x_1, p_0, p_1) and $\hbar = 2$. Its ellipse is a one-standard-deviation contour of a selected two-coordinate Gaussian marginal, with the marginal covariance taken from the actual intermediate record. It is not a plot of a trained interacting field theory. The Harper page separately displays the learned proposal and the fifteen exact rational principal-minor certificates. Rejected proposals are labelled as repaired. The carrier uses backward V, so $VU = iUV$; the live forward-clock example instead has $UV = \omega VU$.

28.4 A reproducible way to explore

Choose an evidence category and example, then a step. Clicking an input follows its recorded dependency. Head/token/coordinate controls select literal array slices; paging exposes large states without pretending that a small viewport is the entire tensor. Color is a display convention and the selected coordinate gives its numerical real and imaginary parts. A zero amplitude has undefined

phase and is labelled accordingly. Playback visits the forward pass, and stops when cases change or the tab becomes hidden. Orbit and zoom change the view, not the state or physical model.

The shared numerical records have independent quantum-backend and scalar reference checks. Browser checks exercise all saved cases, coordinate selection, masks, playback, view controls, circuits, covariance and certificate panels. These computational checks do not imply complete Lean formalization, generalization beyond the samples, quantum advantage or the arithmetic positive-determinant comparison. The general path/category and arithmetic interfaces remain the research obligations described in the final chapters.

Chapter 29

ToricGT: quantum features, tropical certificates and mirror interfaces

ToricGT organizes graph tokens, trajectories and auxiliary algebraic data; its companion research program separates architecture, dynamics, memory, oracle labels and toric training objectives [104, 105]. This chapter extends that organization to the present surface-order problem. A quantum feature map can replace a specified finite feature block, and exact algebra can check a proposed relation or parameter. Neither operation supplies the arithmetic identification needed by Hilbert–Pólya.

The inspected ToricGT revision is `a5199bc989a6fb170e21220e5e87bf95bc2fc0b4`; the program revision is `f91ba1811b5a4f70d53c7120acf8ca2a0049f85b`. The files `toric_algebra.py`, `tropical_toric_certificates.py`, `cas_certificates.py` and the graph-token schemas are the concrete interfaces used here. Their mathematical objects are not inferred from repository names. In particular, we do not claim to have reproduced the upstream language-model training campaigns or their compression scores.

29.1 A toric chart inside an actual dessin representation stratum

Let A be the infinite polynomial elliptic-square dessin algebra of the preceding chapters. It has two quiver vertices, two arrows in each direction, and the specified locally gentle zero relations. Retain one arrow $X : V_0 \to V_1$, put the other arrows equal to zero, and frame $V_0 = V_1 = k^2$. Every path of length two now acts by zero, so these are genuine continuous finite-dimensional modules of the completed order as well. Impose the additional closed condition $\operatorname{rank} X \leq 1$. In the chosen frames write

$$X = \begin{pmatrix} z_{00} & z_{01} \\ z_{10} & z_{11} \end{pmatrix}, \qquad f = z_{00}z_{11} - z_{10}z_{01}.$$

The coordinate ring is $B = k[z_{00}, z_{01}, z_{10}, z_{11}]/(f)$. This is a representation *stratum with a rank condition*; it is not an assertion that the full order or its invariant ring is a toric variety.

Proposition 29.1 (The Segre semigroup and its full relation)**.** *The homomorphism*

$$\begin{aligned} \pi : k[z_{00}, z_{10}, z_{01}, z_{11}] &\longrightarrow k[s, x, y], \\ z_{ij} &\longmapsto s x^i y^j \end{aligned}$$

has kernel (f)*. Its image is the semigroup ring generated by* $(1,0,0),(1,1,0),(1,0,1),(1,1,1)$ *in* $\mathbb{Z}^3$*. With all* z *of degree one,*

$$H_B(t) = \frac{1-t^2}{(1-t)^4} = \frac{1+t}{(1-t)^3}, \qquad \dim_k B_n = (n+1)^2.$$

Proof. The displayed binomial maps to zero. Replace any factor $z_{00}z_{11}$ in a monomial by $z_{10}z_{01}$ until no such factor remains. The integer kernel of the exponent matrix is generated by $(1,-1,-1,1)$ in the order $(00,10,01,11)$. Two distinct reduced monomials cannot have the same image: their exponent difference would be a nonzero multiple of this vector, and the monomial at the positive end would contain both z_{00} and z_{11}. Distinct reduced monomials therefore have linearly independent images. Reducing a polynomial proves $\ker \pi = (f)$. Multiplication by f is injective in the polynomial domain $P = k[z]$, so the graded exact sequence $0 \to P(-2) \xrightarrow{f} P \to B \to 0$ gives the Hilbert series. In degree n, image monomials are exactly $s^n x^a y^b$ with $0 \le a, b \le n$: choose the overlap of the a x-positions and b y-positions between $\max(0, a+b-n)$ and $\min(a,b)$. This gives $(n+1)^2$ monomials. □

This supplies a small complete certificate rather than an unspecified "syzygy loss." A learned output $\widehat{z}$ can be penalized by $|f(\widehat{z})|^2$. After rationalization, checking $f(z) = 0$ exactly certifies the rank-one relation. It does not certify the learned point's interpretation as a module with other nonzero arrows, an arithmetic endomorphism, or a positive spectral family. The multiplicative parametrization $z_{ij} = a_i b_j$ enforces this one relation identically, including degenerate rank-zero points.

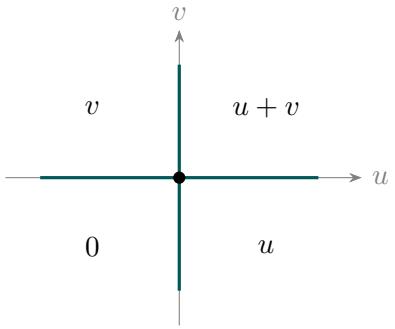
The projectivization of this cone is the Segre surface $\mathbb{P}^1 \times \mathbb{P}^1$. The full framed representation space, the rank-one cone, its projectivization and a quotient by the vertex group are different objects. Under the full left-right action on X, its individual entries are not all semi-invariants. The determinant has its usual determinant character, but is identically zero on this stratum. Framing is thus part of the example, not an invisible choice of gauge.

29.2 Tropical charts and exactly what their certificates prove

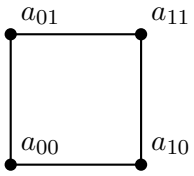
The same square of exponents defines the max-plus function

$$\psi(u,v) = \max\{0, u, v, u+v\} = \max\{0,u\} + \max\{0,v\}.$$

Its corner locus is the union of the coordinate axes. There are four one-dimensional rays of weight one, with primitive outgoing directions $(1,0),(-1,0),(0,1),(0,-1)$. Their sum is zero. This is an exact local balancing calculation. The Newton polygon is the unit square.



Four rays, weight one
local balance = 0



Dual Newton square
$a_{00} + a_{11} = a_{10} + a_{01}$

There are two associated tropical constructions, and they should not be conflated. The hypersurface of a four-term polynomial with this support lives in $\mathbb{R}^2$. The tropicalization of the two-term Segre ideal lives in $\mathbb{R}^4$ and satisfies

$$w_{00} + w_{11} = w_{10} + w_{01}.$$

The latter follows because the two terms of the binomial must have equal valuation. Over a valued field, tropicalization records valuations and initial degenerations; it is not a literal embedding of the tropical set as an ordinary algebraic subvariety of the complex toric variety. A fan specifies a toric compactification, while compatibility with that fan and the valued-field data specifies the relevant tropical geometry.

The inspected helper in `tropical_toric_certificates.py` uses integer supports, rational wall intersections and polygon edge normals. At this revision its input conversion applies `int` to floats. Consequently the exact claim here is restricted to inputs already verified to be integral; a nonintegral learned exponent cannot be silently rounded into an exact certificate. Pairwise ties also need their dimension checked. In the square example the pair $(00, 11)$ ties maximally only at the origin, as does $(10, 01)$. These zero-dimensional sets are not extra facets. A balancing sum for one dual polygon is not by itself a certified Chow class on an unspecified complete fan, nor an exact description of a higher-dimensional tropical variety. The stronger boolean names in the upstream record are not used as additional mathematical premises here.

A useful hybrid pipeline learns affine scores with real parameters, exports integer/rational candidates with an explicit approximation rule, and then recomputes active cells, dimensions, weights and balancing exactly. The approximation error belongs to the learned model; the certificate belongs to the exported data. Recording both prevents a small floating residual from becoming an alleged exact identity.

29.3 The quantum phase block and its cocycle

The ToricGT shift convention differs from the forward shift used in some other chapters. At $\theta = p/q$, with $\omega = e^{2\pi i\theta}$, its untwisted matrices act by

$$Ue_j = \omega^j e_j, \qquad Ve_j = e_{j-1 \bmod q}, \qquad VU = \omega UV.$$

If p, q are coprime they are a finite irreducible carrier of a central quotient. They do not faithfully represent the full rational rotation algebra. A boundary twist has to have modulus one if V is to be unitary. Rational approximants to an irrational θ produce a sequence of finite models, but not an automatically nested sequence of matrix-algebra inclusions.

Proposition 29.2 (Ordered and symmetric Weyl multipliers)**.** *For integer pairs $x = (a, b)$ and $y = (c, d)$ set $W(x) = U^a V^b$. Then*

$$W(x)W(y) = \omega^{bc} W(x + y).$$

Choose a real lift of θ and set $D(a, b) = e^{\pi i\theta ab} U^a V^b$. Then

$$D(x)D(y) = e^{\pi i\theta(bc-ad)} D(x + y), \qquad D(x)D(y) = e^{2\pi i\theta(bc-ad)} D(y)D(x).$$

These formulas are on $\mathbb{Z}^2$; descent to exponents modulo q requires retaining the resulting central phases.

Proof. Repeated use of $VU = \omega UV$, including inverses, gives $V^b U^c = \omega^{bc} U^c V^b$ and proves the first identity. The exponent introduced by the symmetric gauge is

$$\tfrac{1}{2}ab + \tfrac{1}{2}cd + bc - \tfrac{1}{2}(a+c)(b+d) = \tfrac{1}{2}(bc - ad).$$

Reversing x, y negates this last expression. Dividing the two multipliers proves the commutator formula. Associativity also follows directly from bilinearity: $bc + (b+d)e = de + b(c+e)$ for a third pair (e, f). □

The generic routine `projective_phase` instead evaluates $\exp(2\pi i x^T \Theta y)$. For skew Θ its commutator phase is $\exp(4\pi i x^T \Theta y)$. To match the preceding symmetric gauge, use the explicit matrix

$$\Theta = \begin{pmatrix} 0 & -\theta/2 \\ \theta/2 & 0 \end{pmatrix}.$$

This factor of two and the orientation sign matter for learned phase features, circuit comparisons and transport around a surface. A separate numerical skew matrix is not automatically the ordered multiplier ω^{bc}. For even q, writing $1/2$ modulo q would introduce another error; the real-lift formula avoids that operation.

For a density matrix ρ, one physically measurable feature is

$$f_x(\rho) = \operatorname{Tr}(\rho D(x)).$$

Its real and imaginary parts are expectations of the Hermitian operators $(D(x) + D(x)^*)/2$ and $(D(x) - D(x)^*)/(2i)$. They lie in $[-1, 1]$. Thus a classical graph-token block can receive measured finite-dimensional phase features; a variational circuit can learn the state preparation or bounded observables; and a classical head can combine these with exact combinatorial records. No claim of quantum speedup follows from replacing an already inexpensive matrix calculation by a circuit.

Proposition 29.3 (A genuine positive feature kernel)**.** *For any finite list x_i and any density matrix ρ, the matrix $K_{ij} = \operatorname{Tr}(\rho D(x_i)^* D(x_j))$ is positive semidefinite.*

Proof. For complex c_i put $B = \sum_i c_i D(x_i)$. Then $\sum_{ij} \overline{c}_i K_{ij} c_j = \operatorname{Tr}(\rho B^* B) = \operatorname{Tr}((B\sqrt{\rho})^* B\sqrt{\rho}) \geq 0$. □

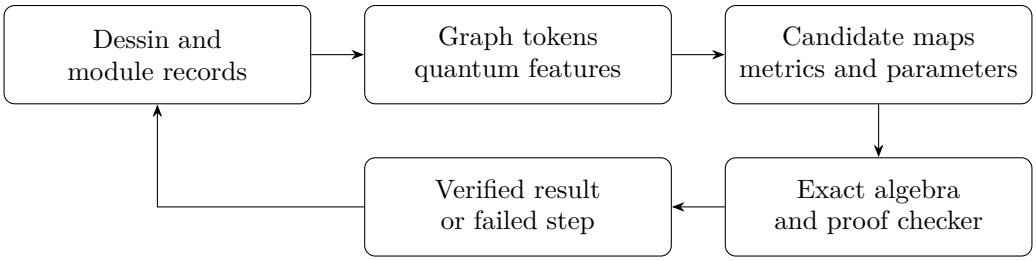
This proof applies to a physical Gram kernel. It does not identify that kernel with a Jensen metric, the arithmetic explicit-formula pairing, or the determinant family of a completed L-function. The latter comparison is an additional theorem, not a consequence of positivity of ρ.

29.4 A quantum–classical ToricGT extension with auditable interfaces

A typed record for a surface example should contain dart permutations, arrow endpoints, zero relations, central parameters, coefficient fields, module dimensions and category labels. Neural embeddings can be added to these records; they should not replace them. A candidate hybrid block is

$$h'_v = F_\eta\left(h_v, \sum_{w \sim v} \alpha_{vw} h_w, \{\Re f_x(\rho_v), \Im f_x(\rho_v)\}_x,\ c_v\right),$$

where α is a normalized classical attention rule, ρ_v is prepared by a specified finite circuit, and c_v is an exact certificate summary. Its architecture is meaningful even if the quantum features are first simulated classically, as in the finite training examples of the preceding computational chapter.



There are three distinct uses. First, supervised regression can learn a coarse-graining map against exact Ising or Gaussian targets, with a held-out parameter set. Second, a proposal model can search for a small module map or a polynomial identity, checked by exact substitution and stated presentation relations. Third, a proof-search policy can select candidate lemmas or tactics, whose accepted output must pass the proof assistant. Only the last two have the potential to yield new exact mathematical results, and only when the checker verifies the full exported statement. A classifier predicting that a proof is valid is not such a checker.

The ToricGT program's certificate distinction is particularly useful here: store exact CAS data, validated surrogates and unavailable backends as different provenance classes. When comparing a quantum feature head with a classical one, hold the data split, parameter budget and objective fixed; report sample or simulation cost as well as error. A network with extra features need not improve the objective. None of the finite experiments in this book is an upstream full-transformer language-model ablation.

An exact candidate metric G for a finite matrix M illustrates the proposal/check distinction. For rational real entries, symmetry and all principal-minor nonnegativity can certify $G \succeq 0$. Positive definiteness and $GM = M^*G$ then imply similarity to a Hermitian matrix: $G^{1/2}MG^{-1/2}$ is Hermitian. Checking a few sampled minors or a small floating residual does not prove these assertions, and applying the result to an arithmetic family still requires uniform positivity and the exact arithmetic determinant comparison.

29.5 Mirror symmetry supplies category comparisons, not a spectral shortcut

Kontsevich–Soibelman develop torus-fibration and homological-perturbation methods, including an A_∞ comparison for specified subcategories of abelian mirrors [106, Sections 6–8, Theorem 4]. A lattice of phase labels is useful computational data for such geometry, but a finite Weyl pair alone does not construct either Fukaya objects or an A_∞ functor. The latter must carry complexes of morphisms and all higher composition maps with their signs and convergence rules.

A concrete toric theorem provides a better model for the proposed link. For a smooth projective toric X, Abouzaid constructs a full A_∞ sub-pre-category of the relative Fukaya category of its Landau–Ginzburg mirror that is quasi-equivalent to the DG category of line bundles on X [107, Theorem 1.2 and Corollary 1.4]. The theorem gives an embedding after the indicated derived/split closure; it does not assert an unrestricted equivalence for every non-Fano toric variety or for our singular completed order.

For intuition, the customary toric mirror potential for the Segre surface is

$$W(z,w) = z + z^{-1} + w + w^{-1}, \qquad (z,w) \in (\mathbb{C}^*)^2.$$

Solving $z\partial_z W = w\partial_w W = 0$ gives $(z,w) = (\pm 1, \pm 1)$, with critical values $4, 0, 0, -4$. The logarithmic Hessian is diagonal with entries $z + z^{-1}$ and $w + w^{-1}$ and is nonsingular at those four points. This

is a direct calculation of a mirror *potential*. By contrast the two-level clock matrix $2Z + 2X$ has eigenvalues $\pm 2\sqrt{2}$, since $XZ + ZX = 0$. The critical values and matrix eigenvalues are plainly different data. A proposal to relate them must specify a functor or quantization and its normalization; visual resemblance of their torus coordinates is not that comparison.

For gentle surfaces, the precise positive input already used in this book is the equivalence for a homologically smooth graded locally gentle algebra with its prescribed stopped surface and line field [66, Theorem 3.2.2]. Our nonempty closed dessin algebras have forbidden cycles and fail that smoothness hypothesis. A stopped quotient or categorical resolution can be a meaningful intermediate object, but it requires its own functor. Ordinary exact quotient inflation does not automatically preserve derived Hom spaces, AR triangles or determinants. In particular, the polynomial node branch $k[x]$ has a valid stopped-annulus interpretation, whereas inflation into $k[x, y]/(xy)$ is not derived fully faithful, as proved earlier.

A finite training loss for $m_1^2 = 0$ or a chain-map residual checks only the displayed truncation. A genuine category construction additionally needs the remaining A_∞ equations, admissible geometric objects, grading, compactness or convergence, and the claimed generation result. Tropical gradient trees can guide computations of compositions under appropriate geometric hypotheses; they do not remove those obligations. This is a constructive research interface: learn a small candidate, calculate its maps exactly, and then prove the specified comparison.

29.6 What would make the learned component relevant to GRH?

For the present Hilbert–Pólya program the useful learned outputs are candidates for arithmetic correspondences, controlled positive metrics, coefficient kernels outside the classes already refuted, or estimates with an exact symbolic witness. Every candidate must be checked against local Frobenius factors, central and deck actions, determinant multiplicities, conductor and gamma normalizations, and compact-uniform convergence on the stated complex domain. A Dirichlet comparison is not a proof for Hecke or general Artin L-functions.

An exact finite witness can disprove a proposed universal construction, and that is valuable progress. A successful finite fit only proposes its next theorem. Neither a toric binomial certificate nor a positive quantum feature kernel proves that all arithmetic approximants have their zeros on the critical line. The small-gcd analytic estimate and the independent positive arithmetic determinant comparison therefore remain separate load-bearing research tasks. The role of the hybrid architecture is to make their search more explicit, testable and reproducible.

29.7 An executed hybrid spectral experiment

The executable example `examples/quantum/toricgt_hybrid.py` imports the pinned ToricGT rotation-algebra module. For $\theta = 1/4$ it obtains four-dimensional Weyl generators and forms the Hermitian Harper family

$$H_\lambda = U + U^* + \lambda(V + V^*), \qquad 0.1 \le \lambda \le 1.8.$$

A two-qubit PennyLane feature circuit and affine head, with twelve parameters, learn the smallest eigenvalue divided by six. Seed 20261104 specifies one hundred sampled couplings, with seventy-five training and twenty-five held-out samples. After 160 Adam updates the held-out mean-square loss is 4.01905×10^{-5}, compared with 0.220518 at initialization. These are finite simulated spectra, not L-function zeros.

For six rational test couplings, the program forms the exact rational matrix

$$H_\lambda = \operatorname{diag}(2, 0, -2, 0) + \lambda \operatorname{Adj}(C_4).$$

It rounds a learned proposal to a rational number E and checks all fifteen nonempty principal minors of $H_\lambda - EI$. A real symmetric matrix is positive semidefinite exactly when every principal minor is nonnegative, so an accepted proposal is an exact lower eigenvalue bound. A rejected proposal is repaired by rational bisection, starting from the independent Gershgorin lower bound $-2 - 2\lambda$ and the Rayleigh upper bound -2. The saved rational minors certify the resulting lower bounds. Thus learning proposes; an independent exact criterion accepts or repairs. The certificate does not depend on the neural prediction being accurate. No arithmetic operator or completed-order equivalence is claimed by this experiment.

This benchmark has an exact closed-form solution: $\det(tI - H_\lambda) = t^2(t^2 - 4(1 + \lambda^2))$, so $E_0 = -2\sqrt{1 + \lambda^2}$. For $\lambda > 0$, rows one and three of H_λ agree, and the difference of rows zero and two is $2/\lambda$ times row one. Thus its rank is at most two. Its trace is zero and $\operatorname{tr}(H_\lambda^2) = 8(1 + \lambda^2)$, giving the stated eigenvalues (and hence characteristic polynomial). The experiment deliberately has an exact oracle: it tests learning, certificate acceptance and repair on a known problem rather than presenting a numerical fit as a new spectral theorem.

Chapter 30

Two-parameter persistence, Schur syzygies and surface-order coefficients

A persistence diagram records how objects and maps change under an ordered family of inclusions. A projective resolution records the generators, relations and relations among relations of one module. These descriptions meet naturally over a polynomial parameter ring. They must still be distinguished from the infinite locally gentle algebra attached to the surface: the parameter ring describes the filtration, not a replacement surface algebra.

The three supplied papers motivate these interfaces [108, 109, 110]. We use their surface-order, lattice and finite-presentation ideas with the corrections already proved in this book. In particular, a matrix satisfying consecutive-product zero is a complex, not automatically a resolution; Fitting base change is not exact base change; and the existence of a generic resolution is not a finite terminating algorithm. The worked example below proves its own exactness and exhibits an actual topological realization.

30.1 The polynomial ring encodes the two commuting directions

Let k be a field and give $\mathbb{N}^2$ its coordinatewise order. A persistence module is a covariant functor $P : \mathbb{N}^2 \to \mathrm{Vect}_k$. Writing $P_{a,b} = P(a, b)$, put

$$R = k[x, y], \quad \deg x = (1, 0), \quad \deg y = (0, 1), \qquad \mathcal{P} = \bigoplus_{a,b \geq 0} P_{a,b}.$$

Multiplication by x and y is the map to the next coordinate in the corresponding direction. Functoriality makes their square commute. Conversely the graded pieces and multiplication maps of an $\mathbb{N}^2$-graded R-module define such a functor. General transition maps are multiplication by $x^{c-a}y^{d-b}$ when $(a, b) \leq (c, d)$. The exponents are nonnegative in this direction, a useful sign check.

This gives an equivalence for the stated discrete diagrams. Finite generation is an additional condition, not a consequence of having finite-dimensional pieces. A module can have infinitely many independent births. Finite cellular input with finitely many birth grades gives the finite presentations used here. A real-parameter filtration needs an appropriate discretization or a different coefficient algebra before being treated as a module over $k[x, y]$.

We write $R(-a, -b)$ for a free module whose generator has grade (a, b). Over a field, a finite graded module has a minimal graded free resolution of length at most two. Minimality means that every differential entry belongs to (x, y); tensoring with $k = R/(x, y)$ then kills the differentials. The numbers and grades of the surviving generators are the multigraded Betti numbers. The field

hypothesis matters: $\mathbb{Z}[x, y]$ has homological dimension three, and the regular sequence $(2, x, y)$ supplies a length-three example. The particular resolution below also works integrally, but that does not make every integral resolution length two.

30.2 A complete three-grade example

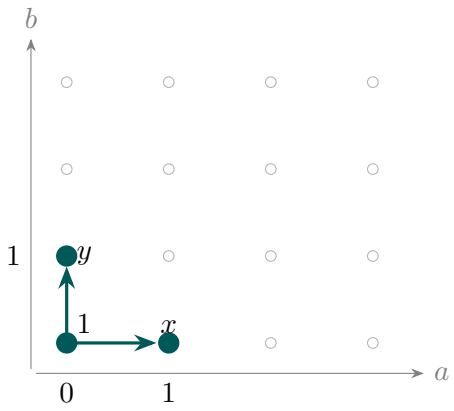
Set

$$M = R/(x^2, xy, y^2).$$

Its basis is $1, x, y$, in grades $(0,0), (1,0), (0,1)$. All other graded pieces vanish. The two nonzero proper transition maps send 1 to x and 1 to y. In particular, for comparable grades $\nu \leq \mu$,

$$\rho_M(\nu, \mu) = \operatorname{rank}(M_\nu \to M_\mu) = \begin{cases} 1 & \nu, \mu \in \{(0,0), (1,0), (0,1)\}, \\ 0 & \text{otherwise.} \end{cases}$$

This computes this module's rank invariant; rank invariants do not in general classify all two-parameter modules.



Filled point: one copy of k
Open point: zero
Only the two thick arrows survive

Figure 30.1: The actual rank diagram of $R/(x, y)^2$. Moving right or up outside its three-point support kills the class. Incomparable grades are not joined by a persistence map.

Theorem 30.1 (A minimal bigraded resolution)**.** *The following complex is a minimal graded free resolution of M:*

$$0 \longrightarrow F_2 \xrightarrow{d_2} F_1 \xrightarrow{d_1} R \longrightarrow M \longrightarrow 0,$$

where

$$F_1 = R(-2,0) \oplus R(-1,-1) \oplus R(0,-2),$$

$$F_2 = R(-2,-1) \oplus R(-1,-2), \qquad d_1 = \begin{pmatrix} x^2 & xy & y^2 \end{pmatrix}, \quad d_2 = \begin{pmatrix} -y & 0 \\ x & -y \\ 0 & x \end{pmatrix}.$$

Proof. Each column of d_2 has the indicated source grade, and multiplication gives $d_1 d_2 = 0$. If $d_2(a, b) = 0$, its first and last entries give $ya = xb = 0$. Since R is a domain, $a = b = 0$.

Suppose $x^2 f + xyg + y^2 h = 0$. Reducing modulo x shows $x \mid h$, say $h = xb$. Dividing the relation by x and reducing modulo y shows $y \mid f$, so write $f = -ya$. Substitution now gives

$y(-xa+g+yb)=0$, hence $g=xa-yb$. Thus $(f,g,h)=d_2(a,b)$. This proves exactness at F_1. The image of d_1 is its displayed ideal, giving exactness at R and the augmentation. Every entry of both matrices has positive grade, so the resolution is minimal. The divisibility argument also works in $\mathbb{Z}[x,y]$, where x,y remain prime and nonzerodivisors. □

The multigraded Betti numbers are consequently

$$\beta_{0,(0,0)}=1,\quad \beta_{1,(2,0)}=\beta_{1,(1,1)}=\beta_{1,(0,2)}=1,\quad \beta_{2,(2,1)}=\beta_{2,(1,2)}=1.$$

They are dimensions of the corresponding graded $\operatorname{Tor}_i^R(M,k)$. Taking degreewise Euler characteristics gives the same Hilbert series by two independent descriptions:

$$H_M(s,t)=1+s+t=\frac{1-s^2-st-t^2+s^2t+st^2}{(1-s)(1-t)}.$$

After forgetting the separate grading this becomes $(1-3z^2+2z^3)/(1-z)^2=1+2z$. The numerator records the alternating resolution shifts, not a characteristic polynomial whose zeros locate arithmetic L-zeros.

One can also see the dual calculation directly. Ignoring shifts only to write the matrices, the dual differential is

$$d_2^*=\begin{pmatrix}-y & x & 0\\ 0 & -y & x\end{pmatrix}.$$

The equations $-yA+xB=-yB+xC=0$ imply $(A,B,C)=(x^2q,xyq,y^2q)$, again using coprimality. Thus $\operatorname{Ext}^0_R(M,R)=\operatorname{Ext}^1_R(M,R)=0$. The second Ext is generated by e_1,e_2 with relations $ye_1=0$, $xe_1=ye_2$ and $xe_2=0$; its k-basis is e_1,e_2,xe_1. This keeps the syzygy calculation distinct from merely counting the support points of M.

30.3 A bifiltered CW realization and the syzygies it creates

Start with one vertex and an oriented loop e, both born at grade $(0,0)$. Attach three two-cells by degree-one maps to that loop, at grades

$$D_{20}:(2,0),\qquad D_{11}:(1,1),\qquad D_{02}:(0,2).$$

At grade (a,b) include precisely the cells whose birth grades are at most (a,b). Every boundary is already present, so these are actual subcomplex inclusions and their squares commute. The persistence cellular chain modules are

$$C_2=F_1,\qquad C_1=Re,\qquad C_0=R,\qquad \partial_2=(x^2\ \ xy\ \ y^2),\quad \partial_1=0.$$

Therefore $H_1=M$ exactly. No assertion that every abstract persistence module has this particular geometric form is needed.

The loop survives before any disk is present and dies when the first disk appears. Two disks attached along their common boundary form a sphere. The first two persistent H_2 generators are the compatible cell cycles

$$-yD_{20}+xD_{11}\quad\text{at }(2,1),\qquad -yD_{11}+xD_{02}\quad\text{at }(1,2).$$

Theorem 30.1 proves $H_2\cong R(-2,-1)\oplus R(-1,-2)$; there are no three-cells to kill them. At $(2,2)$ all three disks give two independent sphere classes. Thus the same matrix d_2 is a relation among the killing cells and a positive persistent homology class in the larger CW complex. This auxiliary CW space need not be a Riemann surface: three disks along one circle are not a manifold. It is a topological realization of the filtration module, not another claimed cellular dessin.

30.4 Schur modules organize the generators and the first syzygies

Assume now $\operatorname{char} k = 0$, put $V = kx \oplus ky$, and let $\mathrm{GL}(V)$ act on $R = \operatorname{Sym} V$ by linear substitution. The ideal $(x, y)^2$ is invariant. Its degree-two generators form $\operatorname{Sym}^2 V$, with basis x^2, xy, y^2. The linear syzygies form the kernel of multiplication

$$V \otimes \operatorname{Sym}^2 V \longrightarrow \operatorname{Sym}^3 V.$$

There is a natural equivariant map

$$\delta : \bigwedge^2 V \otimes V \longrightarrow V \otimes \operatorname{Sym}^2 V, \qquad (v \wedge w) \otimes z \longmapsto v \otimes wz - w \otimes vz.$$

Multiplication kills its image. On the two basis vectors its values are exactly the two columns of d_2:

$$(x \wedge y) \otimes x \mapsto x \otimes xy - y \otimes x^2, \qquad (x \wedge y) \otimes y \mapsto x \otimes y^2 - y \otimes xy.$$

They are independent, and the kernel has dimension $6 - 4 = 2$, so δ identifies it with $\det V \otimes V = S_{(2,1)}V$. The Littlewood–Richardson rule gives

$$V \otimes \operatorname{Sym}^2 V \cong \operatorname{Sym}^3 V \oplus S_{(2,1)}V.$$

Here the partition convention uses symmetric rows. The equivariant total- graded resolution has multiplicity spaces $k, \operatorname{Sym}^2 V, S_{(2,1)}V$ in generator degrees $0, 2, 3$. The diagonal torus recovers their bigraded weights $(2,0), (1,1), (0,2)$ and $(2,1), (1,2)$.

The grading distinction is essential. General GL_2 substitutions mix x and y, so they preserve total degree, not the two original filtration coordinates. Only the diagonal torus preserves each bigrading. In positive characteristic the displayed maps and our exact resolution remain valid, but semisimple direct-sum representation theory must not be assumed. Nor does applying a Schur functor termwise to an arbitrary exact sequence preserve exactness, as the earlier counterexamples show.

30.5 Fitting data, exactness and specialization

For a finite presentation $R^p \xrightarrow{D} R^q \to N \to 0$ define

$$\operatorname{Fitt}_j(N) = I_{q-j}(D), \qquad I_0(D) = R,$$

where I_r is the ideal of r-minors, $I_r = R$ for $r \le 0$, and $I_r = 0$ for $r > \min(p, q)$. The ideal is independent of the chosen finite presentation. Its formation commutes with arbitrary base change because minors are polynomials in the presentation entries. For our cyclic module,

$$\operatorname{Fitt}_0(M) = (x, y)^2, \qquad \operatorname{Fitt}_j(M) = R \quad (j \ge 1).$$

The support is the origin, but its Fitting scheme retains a length-three thickening. This is not the same datum as the rank invariant of the ordered graded pieces.

The 2-minors of d_2 are $y^2, -xy, x^2$, so their ideal is also $(x, y)^2$. Expected ranks are $1, 2$ and this ideal has grade two; x^2, y^2 is a regular sequence. These facts agree with the finite-free acyclicity criterion of Buchsbaum–Eisenbud [111], independently of our direct proof. Over a general Noetherian ring, a height count cannot silently replace the criterion’s grade/regular-sequence hypotheses. The

criterion concerns a finite complex over a commutative ring; it does not certify every infinite path-algebra resolution.

Arbitrary specialization does not preserve that exactness. At $x = y = 0$ both matrices vanish, leaving positive homology k^3, k^2 in degrees one and two. More informatively, after setting only $y = 0$, the complex has

$$H_0 = k[x]/(x^2), \qquad H_1 = (k[x]/(x))^2, \qquad H_2 = 0.$$

These follow immediately from $d_1 = (x^2, 0, 0)$ and the two surviving entries x of d_2. They are the derived-base-change Tor groups. Flat base change, suitable Tor vanishing, or a new exactness proof is needed to obtain an ordinary specialized resolution. A generic parameter space and a small differential residual do not remove those conditions.

30.6 A correct finite space of complexes and its incidence model

For variety language in this section take k algebraically closed. Let V_i have dimensions n_i for $i = 0, 1, 2$. Fix feasible ranks a, b with

$$a \le \min(n_0, n_1), \qquad b \le \min(n_1, n_2), \qquad a + b \le n_1.$$

The reduced variety

$$\mathcal{C}_{a,b} = \{(A, B) : V_2 \xrightarrow{B} V_1 \xrightarrow{A} V_0,\ AB = 0,\ \operatorname{rank} A \le a,\ \operatorname{rank} B \le b\}$$

parametrizes based complexes with these upper bounds, not isomorphism classes of resolutions. Consider flags $L \subset K \subset V_1$ with $\dim L = b$, $\dim K = n_1 - a$, and impose

$$\operatorname{im} B \subset L \subset K \subset \ker A.$$

Over the flag variety the allowed maps form the vector bundle

$$\operatorname{Hom}(V_1/K, V_0) \oplus \operatorname{Hom}(V_2, L).$$

Its total space is smooth and irreducible. Projection is proper onto $\mathcal{C}_{a,b}$: for any complex with the rank bounds, first extend $\operatorname{im} B$ to a b-plane inside $\ker A$, then extend it to an $(n_1 - a)$-plane there. On the full-rank locus the choices are uniquely $L = \operatorname{im} B$, $K = \ker A$. Full ranks exist by feasibility, so this projection is birational. This proves the asserted incidence resolution of the reduced rank-bound variety.

For $(n_0, n_1, n_2) = (1, 3, 2)$ and $(a, b) = (1, 2)$, the two flag spaces coincide in a plane, and the base is $\mathbf{Gr}(2, k^3)$. The vector bundle has fiber dimension $1 + 4$ and base dimension two, giving total dimension seven. Full-rank points are exact in positive degrees because $\operatorname{im} B = \ker A$ and B is injective. The zero pair is also a point of the complex variety and is not exact. Our polynomial matrices map the punctured plane to the exact locus; the origin maps to that nonexact boundary. A specialization of the universal matrices therefore still needs its own homological check.

Bruns' generic-resolution construction uses a directed limit of rings which kills positive homology in the limit [112, Theorem 1, pp.35–36]. It does not assert that some finite stage always becomes exact, nor that every specialization of the resulting matrices preserves exactness. The finite example in this chapter makes its stronger local claims by explicit calculation rather than by invoking an unspecified generic object.

30.7 Coefficients in a genuine infinite dessin order

Return to the actual nodal interval dessin and distinguish its variables:

$$A = k[\xi, \eta]/(\xi\eta), \qquad \Lambda = k[[\xi, \eta]]/(\xi\eta), \qquad u = \xi + \eta.$$

Both A and Λ are infinite dimensional over k. The completed order is free of rank two over $k[[u]]$, on $1, \xi$, with $\xi^2 = u\xi$ and $\eta = u - \xi$. Multiplication by u is injective, as is clear in its two normalization branches. Set $E = \Lambda/u^2\Lambda$. It has basis $1, \xi, u, u\xi$ over k, nontrivial arrow actions, and

$$0 \longrightarrow \Lambda \xrightarrow{u^2} \Lambda \longrightarrow E \longrightarrow 0$$

is a genuine free resolution. The module E is finite length, not a $k[[u]]$-free lattice. Its finite dimension does not turn its parent order into a finite-dimensional substitute.

The family $E \otimes_k M$ is a persistence module valued in continuous Λ-modules: each nonzero piece is E, and its structure maps are the identity or zero as in Figure 30.1. The actions of x, y commute with the fixed action of u, ξ, η. There is no identification of either persistence direction with the order uniformizer.

More precisely put $\mathcal{B} = \Lambda \otimes_k k[x, y]$. Tensor the length-one Λ-resolution just displayed with Theorem 30.1 over k. Its finite total complex has free $\mathcal{B}$-ranks

$$(\operatorname{rank} F_0, \operatorname{rank} F_1, \operatorname{rank} F_2, \operatorname{rank} F_3) = (1, 4, 5, 2)$$

and resolves $E \otimes_k M$. If Q_i denotes the order resolution and P_j the persistence resolution, the total differential on $Q_i \otimes P_j$ is

$$D = \partial_Q \otimes 1 + (-1)^i 1 \otimes d_P.$$

Order the degree-n terms as $P_n \oplus (Q_1 \otimes P_{n-1})$. Suppressing only the tensor symbols, the three matrices are

$$D_1 = \begin{pmatrix} d_1 & u^2 \end{pmatrix}, \qquad D_2 = \begin{pmatrix} d_2 & u^2 I_3 \\ 0 & -d_1 \end{pmatrix}, \qquad D_3 = \begin{pmatrix} u^2 I_2 \\ -d_2 \end{pmatrix}.$$

The cross terms cancel, so consecutive products are zero. For exactness, write a positive-degree cycle as (a, b), with $da + u^2 b = 0$ and $db = 0$. Modulo u^2 the first component is a cycle in $E \otimes_k P_\bullet$. That complex is exact in positive degrees, since tensoring vector spaces over k is exact. Hence $\bar{a} = d\bar{c}$. Lift $\bar{c}$ to c in the free term and write $a - dc = u^2 e$. The original equation becomes $u^2(de + b) = 0$. Multiplication by u^2 is injective on every free term, so $b = -de$. Consequently $(a, b) = D(c, e)$, as required. The same argument includes the top degree by putting missing terms equal to zero. Degree-zero homology is $\mathcal{B}/(u^2, x^2, xy, y^2) = E \otimes_k M$. This proves exactness without assuming that an arbitrary tensor product of resolutions stays exact. The ranks follow from $(1 + z)(1 + 3z + 2z^2) = 1 + 4z + 5z^2 + 2z^3$.

This construction is compatible with the original center action. It does not shorten the infinite resolution of the nodal residue simple, make E a lattice, or construct Frobenius. Restriction to the commutative coefficient ring $k[[u]][x, y]$ doubles these free ranks. Determinant functors apply to the specified finite perfect complex, with virtual rank $1 - 4 + 5 - 2 = 0$, but determinant multiplicity and arithmetic comparison remain additional questions.

30.8 Quantum data, monotonicity and the next usable interface

An arbitrary time-varying interaction graph is not a bifiltration: an edge can disappear at a later time. One valid alternative fixes finite birth pairs (a_e, b_e) for edges and includes an edge exactly

when both coordinates are reached, then takes the flag complex. Another uses a cumulative union over times up to the first coordinate and decreasing interaction thresholds in the second. Both give inclusions; the second records cumulative history rather than individual snapshots. Vertices and higher cells must have compatible birth grades. Zigzag persistence is a different appropriate tool when deletions must be retained.

Quantum channels are linear on density matrices, but two unrelated channels need not commute. Assigning them to the two directions does not define a persistence functor unless the required squares commute. One may instead extract a declared classical correlation graph, choose a monotone filtration, and compute its homology with the conventions above. The data-to-distance rule and its statistical uncertainty still need justification.

For the research program this yields concrete trainable and checkable targets: propose birth grades, presentation entries or a low-degree syzygy; test inclusions, gradings, the matrix identities and actual exactness; record Betti and Fitting data before interpreting them. A loss penalizing D^2 checks only the complex equation. The exact module example, the flag-incidence model and the surface-order coefficient resolution show how to add the missing checks. None identifies those finite invariants with a completed arithmetic L-function, provides a positive arithmetic spectral limit, or proves GRH.

30.9 The executable staircase and quantum laboratory

The canonical on-page tool accompanies this chapter; other chapters link here. Proposition 2.3 generalizes the worked square-ideal resolution to every strict Artinian staircase. Four displayed ideals have lengths $3, 6, 13, 16$, with custom natural exponents also allowed. The tool shows generator and adjacent-syzygy degrees, standard monomials, exact comparable-grade maps and cellular incidence ranks. Coefficients in E multiply field dimensions by four; the displayed free ranks over $\Lambda[x, y]$ instead come from the u^2 cone. These are different rings.

Specializing the free resolution at a point over $\mathbb{F}_2$ computes its fiber-complex homology. At the origin its dimensions are $(1, m, m-1)$; at each of the three nonzero $\mathbb{F}_2$ points they vanish, since a pure-power generator becomes a unit. This is derived fiber data, not the CW homology at a filtration grade. The signed maximal minors and Hilbert numerator are checked independently with symbolic polynomials.

The quantum tab uses the single saved trained four-qubit QCNN unitary U and the channel $T(\rho) = \mathrm{Tr}_{2,3}(U\rho U^*)$. The designed inputs have their last three qubits in $|000\rangle$. Their first qubit runs over eight equatorial states $(|0\rangle + e^{2\pi ij/8}|1\rangle)/\sqrt{2}$ and both poles. This is an out-of-training-distribution illustrative family, not new training or a held-out performance claim. Birth labels are fixed from the input population $(I + Z_0)/2$: south 0, equator 500, north 1000. They are retained at the output; recomputing them from output energy would not justify the same identity inclusion.

Floating-point input/output trace distances are scaled by 1000 and rounded upward. Their frozen integer graph labels are saved with a SHA-256 checksum and checkpoint provenance. The inequality of output and input labels is verified directly, never imposed by taking their minimum. Proposition 2.4 explains the exact-channel contractivity; the numerical records are not interval certificates. For the recorded labels, all finite boundary matrices and induced ranks are computed exactly over $\mathbb{F}_2$. At grade $(400, 500)$ the input contains an equatorial 8-cycle and an isolated south pole, with $(\beta_0, \beta_1) = (2, 1)$. Its H_1 map to $(600, 500)$ has rank one, whereas its map to $(800, 1000)$ has rank zero. The plotted two-parameter rank table fixes the target and varies comparable source grades; it is not a complete multiparameter barcode. Rotating the drawing changes the input-qubit Bloch layout only. The same layout is used at output and does not embed its two-qubit state space.

Independent row elimination and the browser's bit-column implementation agree on both Betti numbers and inclusion ranks for the saved cases. These executable finite examples complement the proofs and the cited persistence theory. The worked square-ideal module now has a kernel-checked ungraded free resolution, described below. General staircase exactness, graded and categorical interfaces, the coefficient cone and channel norm contractivity remain separate formalization tasks. No learned component or finite homology calculation supplies the unresolved positive arithmetic determinant bridge.

30.10 What the polynomial Lean resolution actually verifies

The file `PersistenceResolution.lean` uses $R = \mathrm{MvPolynomial}(\mathrm{Fin}\,2, k)$ for an arbitrary field. Its free coordinate modules are $F_1 = R \times R \times R$ and $F_2 = R \times R$, with the matrices of Theorem 30.1. The augmentation is the actual R-linear quotient map to $R/(x^2, xy, y^2)$. The proof of `d1_kernel` quantifies over all polynomial coefficient vectors; it is not a sampled scalar-rank test.

The reusable lemma `relation_decomposition` starts with a commutative domain and a pair x, y such that $x \neq 0$ and $x \mid yz$ implies $x \mid z$. From $x^2a + xyb + y^2c = 0$ it constructs p, q with $(a, b, c) = (-yp, xp - yq, xq)$. For the actual polynomial variables, the regularity premise is discharged by viewing R as polynomials in x with coefficients in $k[y]$. A polynomial is divisible by x exactly when its constant coefficient vanishes; multiplication by the nonzero y cannot kill such a coefficient. Thus the final kernel theorem has no unproved regularity hypothesis. The first and last coordinates then prove `d2_injective`; `d1_range` identifies the image with the specified ideal. Augmentation exactness and surjectivity complete the sequence, and the free terms are proved projective.

The field k carries the residue action $r \cdot a = \mathrm{constantCoeff}(r)a$. The free modules retain ordinary polynomial multiplication. Both actions are explicit in the source, with action identities checked in the kernel. For every R-linear map from the appropriate free term to this residue field, precomposition with d_1 or d_2 is zero, since all displayed entries contain x or y. These are actual Hom statements with their module actions, not merely entrywise evaluations. Abstract Ext-group identifications and the graded minimal-resolution interface are not yet packaged as formal theorems.

All theorem types and axiom dependencies are printed. Only the standard Lean foundations occur; no project axioms or incomplete proofs are used. This establishes the entire underlying ungraded worked resolution and its residue Hom differential calculation. It does not formalize every staircase, the CW model, Schur equivariance, the E-coefficient tensor cone or any arithmetic determinant comparison. Those distinctions remain visible on the local verification page beside the executable laboratory.

Chapter 31

The remaining arithmetic theorem and a verifiable research loop

31.1 A precise positive spectral target

For a primitive nonprincipal Dirichlet character χ, choose the completed L-function with its parity, conductor and gamma factor fixed. A sufficient spectral program requires holomorphic functions F_N on a connected domain containing the critical strip, compact-uniform convergence $F_N \to \Lambda(s,\chi)$, and a proved zero-location property for every F_N. A self-adjoint determinant representation must include its domain, determinant regularization, nowhere-zero normalization and coefficient comparison. Self-adjointness of an unrelated finite matrix does not establish any of these identifications.

If every F_N is zero-free off the critical line, Hurwitz on each connected off-line component shows that the nonzero holomorphic limit is zero-free there. The limit cannot be identically zero on an open component if it is a nonzero holomorphic function on the connected containing domain. This last connected-domain hypothesis matters: componentwise convergence alone does not rule out an identically zero limit on a disconnected component. The conclusion is conditional until the arithmetic approximants and positive representation are constructed. General Artin functions also need their own holomorphy statement; Brauer quotients do not supply it. Hecke, Dedekind and elliptic-curve normalization and critical centers must remain distinct.

The original modular-geodesic argument fails at its prime identification. Integer hyperbolic traces and $e^{\ell/2} + e^{-\ell/2} = \sqrt{p} + 1/\sqrt{p}$ cannot match all rational primes. An asymptotic similarity between prime and geodesic counts does not give a termwise equality of Euler factors. Nor does finite-field purity for a family of curves prove number-field GRH. The local Jacobi theta construction is a positive arithmetic benchmark with its own coefficient and symmetry obstructions, treated in the earlier chapters.

31.2 A surviving short-row reduction

Write $\mathcal{O} = \mathbb{Z}[\omega]$. For squarefree primary column ideals n of norm about D, good away from a fixed set containing 2 and 3, let χ_n be the cubic residue character extended by zero at nonunits. At the unit ideal take $\chi_1(u) = 1$, including $u = 0$. For row-independent coefficients $|c_n| \le 1$, the normalized Gram expansion is

$$\mathcal{E} = D^{-1} \sum_{n,m} c_n \overline{c_m} K_H(n,m), \qquad K_H(n,m) = \sum_{0<Nu\le H} \chi_n(u)\overline{\chi_m(u)}.$$

For $n = ha, m = hb$ with h, a, b squarefree and pairwise coprime,

$$\chi_n(u)\overline{\chi_m(u)} = \mathbf{1}_{(u,h)=1}\psi_{a,b}(u), \qquad \psi_{a,b} = \chi_a\overline{\chi_b}.$$

The oscillatory primitive conductor is ab; its zero-preserving periodic modulus is hab. Cancelling h without its coprimality mask would change the row family. Möbius inversion gives

$$K_H(ha, hb) = \sum_{d|h} \mu(d)\psi_{a,b}(d) \sum_{0<Nv\le H/Nd} \psi_{a,b}(v).$$

A row weight has to be evaluated at dv in this identity.

For a fixed nonnegative radial Schwartz profile Φ, replace the sharp row cutoff by $\Phi(u/\sqrt{H})$. For an off-diagonal pair, $\psi_{a,b}$ is primitive nonprincipal. If $Q = N(ab)$, its complete Gauss transform and lattice Poisson summation give, for every $X > 0$,

$$\left|\sum_{v\in\mathcal{O}} \psi_{a,b}(v)\Phi(v/\sqrt{X})\right| \ll_\Phi \sqrt{Q}.$$

Indeed the nonzero dual frequencies have Gauss modulus at most $\sqrt{Q}$ and the zero frequency vanishes. After scaling, Fourier decay bounds their sum by $O_\Phi(Q/X)$: for $X \le Q$ use planar annuli; for $X > Q$ use a decaying power greater than two. The Poisson prefactor is $X/\sqrt{Q}$. The ideal generator rotates and scales the Eisenstein lattice, so radiality preserves the profile in the divisor sum. This includes scales $X < 1$.

For $Nh \asymp G$, the mask costs $\tau(h) \ll_\varepsilon D^\varepsilon$, each kernel is $O_\varepsilon(D^{1+\varepsilon}/G)$, and there are $O(D^2/G)$ pairs. Thus the off-diagonal normalized contribution satisfies

$$|\mathcal{E}_G^{\mathrm{off}}| \ll_{\Phi,\varepsilon} \frac{D^{2+\varepsilon}}{G^2}.$$

The diagonal is $O_\Phi(H)$ for $H \ge 1$. At $H = D^{3/4}$, the dyadic range $G \ge D^{5/8}$ is therefore controlled by $O_{\Phi,\varepsilon}(HD^\varepsilon)$. The small-gcd range, especially $h = 1$, remains open. This result is for one fixed smooth radial profile; it does not prove the source's stronger moving-profile or supremum uniformity. The derivation and surviving alternative routes are recorded in `research/short-row-high-gcd.md`.

31.3 Formal coverage and the next research step

The new Lean interface checks actual Plücker polynomials, nodal matrix relations and determinants, finite clock/shift identities and Kraus positivity/trace preservation. The short-row module checks finite-character masks, original finite Gram expansion, matched-row lower bounds and actual Eisenstein norm/evaluation identities. Full planar Poisson summation, ideal-character interfaces, the AR classifications, the derived functors and the complete physics program are not all formalized by these components. Their theorem types and axiom dependencies are printed in the verification logs.

The next load-bearing research task is an estimate for the remaining low-gcd original arithmetic kernel, or an independently constructed positive arithmetic determinant family escaping the precisely stated finite-kernel obstructions. A learned ToricGT-style component may propose a family, choose a search ordering or discover a candidate identity. Acceptance still requires exact algebra, analytic estimates, Lean interfaces where available and independent critical review. A fitted zero picture or a small training loss is an observation, not the desired proof.

Bibliography

[1] A. Schreiber, *Surface Algebras I: Dessins D'enfants, Surface Algebras, and Dessin Orders*, arXiv:1810.06750v5 (2023).

[2] A. Schreiber, *Surface Algebras and Surface Orders II: Affine Bundles on Curves*, arXiv:1812.00621v1; record 2018, downloaded PDF internally dated November 27, 2024.

[3] A. Schreiber, *Multiparameter Persistent Homology—Generic Structures and Quantum Computing*, arXiv:2210.11433v1 (2022).

[4] C. Bessenrodt and T. Holm, *Weighted locally gentle quivers and Cartan matrices*, arXiv:math/0511610v1 (2005).

[5] S. Ford, A. Oswald and J.J. Zhang, *Homological conditions on locally gentle algebras*, arXiv:2409.08333v3 (5 September 2026).

[6] S. Opper, P.-G. Plamondon and S. Schroll, *A geometric model for the derived category of gentle algebras*, arXiv:1801.09659v7 (10 July 2025).

[7] W. Chang, *Geometric models for the algebraic hearts in the derived category of a gentle algebra*, arXiv:2305.03998v2 (4 September 2024).

[8] Y.-Z. Liu, C. Zhang and H. Zhang, *Constructing projective resolution and taking cohomology for gentle algebras in the geometric model*, arXiv:2308.07220v1 (2023).

[9] Z. Li, Y. Qiu and Y. Zhou, *A geometric realization of Koszul duality for graded gentle algebras*, arXiv:2403.15190v1 (2024).

[10] W. Chang, H. Jin and S. Schroll, *Recollements of partially wrapped Fukaya categories and surface cuts*, arXiv:2206.11196v2 (2023).

[11] B. Dequêne and M. Schoonheere, *Resolving subcategories for gentle algebras I: Monogeneous resolving subcategories for gentle trees*, arXiv:2502.20994v3 (2026).

[12] D.A. Buchsbaum and D. Eisenbud, *What makes a complex exact?*, J. Algebra 25 (1973), 259–268, primary paper.

[13] W. Bruns, *The existence of generic free resolutions and related objects*, Math. Scand. 55 (1984), 33–46.

[14] D. Eisenbud, *Homological algebra on a complete intersection, with an application to group representations*, Trans. AMS 260 (1980), 35–64.

[15] A. Beauville and Y. Laszlo, *Un lemme de descente*, C. R. Acad. Sci. Paris 320 (1995), 335–340, author text.

[16] P. Magyar, *Affine Schubert varieties and circular complexes*, arXiv:math/0210151v1 (2002).

[17] F.F. Knudsen and D. Mumford, *The projectivity of the moduli space of stable curves I: Preliminaries on 'det' and 'Div'*, Math. Scand. 39 (1976), 19–55, doi:10.7146/math.scand.a-11642.

[18] M. Breuning, *Determinant functors on triangulated categories*, arXiv:math/0610435v1 (2006).

[19] D. Eriksson and G. Freixas i Montplet, *Deligne–Riemann–Roch and intersection bundles*, J. Éc. polytech. Math. 11 (2024), 247–361, doi:10.5802/jep.254.

[20] P. Deligne, *La conjecture de Weil I*, Publ. Math. IHÉS 43 (1974), 273–307, primary text.

[21] P. Deligne, *La conjecture de Weil II*, Publ. Math. IHÉS 52 (1980), 137–252, primary text.

[22] P. Deligne, *Cohomologie étale (SGA* $4\frac{1}{2}$*)*, Lecture Notes in Mathematics 569, Springer-Verlag, 1977, doi:10.1007/BFb0091516.

[23] G. Tabuada, *Noncommutative Weil conjecture*, arXiv:1808.00950v7 (6 December 2019).

[24] C. Deninger, *Arithmetic geometry and analysis on foliated spaces*, arXiv:math/0505354v1 (2005).

[25] D.B. Ray and I.M. Singer, *R-Torsion and the Laplacian on Riemannian manifolds*, Advances in Math. 7 (1971), 145–210, doi:10.1016/0001-8708(71)90045-4.

[26] F. Bornemann, *On the numerical evaluation of Fredholm determinants*, arXiv:0804.2543v1; Math. Comp. 79 (2010), 871–915.

[27] B. Dequêne and M. Schoonheere, *Resolving subcategories for gentle algebras II: Resolving subcategories for gentle trees*, arXiv:2509.10120v2 (3 October 2025).

[28] C. Procesi, *The invariant theory of* $n \times n$ *matrices*, Advances in Math. 19 (1976), 306–381, doi:10.1016/0001-8708(76)90027-X.

[29] L. Le Bruyn and C. Procesi, *Semisimple representations of quivers*, Trans. Amer. Math. Soc. 317 (1990), 585–598, doi:10.1090/S0002-9947-1990-0958897-0.

[30] A. Craw and R. Yamagishi, *The Le Bruyn–Procesi theorem following Lusztig*, arXiv:2312.08527v4 (14 November 2025).

[31] P. Etingof, O. Golberg, S. Hensel, T. Liu, A. Schwendner, D. Vaintrob and E. Yudovina, *Introduction to Representation Theory*, with historical interludes by S. Gerovitch, Student Mathematical Library 59, American Mathematical Society, 2011, doi:10.1090/stml/059; author-hosted text.

[32] K. Akin, D.A. Buchsbaum and J. Weyman, *Schur functors and Schur complexes*, Advances in Math. 44 (1982), 207–278, doi:10.1016/0001-8708(82)90039-1.

[33] Q. Jiang, *Derived Grassmannians and derived Schur functors*, arXiv:2212.10488v3 (6 July 2023).

[34] A. Price, A. Stelzer and A. Yong, *Representations from matrix varieties, and filtered RSK*, arXiv:2403.09938v2 (6 October 2025).

[35] H. Derksen and J. Weyman, *Semi-invariants of quivers and saturation for Littlewood–Richardson coefficients*, J. Amer. Math. Soc. 13 (2000), 467–479, primary author text.

[36] A. Schofield and M. Van den Bergh, *Semi-invariants of quivers for arbitrary dimension vectors*, arXiv:math/9907174v1 (27 July 1999); Indag. Math. 12 (2001), 125–138.

[37] J.S. Milne, *Reductive Groups*, author's course notes, version 2.00 (10 March 2018), §5.18.

[38] W. Crawley-Boevey, *Classification of modules for infinite-dimensional string algebras*, Trans. Amer. Math. Soc. 370 (2018), 3289–3313, doi:10.1090/tran/7032; arXiv:1308.6410v2 (4 March 2016).

[39] M. Auslander and I. Reiten, *Almost split sequences for Cohen–Macaulay-modules*, Math. Ann. 277 (1987), 345–349, doi:10.1007/BF01457367.

[40] R. Bautista, S. Liu and C. Paquette, *Representation Theory of an Infinite Quiver*, arXiv:1109.3176v1 (14 September 2011); published as *Representation theory of strongly locally finite quivers*, Proc. London Math. Soc. 106 (2013), 97–162.

[41] C. Paquette, *Irreducible morphisms and locally finite dimensional representations*, arXiv:1508.04353v3 (4 January 2016).

[42] S. Liu, P. Ng and C. Paquette, *Almost split sequences and approximations*, arXiv:1203.5009v1 (19 March 2012).

[43] C.M. Ringel, *The Minimal Representation-Infinite Algebras which are Special Biserial*, author-hosted ICRA Tokyo 2010 proceedings text, 57 pages, retrieved 3 November 2026.

[44] G. Puninski and M. Prest, *Ringel's conjecture for domestic string algebras*, arXiv:1407.7470v2 (13 September 2015).

[45] R. Laking, M. Prest and G. Puninski, *Krull–Gabriel dimension of domestic string algebras*, arXiv:1506.07703v2 (29 June 2015).

[46] H.-L. Huang, Z. Lin and X. Su, *The Auslander–Reiten quivers of string algebras of affine type $\widetilde{C}$ and a conjecture by Geiss–Leclerc–Schröer*, arXiv:2110.02777v2 (16 November 2021).

[47] G. Abrams and G. Aranda Pino, *The Leavitt path algebra of a graph*, arXiv:math/0509494v1 (21 September 2005).

[48] T. Bates, D. Pask, I. Raeburn and W. Szymański, *The C^*-algebras of row-finite graphs*, New York J. Math. 6 (2000), 307–324, primary journal text.

[49] N.J. Fowler and I. Raeburn, *The Toeplitz algebra of a Hilbert bimodule*, arXiv:math/9806093v3 (23 June 1998).

[50] N. Brownlowe, M. Laca, D. Robertson and A. Sims, *Reconstructing directed graphs from generalised gauge actions on their Toeplitz algebras*, arXiv:1812.08903v1 (21 December 2018).

[51] E.T.A. Kakariadis and O.M. Shalit, *On operator algebras associated with monomial ideals in noncommuting variables*, arXiv:1501.06495v3 (21 November 2018).

[52] O.M. Shalit and B. Solel, *Subproduct systems*, arXiv:0901.1422v3 (26 May 2009).

[53] A. Connes and C. Consani, *Knots, primes and class field theory*, arXiv:2501.06560v1 (11 January 2025).

[54] A. Connes, C. Consani and H. Moscovici, *Zeta zeros and prolate wave operators*, arXiv:2310.18423v2 (4 May 2024).

[55] A. Connes, C. Consani and H. Moscovici, *On q-series and the moment problem associated to local factors*, arXiv:2403.01247v1 (2 March 2024).

[56] A. Connes, C. Consani and H. Moscovici, *Zeta Spectral Triples*, arXiv:2511.22755v1 (27 November 2025).

[57] A. Connes and W.D. van Suijlekom, *Quadratic Forms, Real Zeros and Echoes of the Spectral Action*, arXiv:2511.23257v1 (28 November 2025).

[58] A. Connes, *The Riemann Hypothesis: Past, Present and a Letter Through Time*, arXiv:2602.04022v1 (3 February 2026).

[59] A. Connes and C. Consani, *On the Jacobian of* $\overline{\mathrm{Spec}\,\mathbb{Z}}$, arXiv:2602.15941v1 (17 February 2026).

[60] C. Deninger, *Dynamical systems for arithmetic schemes*, arXiv:1807.06400v4 (7 February 2024); Indag. Math. 37 (2026), 25–136, doi:10.1016/j.indag.2024.05.007.

[61] C. Deninger, *Primes, knots and periodic orbits*, arXiv:2301.11643v1 (27 January 2023).

[62] C. Deninger, *There is no "Weil-"cohomology theory with real coefficients for arithmetic curves*, arXiv:2204.02714v2 (27 January 2023).

[63] R. Ponge and Y. Tian, *Nonclassical Weyl laws and Connes' Integration for weak Lorentz ideals, I*, arXiv:2605.25591v2 (28 September 2026).

[64] F.W.J. Olver et al. (eds.), *NIST Digital Library of Mathematical Functions*, finite summation by parts, Eq. 2.10.9 and Dirichlet L-functions, §25.15, consulted 3 November 2026.

[65] A. Chirvasitu, P.M. Hajac and M. Tobolski, *Non-surjective pullbacks of graph C*-algebras from non-injective pushouts of graphs*, arXiv:1907.10260v3 (31 December 2019), Theorems3.4 and4.3; doi:10.1112/blms.12389.

[66] Y. Lekili and A. Polishchuk, *Derived equivalences of gentle algebras via Fukaya categories*, arXiv:1801.06370v5 (27 August 2019), Definition3.1.1, Lemma3.1.3 and Theorem3.2.2.

[67] M. Tomforde, *Uniqueness theorems and ideal structure for Leavitt path algebras*, arXiv:math/0612628v6 (27 February 2007), Theorem7.3, pp.29–30.

[68] A. Weil, *Numbers of solutions of equations in finite fields*, Bull. Amer. Math. Soc.55 (1949),497–508, especially pp.499–502 and Eq.(4), doi:10.1090/S0002-9904-1949-09219-4; primary-paper scan.

[69] The Stacks Project, Tag0FJ9 and Tag0GSX (retrieved 3 November 2026).

[70] J. Alper, *Good moduli spaces for Artin stacks*, Ann. Inst. Fourier63 (2013),2349–2402, Theorem10.3, pp.2386–2387, doi:10.5802/aif.2833.

[71] The Stacks Project, Tag04V8: Taking roots of line bundles (retrieved 3 November 2026).

[72] G. Zwara, *Smooth morphisms of module schemes*, Proc. London Math. Soc. (3) 84 (2002), 539–558, primary publication.

[73] C. Riedtmann and G. Zwara, *Orbit closures and rank schemes*, Comment. Math. Helv. 88 (2013), 55–84, primary publication, Section6.

[74] Y. Liu, N. Li, B. Xing and P. Chen, *A lecture note on covering theory in representation theory of algebras*, arXiv:2508.01563v5, 28 May 2026.

[75] Q. Dong, Z. Lin, M. Lu and S. Ruan, *Representation-tame Geiß–Leclerc–Schröer algebras and a revised GLS conjecture on root systems*, arXiv:2609.26503v1, 22 September 2026.

[76] S. Schroll, *On geometric models in representation theory*, Proceedings of ICM 2026, vol.3, pp.250–271, doi:10.1137/25M1810647; author preprint arXiv:2601.14396v1.

[77] H.-L. Huang, Z. Lin and X. Su, *The Auslander–Reiten quivers of string algebras of affine type $\widetilde{C}$ and a conjecture by Geiss–Leclerc–Schröer*, arXiv:2110.02777v2, 16 November 2021, especially Propositions 2.9 and 3.12, Lemma 3.10 and Theorem 3.14.

[78] C. M. Ringel, *The Minimal Representation-Infinite Algebras which are Special Biserial*, author-hosted ICRA Tokyo 2010 proceedings text, 57 pages, retrieved 3 November 2026; especially Sections 12.5 and 13.1–13.7. The unversioned retrieval does not determine an additional revision date.

[79] G. Zwara, *Smooth morphisms of module schemes*, Proc. London Math. Soc. (3) **84** (2002), 539–558, doi:10.1112/S0024611502013382.

[80] C. Riedtmann and G. Zwara, *Orbit closures and rank schemes*, Comment. Math. Helv. **88** (2013), 55–84, doi:10.4171/CMH/278, Section 6, especially the definition on p.71 and Theorems 6.1–6.2.

[81] S. Opper, P.-G. Plamondon and S. Schroll, *A geometric model for the derived category of gentle algebras*, arXiv:1801.09659v7, 10 July 2025, especially the category conventions and Section 5.1.

[82] A. Schofield and M. Van den Bergh, *Semi-invariants of quivers for arbitrary dimension vectors*, arXiv:math/9907174v1 (1999), Trans. Amer. Math. Soc. **353** (2001), 1631–1650.

[83] W. Crawley-Boevey, *Classification of modules for infinite-dimensional string algebras*, arXiv:1308.6410v2, 4 March 2016, Trans. Amer. Math. Soc. **370** (2018), 3289–3313.

[84] S. Fomin and A. Zelevinsky, *Total positivity: tests and parametrizations*, Math. Intelligencer **22** (2000), 23–33; arXiv:math/9912128.

[85] K. Osterwalder and R. Schrader, *Axioms for Euclidean Green's functions*, Commun. Math. Phys. **31** (1973), 83–112; doi:10.1007/BF01645738.

[86] G. A. Elliott and D. E. Evans, *The structure of the irrational rotation C^*-algebra*, Ann. Math. **138** (1993), 477–501; doi:10.2307/2946556.

[87] T. Jacobson, *Entanglement equilibrium and the Einstein equation*, Phys. Rev. Lett. **116** (2016), 201101; arXiv:1505.04753.

[88] W. F. Stinespring, *Positive functions on C^*-algebras*, Proc. Amer. Math. Soc. **6** (1955), 211–216.

[89] I. Cong, S. Choi and M. D. Lukin, *Quantum convolutional neural networks*, Nature Physics **15** (2019), 1273–1278; arXiv:1810.03787.

[90] J. Zheng, Q. Gao and Y. Lv, *Quantum graph convolutional neural networks*, arXiv:2107.03257 (2021).

[91] A. Ning, T. Y. Li and N. Y. Chen, *Quantum graph attention network: A novel quantum multi-head attention mechanism for graph learning*, arXiv:2508.17630 (2025).

[92] H. Ramsauer et al., *Hopfield networks is all you need*, ICLR 2021; arXiv:2008.02217.

[93] F. Pastawski, B. Yoshida, D. Harlow and J. Preskill, *Holographic quantum error-correcting codes: toy models for the bulk/boundary correspondence*, JHEP **2015** (6), 149; arXiv:1503.06237.

[94] H.-Y. Huang, R. Kueng and J. Preskill, *Predicting many properties of a quantum system from very few measurements*, Nature Physics **16** (2020), 1050–1057; arXiv:2002.08953.

[95] L. E. Rosaleny et al., *Peptides as versatile scaffolds for quantum computing*, arXiv:1708.09440 (2017), including supplementary information S2. Coherence observations cited here refer to the frozen-solution EPR measurements, not a calibrated QPU.

[96] M. Nitz, M. Sherawat, K. J. Franz, E. Peisach, K. N. Allen and B. Imperiali, *Structural origin of the high affinity of a chemically evolved lanthanide-binding peptide*, Angew. Chem. Int. Ed. **43** (2004), 3682–3685; doi:10.1002/anie.200460028. Structure: PDB 1TJB.

[97] J. Butcher et al., *De novo design of all-atom biomolecular interactions with RFdiffusion3*, bioRxiv 2025.09.18.676967 (2025); primary preprint. Implementation documentation: Rosetta-Commons Foundry RFD3.

[98] A. Schreiber, *ThermoGFN-IF for catalysis: a protein sequence design model tuned with GFlowNets for stable protein design and kinetic-aware enzyme engineering*, author project description (10 March 2026), project description; implementation scaffold. Cited as a proposed methodology, not completed molecular or qubit validation.

[99] *Generalizable and scalable protein stability prediction with rewired protein generative models*, Nature Communications (2026), primary article, doi:10.1038/s41467-025-67609-4; correction doi:10.1038/s41467-026-70929-8.

[100] S. Lewis et al., *Scalable emulation of protein equilibrium ensembles with generative deep learning*, Science (2025), eadv9817; doi:10.1126/science.adv9817. Scope of the cited v1 model: official BioEmu model card.

[101] B. M. Wood et al., *UMA: A family of universal models for atoms*, arXiv:2506.23971 (2025).

[102] N. P. Breuckmann and B. M. Terhal, *Constructions and noise threshold of hyperbolic surface codes*, IEEE Trans. Inform. Theory **62** (2016), 3731–3744; arXiv:1506.04029v2.

[103] N. Delfosse, *Tradeoffs for reliable quantum information storage in surface codes and color codes*, IEEE ISIT (2013), 917–921; arXiv:1301.6588.

[104] A. Schreiber, *ToricGT*, repository revision a5199bc, inspected 3 November 2026; README, graph-token schemas, toric-algebra and finite-certificate modules.

[105] A. Schreiber, *ToricGT Research Program*, repository revision f91ba18, inspected 3 November 2026, README and five-paper program outline.

[106] M. Kontsevich and Y. Soibelman, *Homological mirror symmetry and torus fibrations*, arXiv:math/0011041v2 (4 June 2001), Sections 6–8, especially Theorem 4.

[107] M. Abouzaid, *Morse Homology, Tropical Geometry, and Homological Mirror Symmetry for Toric Varieties*, arXiv:math/0610004v4 (21 April 2009), Theorem 1.2 and Corollary 1.4.

[108] A. Schreiber, *Surface Algebras I: Dessins D'enfants, Surface Algebras, and Dessin Orders*, arXiv:1810.06750v5 (24 February 2023), especially the surface-algebra/order constructions and Section 10 on projective resolutions.

[109] A. Schreiber, *Surface Algebras and Surface Orders II: Affine Bundles on Curves*, arXiv:1812.00621v1, record dated 3 December 2018; cached PDF internally dated 27 November 2024, especially Sections 5–7 and 9–11. The retrieval pin is recorded in the research notes.

[110] A. Schreiber, *Multiparameter Persistent Homology—Generic Structures and Quantum Computing*, arXiv:2210.11433v1 (20 October 2022), especially Sections 3–7 and 10. The chapter uses corrected discrete grading, exactness and specialization hypotheses rather than every literal survey assertion.

[111] D. A. Buchsbaum and D. Eisenbud, *What makes a complex exact?*, J. Algebra **25** (1973), 259–268, primary paper, Theorem and Corollaries 1–2, pp.260–261.

[112] W. Bruns, *The existence of generic free resolutions and related objects*, Math. Scand. **55** (1984), 33–46, author-hosted primary paper, Theorem 1 and direct-limit construction, pp.35–36.